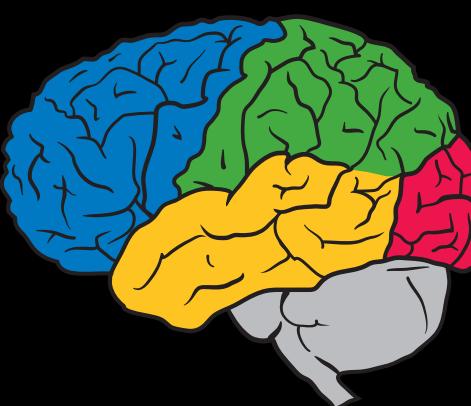


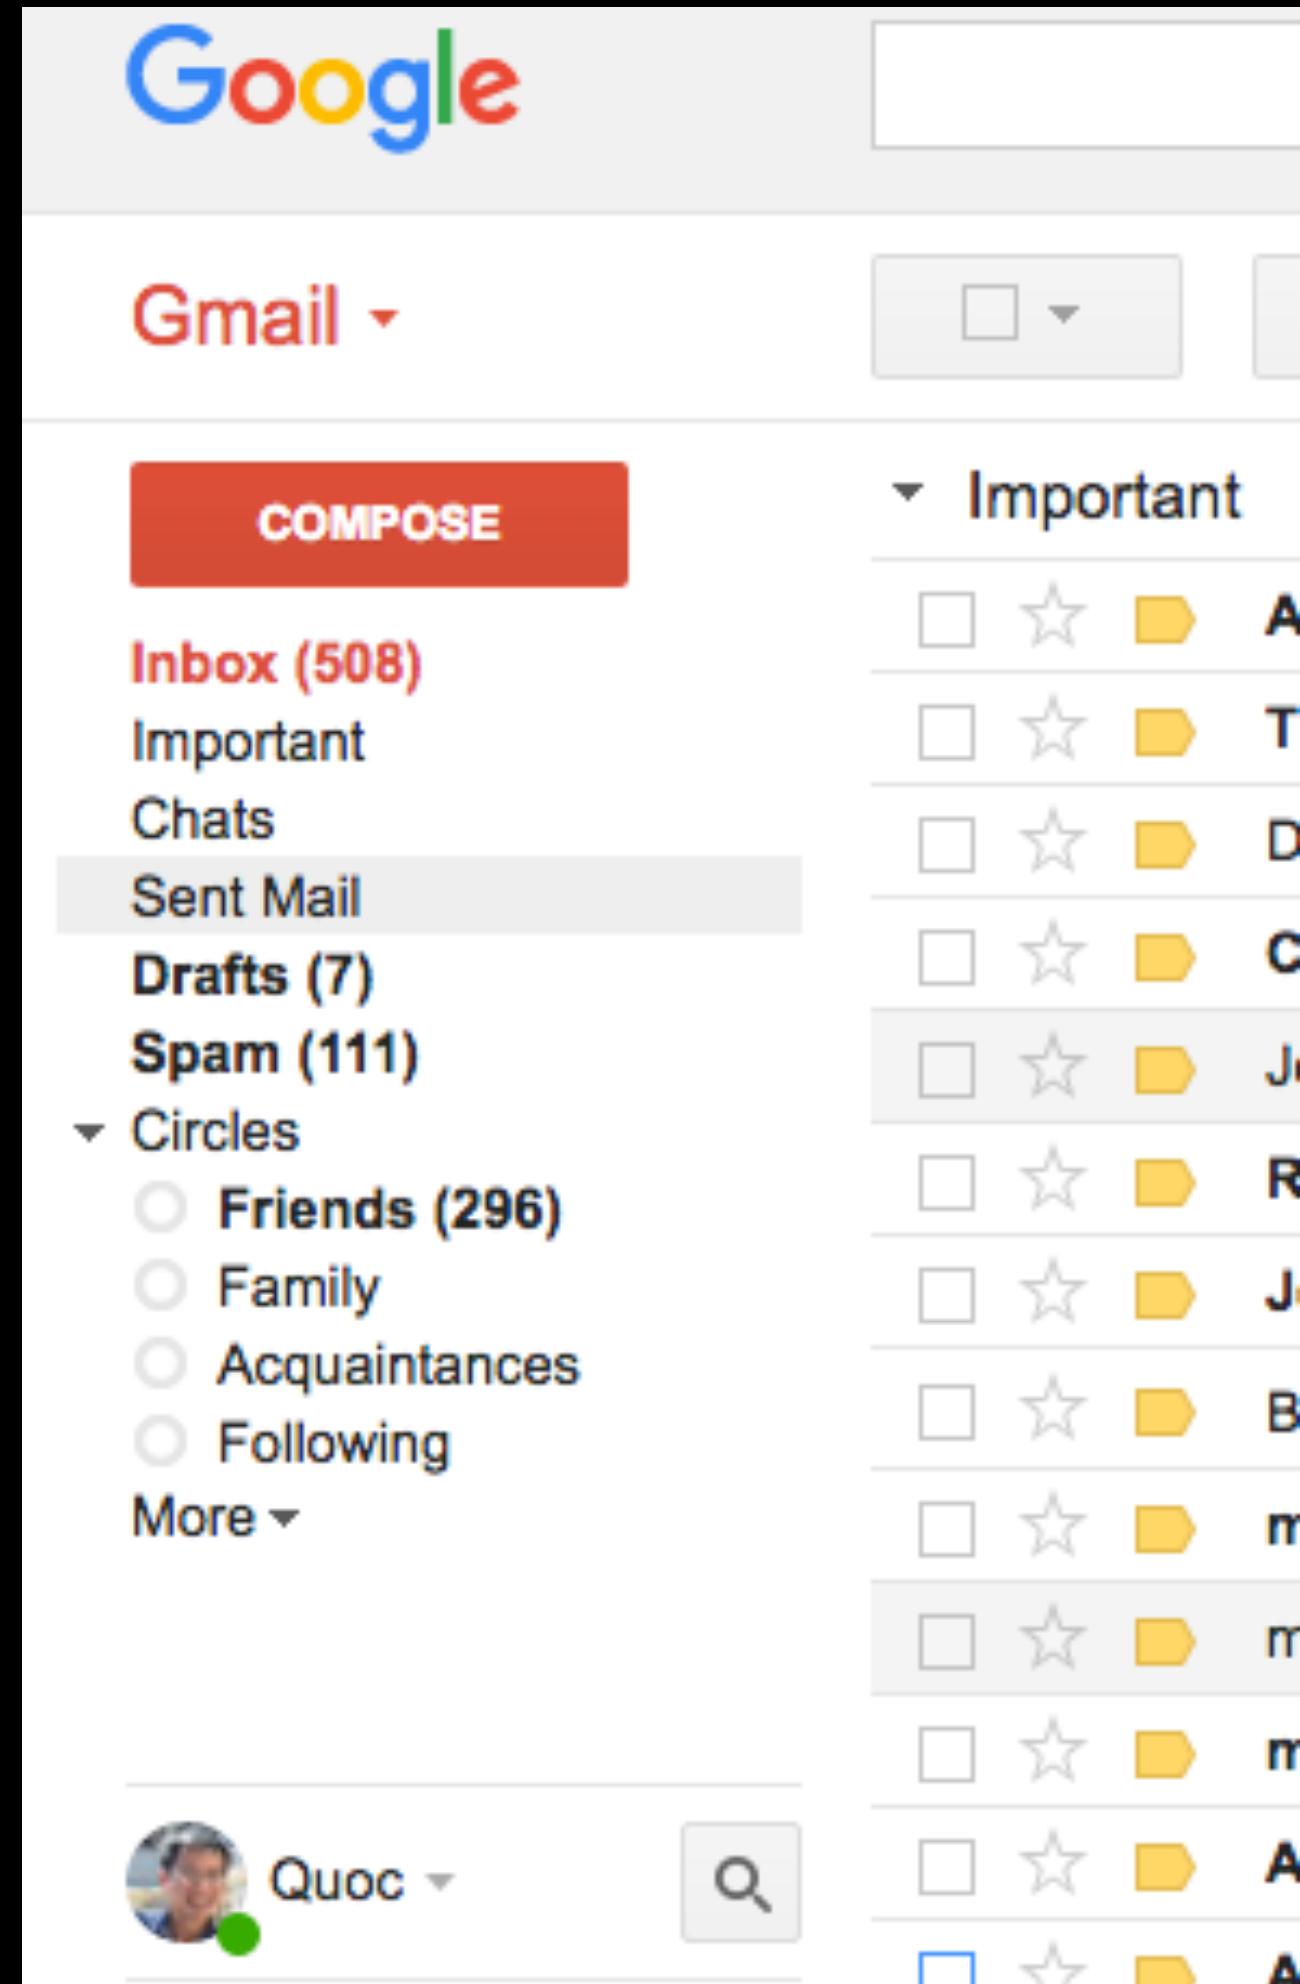
Sequence to Sequence Learning for NLP and Speech

Quoc V. Le

Google Brain team



“AutoReply”



- 508 unread emails!!!
- Some emails just require “Yes” / “No” answers
- Let’s build “AutoReply”

“AutoReply”

- From: Ann
- Subject: Hi
- Content: Are you visiting Vietnam for the new year, Quoc?
- Probable Reply: Yes

Dataset

- Are you visiting Vietnam for the new year, Quoc? -> Yes
- Are you hanging out with us tonight? -> No
- Did you read the cool paper on ResNet? -> Yes
- ...

Preprocessing

- Are you visiting Vietnam for the new year , Quoc ? -> Yes
- Are you hanging out with us tonight ? -> No
- Did you read the cool paper on ResNet ? -> Yes
- ...

Preprocessing

- Are you visiting Vietnam for the new year , Quoc ? -> Yes
- Are you hanging out with us tonight ? -> No
- Did you read the cool paper on ResNet ? -> Yes
- ...

Feature Representation

Are you visiting Vietnam for the new year , Quoc ?

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 0, 2]



20,000 dimensions

Feature Representation

Are you visiting Vietnam for the new year , Quoc ?

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, . . . , 0, 0, 1, 0, 0, 0, 0, 2]



20,000 dimensions

Feature Representation

Are you visiting Vietnam for the new year , Quoc ?

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 2]



20,000 dimensions

Feature Representation

Are you visiting Vietnam for the new year , Quoc ?

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 2]



20,000 dimensions

Feature Representation

Are you visiting Vietnam for the new year , Quoc ?

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 2]

20,000 dimensions

**Special dimension
reserved for out
of vocabulary words**

Formulation

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 2] -> 1

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 1, 0, 0, 0, 0, 0] -> 0

[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, ..., 0, 3, 0, 0, 0, 0, 1] -> 1

Formulation

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 2] -> 1

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 1, 0, 0, 0, 0, 0] -> 0

[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, ..., 0, 3, 0, 0, 0, 0, 1] -> 1



Formulation

- Find \mathbf{W} such that \mathbf{Wx} approximates y
- Since y is in {"Yes", "No"}, this is a "Logistic Regression" problem

$$\begin{bmatrix} \frac{\exp(\mathbf{w}_1^\top \mathbf{x})}{\exp(\mathbf{w}_1^\top \mathbf{x}) + \exp(\mathbf{w}_2^\top \mathbf{x})} \\ \frac{\exp(\mathbf{w}_2^\top \mathbf{x})}{\exp(\mathbf{w}_1^\top \mathbf{x}) + \exp(\mathbf{w}_2^\top \mathbf{x})} \end{bmatrix}$$

Formulation

- Find \mathbf{W} such that \mathbf{Wx} approximates y
- Since y is in {"Yes", "No"}, this is a "Logistic Regression" problem

$$\begin{bmatrix} \frac{\exp(\mathbf{w}_1^T \mathbf{x})}{\exp(\mathbf{w}_1^T \mathbf{x}) + \exp(\mathbf{w}_2^T \mathbf{x})} \\ \frac{\exp(\mathbf{w}_2^T \mathbf{x})}{\exp(\mathbf{w}_1^T \mathbf{x}) + \exp(\mathbf{w}_2^T \mathbf{x})} \end{bmatrix}$$

Positive and sum up to 1

Training with stochastic gradient descent

- For iteration 1, 2, 3, ..., 1000000
 - Sample a random email \mathbf{x} and a reply
 - If reply == Yes, update w_1 and w_2 to increase
 - If reply == No, update w_1 and w_2 to increase

$$\frac{\exp(w_1^T \mathbf{x})}{\exp(w_1^T \mathbf{x}) + \exp(w_2^T \mathbf{x})}$$

$$\frac{\exp(w_2^T \mathbf{x})}{\exp(w_1^T \mathbf{x}) + \exp(w_2^T \mathbf{x})}$$

Training with stochastic gradient descent

- For iteration 1, 2, 3, ..., 1000000
 - Sample a random email \mathbf{x} and a reply
 - If reply == Yes, update w_1 and w_2 to increase
 - If reply == No, update w_1 and w_2 to increase

$$\frac{\exp(w_1^T \mathbf{x})}{\exp(w_1^T \mathbf{x}) + \exp(w_2^T \mathbf{x})}$$

$$\frac{\exp(w_2^T \mathbf{x})}{\exp(w_1^T \mathbf{x}) + \exp(w_2^T \mathbf{x})}$$

Training with stochastic gradient descent

- For iteration 1, 2, 3, ..., 1000000
 - Sample a random email \mathbf{x} and a reply
 - If reply == Yes, update w_1 and w_2

$$w_1 = w_1 + \text{alpha} \frac{d \log(p_1)}{d w_1}$$

$$w_2 = w_2 + \text{alpha} \frac{d \log(p_1)}{d w_2}$$

- If reply == No, update w_1 and w_2

$$w_1 = w_1 + \text{alpha} \frac{d \log(p_2)}{d w_1}$$

$$w_2 = w_2 + \text{alpha} \frac{d \log(p_2)}{d w_2}$$

Prediction

- For any incoming email \mathbf{x}

- Compute
$$\frac{\exp(\mathbf{w}_1^\top \mathbf{x})}{\exp(\mathbf{w}_1^\top \mathbf{x}) + \exp(\mathbf{w}_2^\top \mathbf{x})}$$

- If $> 0.5 \rightarrow \text{reply} = \text{Yes}$

- If $\leq 0.5 \rightarrow \text{reply} = \text{No}$

Information Loss

Are you visiting Vietnam for the new year , Quoc ?

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ..., 0, 0, 1, 0, 0, 0, 0, 2]



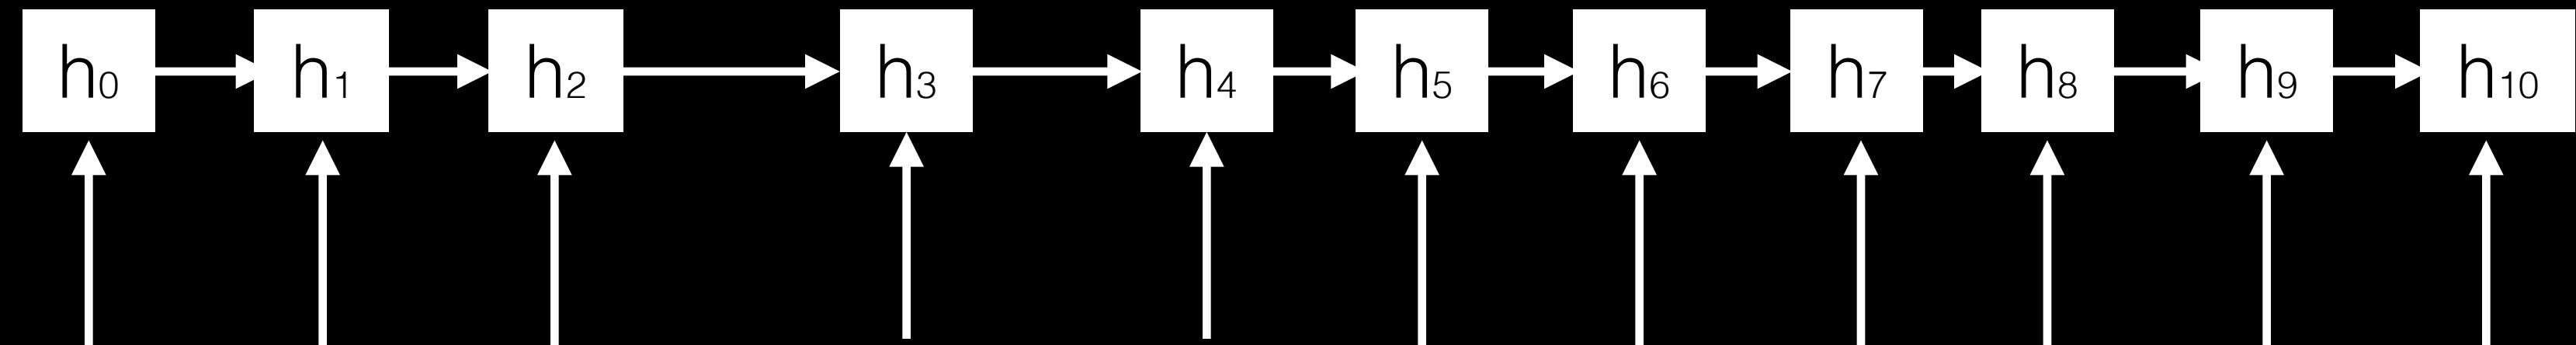
20,000 dimensions

This “bag-of-words representation” does not care about the order of the words!

Recurrent Neural Network

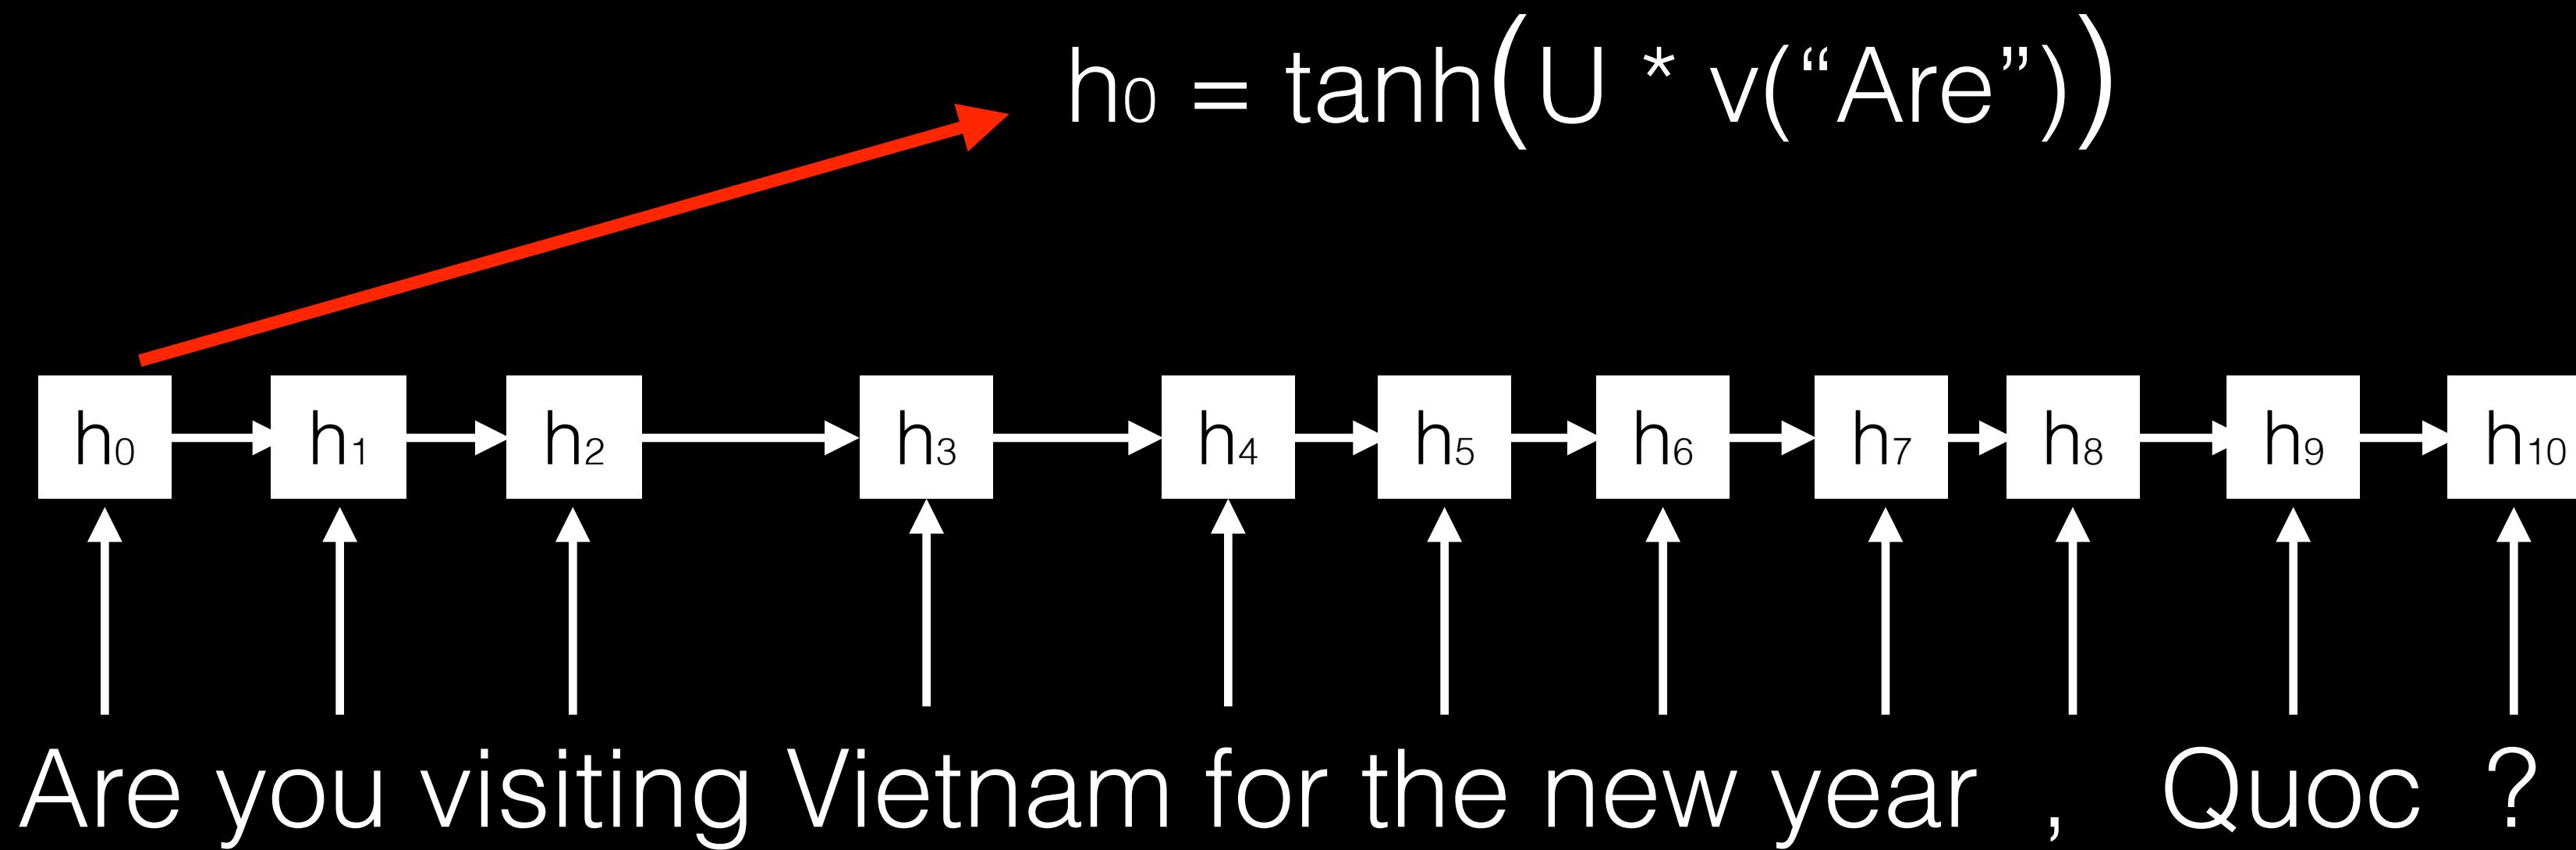
Are you visiting Vietnam for the new year , Quoc ?

Recurrent Neural Network

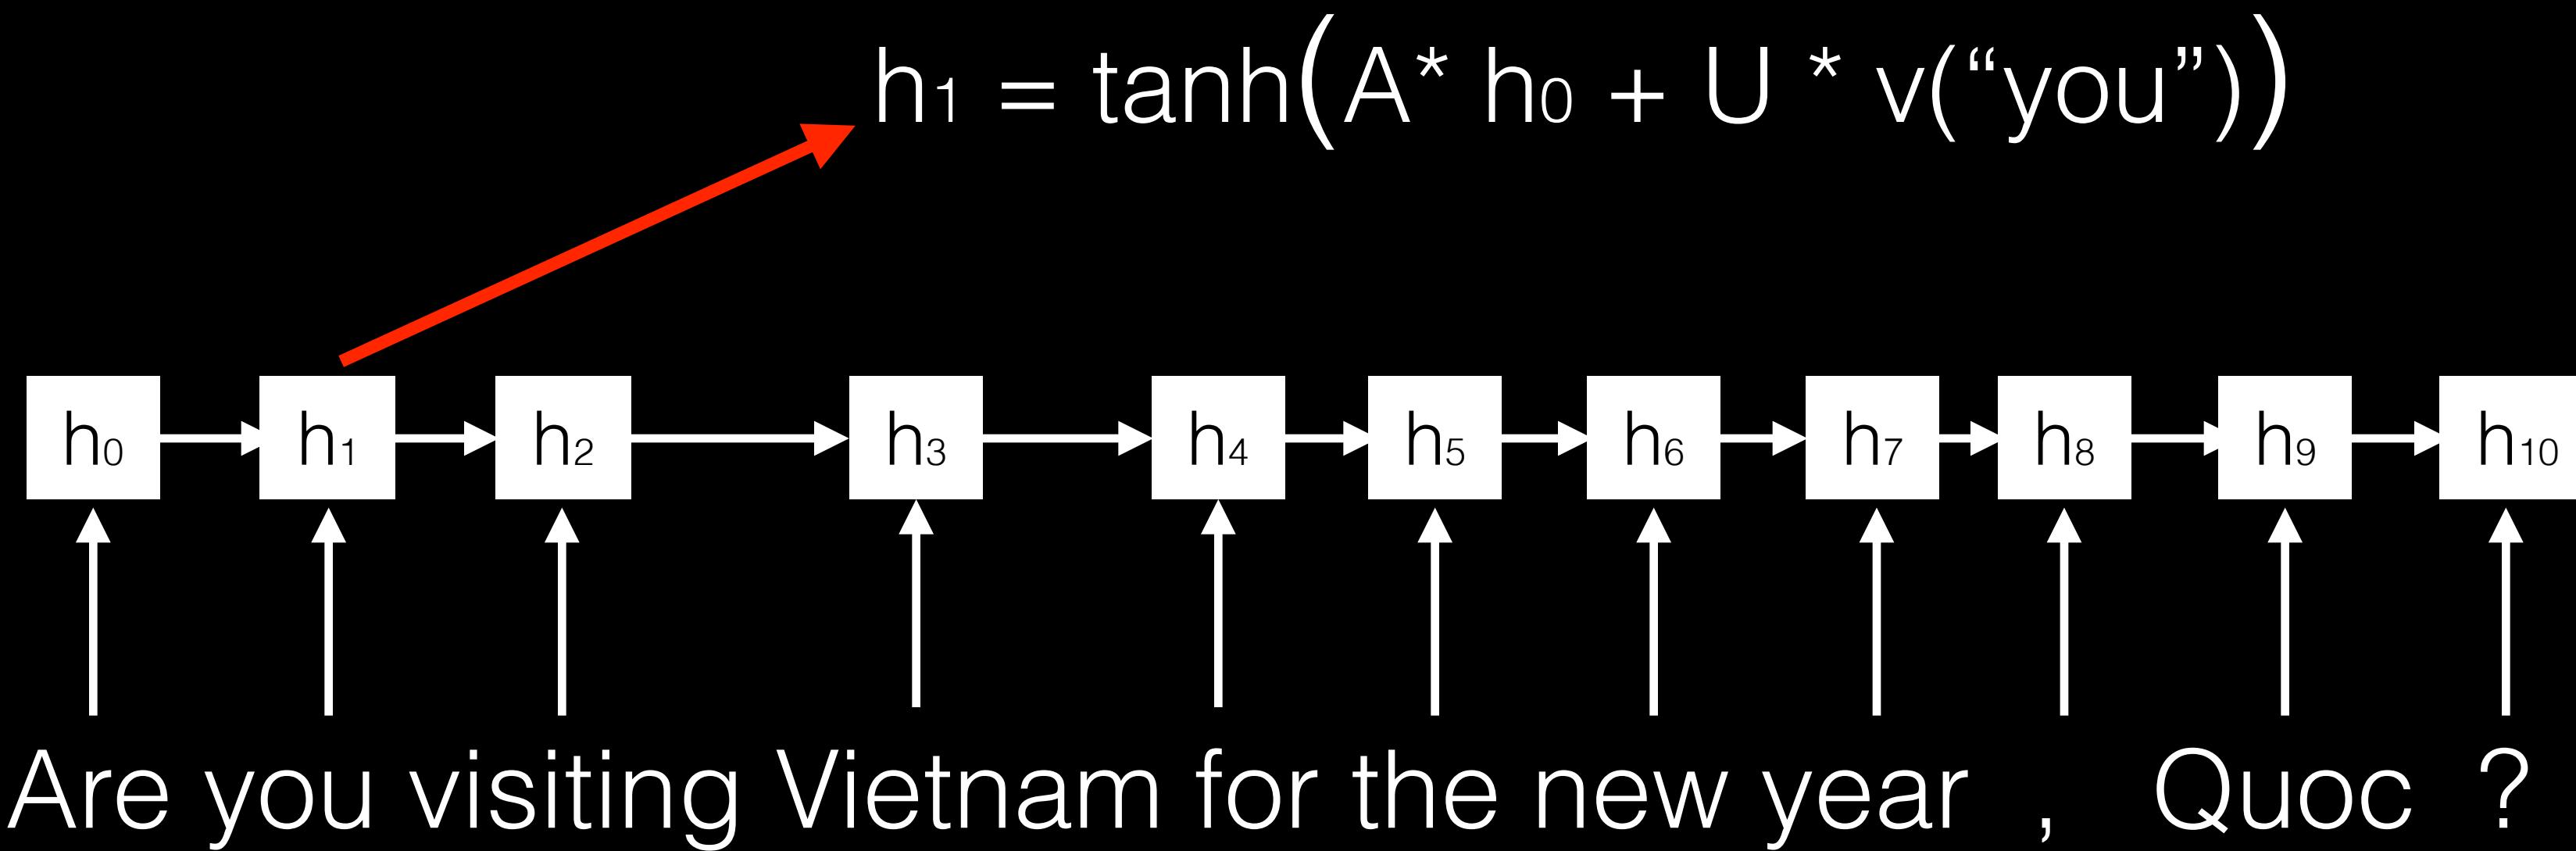


Are you visiting Vietnam for the new year , Quoc ?

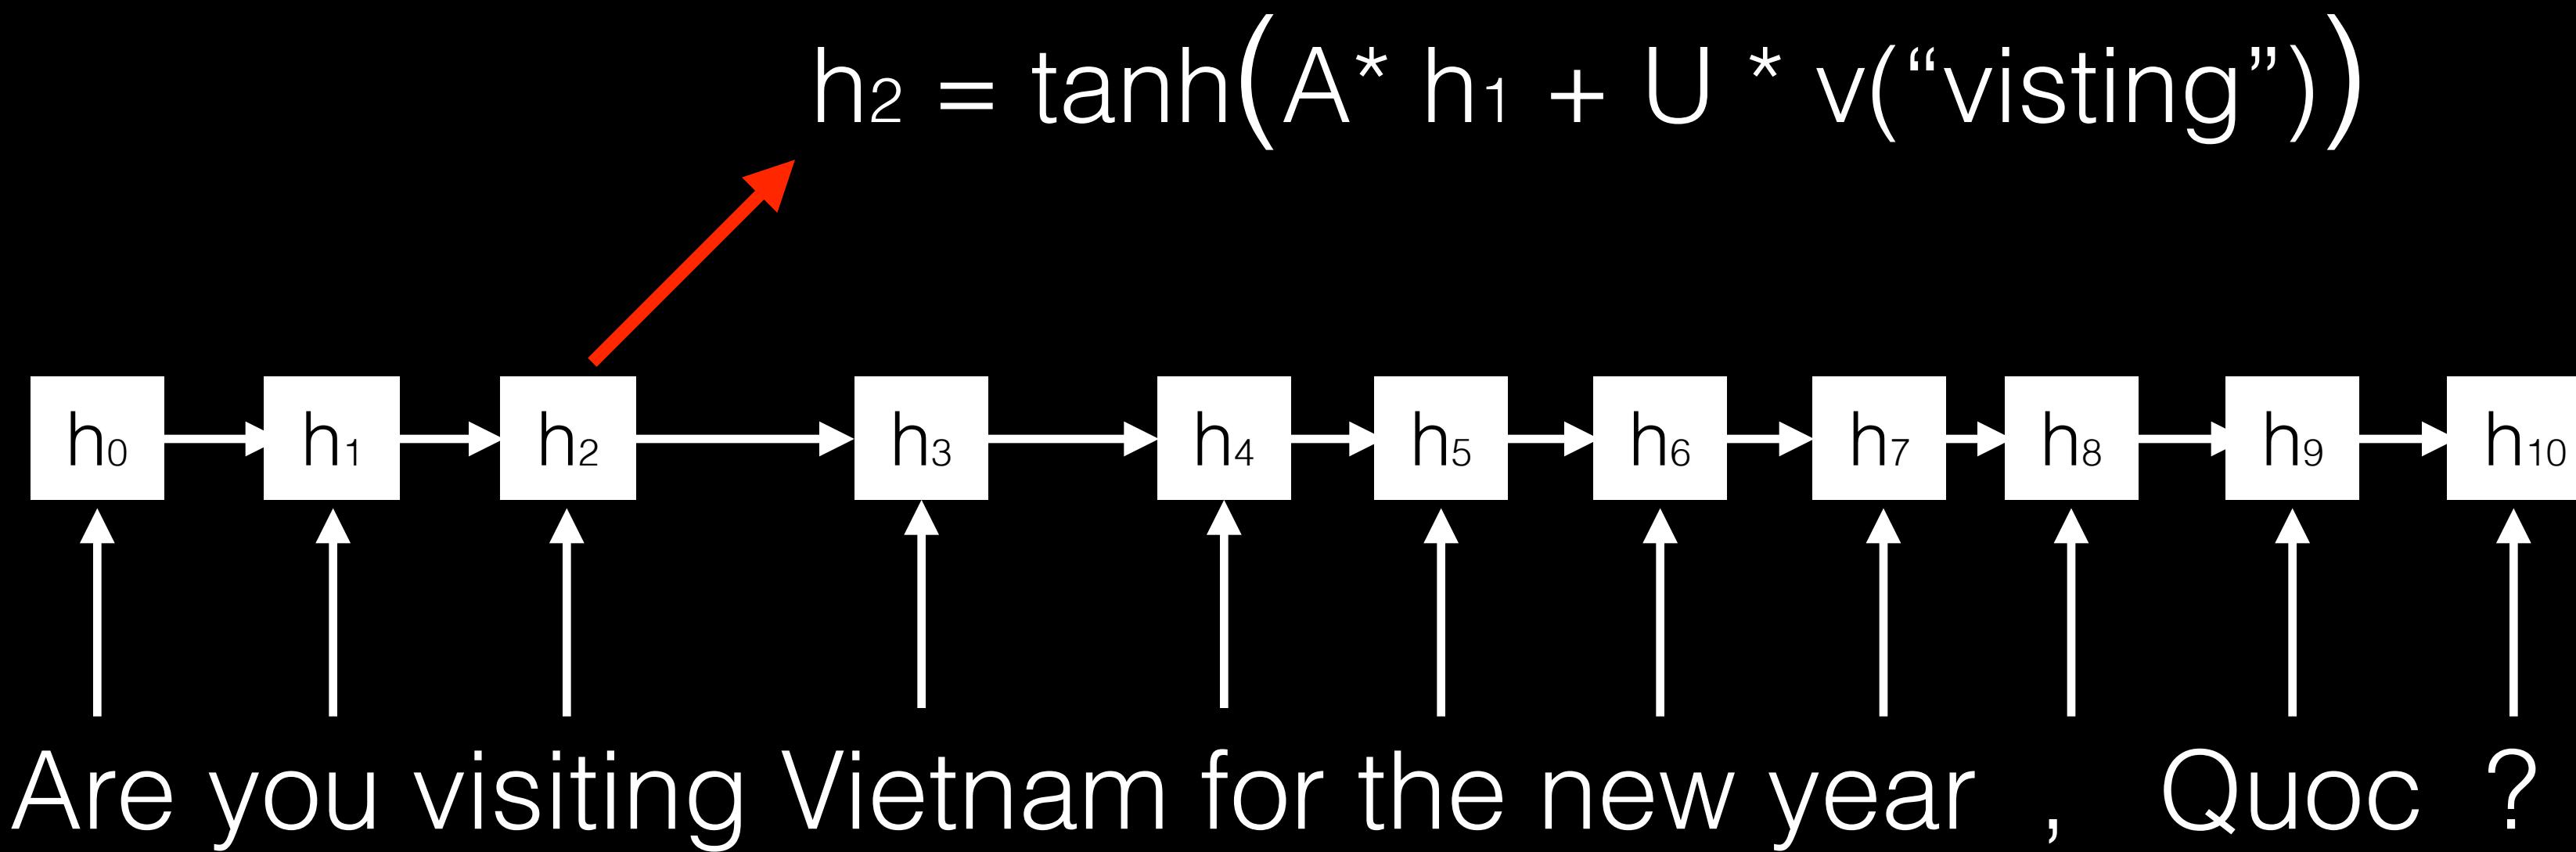
Recurrent Neural Network



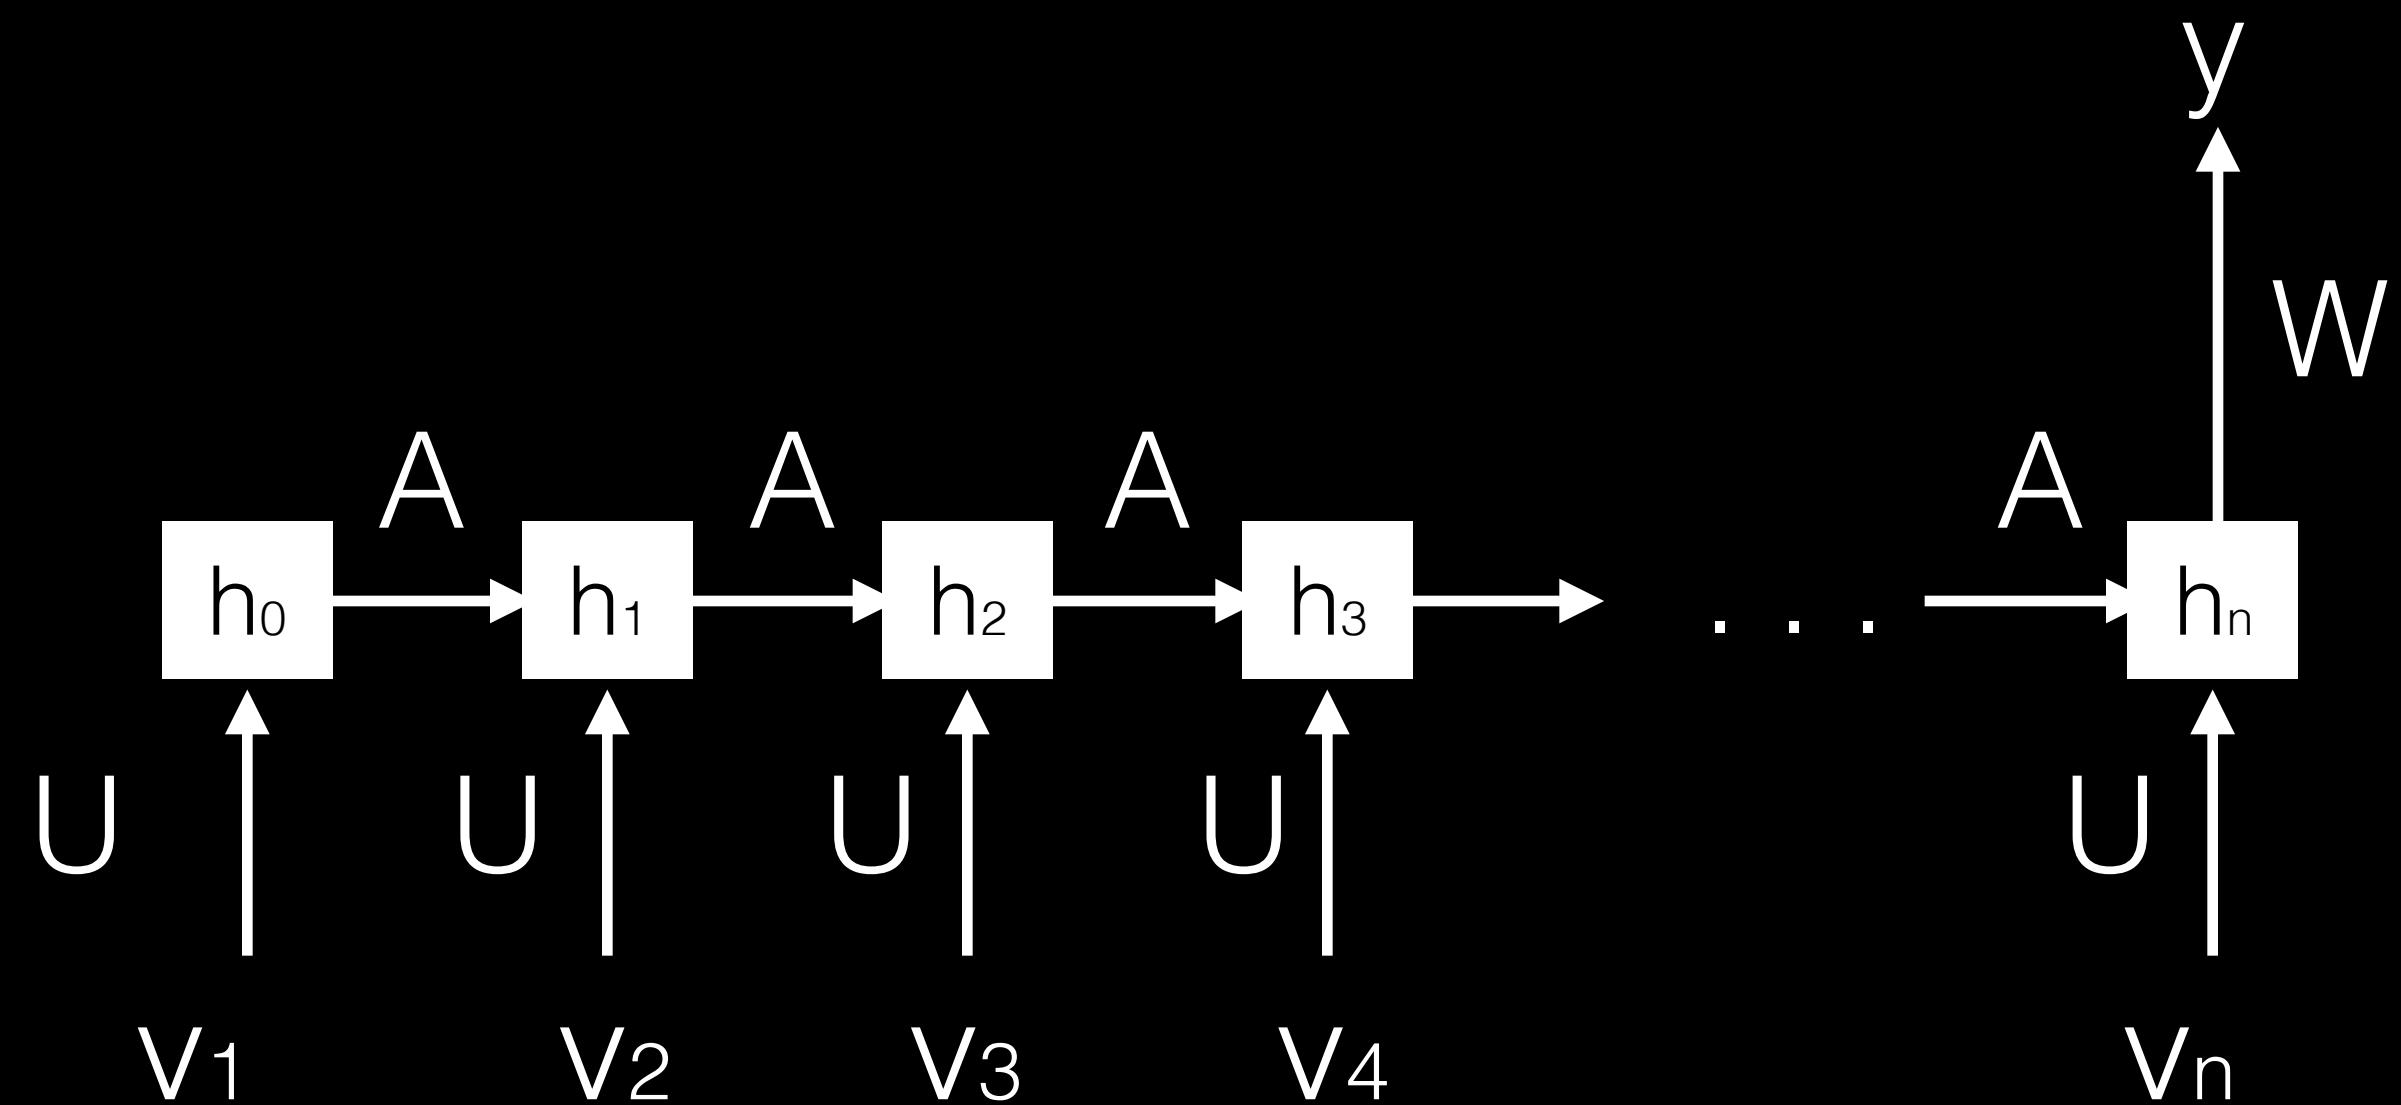
Recurrent Neural Network



Recurrent Neural Network



Recurrent Neural Network



Training RNN with stochastic gradient descent

- For iteration 1, 2, 3, ..., 1000000
 - Sample a random email x and a reply
 - If reply == Yes, update w_1 and w_2

$$w_1 = w_1 + \text{alpha} \frac{d \log(p_1)}{d w_1}$$

$$w_2 = w_2 + \text{alpha} \frac{d \log(p_1)}{d w_2}$$

Training RNN with stochastic gradient descent

- For iteration 1, 2, 3, ..., 1000000
 - Sample a random email x and a reply
 - If reply == Yes, update w_1 and w_2

$$w_1 = w_1 + \text{alpha} \frac{d \log(p_1)}{d w_1}$$

$$w_2 = w_2 + \text{alpha} \frac{d \log(p_1)}{d w_2}$$

Update U , and A

$$A = A + \text{alpha} \frac{d \log(p_1)}{d A}$$

$$U = U + \text{alpha} \frac{d \log(p_1)}{d U}$$

Update all relevant v 's

$$v_i = v_i + \text{alpha} \frac{d \log(p_1)}{d v_i}$$

Training RNN with stochastic gradient descent

- For iteration 1, 2, 3, ..., 1000000
 - Sample a random email x and a reply
 - If reply == Yes, update w_1 and w_2

$$w_1 = w_1 + \text{alpha} \frac{d \log(p_1)}{d w_1}$$

Update U , and A

$$A = A + \text{alpha} \frac{d \log(p_1)}{d A}$$

Update all relevant v 's

$$w_2 = w_2 + \text{alpha} \frac{d \log(p_1)}{d w_2}$$

$$U = U + \text{alpha} \frac{d \log(p_1)}{d U}$$

$$v_i = v_i + \text{alpha} \frac{d \log(p_1)}{d v_i}$$

**Very hard
to derive!**

**Use
autodiff :)**

The big picture so far

- Bag-of-word representation
- RNN representation for variable-sized input
- Autodiff to compute the partial derivatives (TensorFlow, Theano, Torch)
- Stochastic gradient descent for training

More friendly “AutoReply”

- Are you visiting Vietnam for the new year , Quoc ? -> Yes , see you soon !
- Are you hanging out with us tonight ? -> No , I am too busy .
- Did you read the cool paper on ResNet? -> Yes , it's nice !
- ...

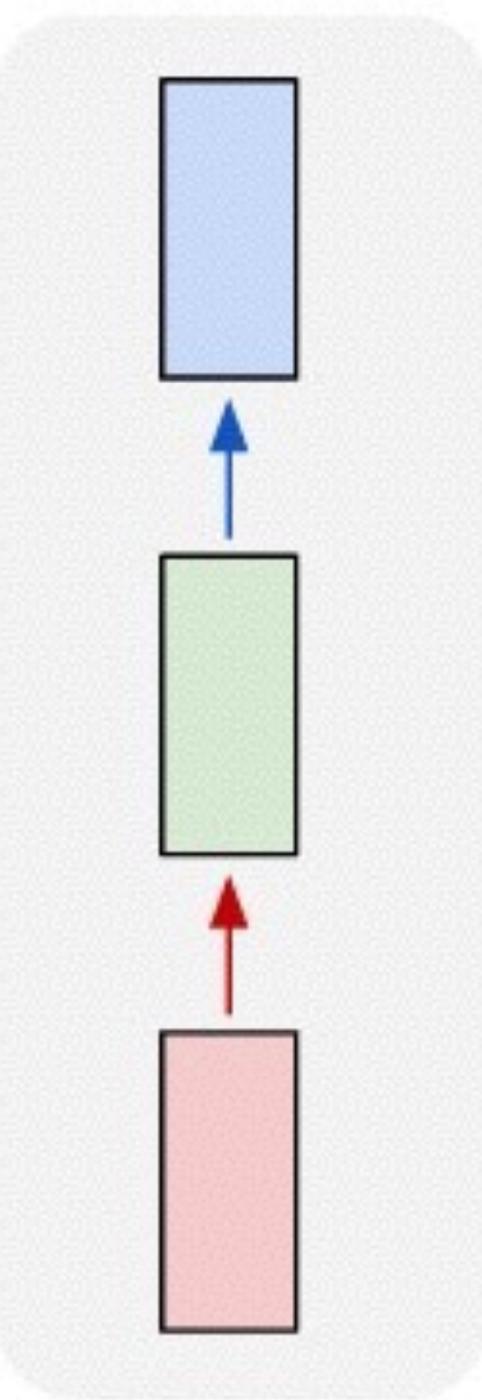
Better Formulation

- Mapping between variable-length input to variable length output

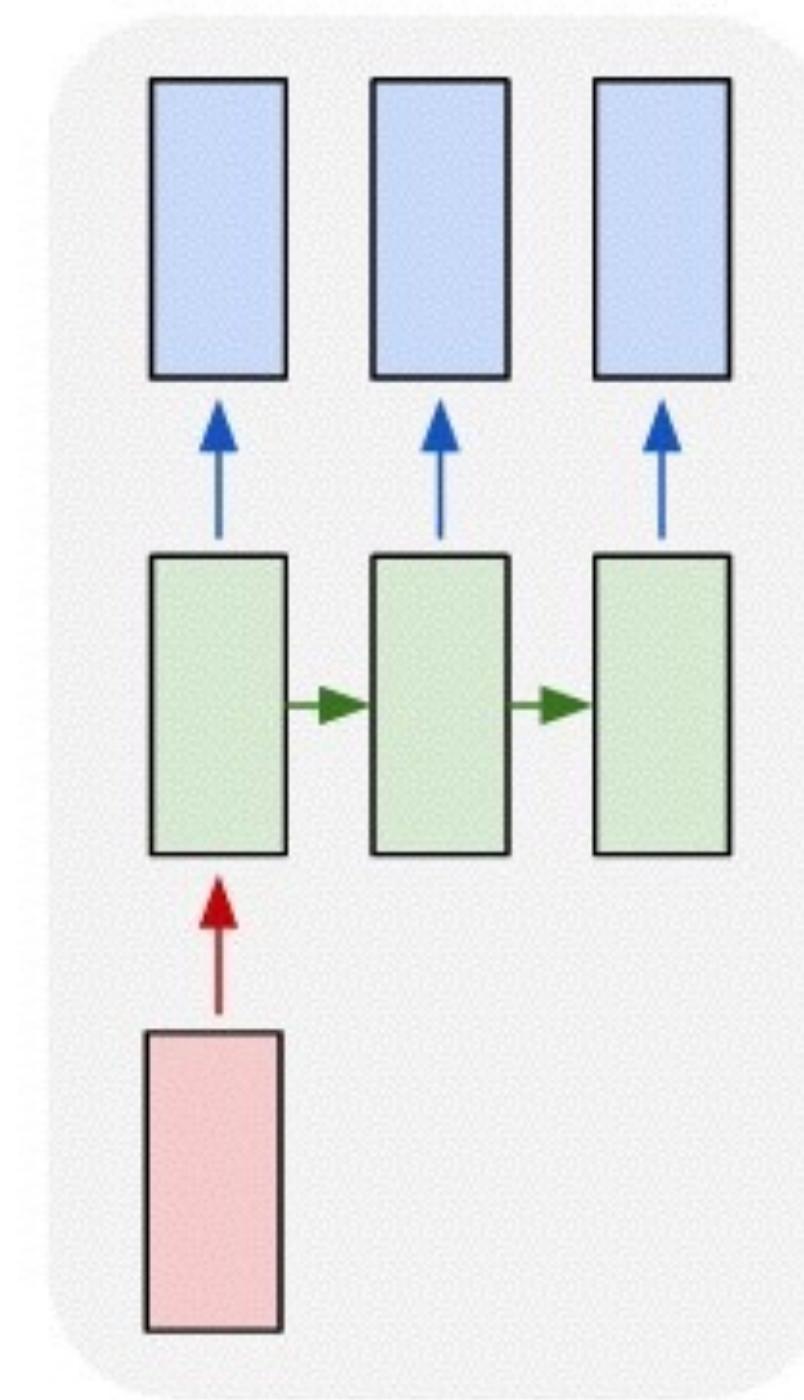
Better Formulation

- Mapping between variable-length input to variable length output
- Applications: AutoReply, Translation, Image Captioning, Summarization, Speech Transcription, Conversation, Q&A, ...

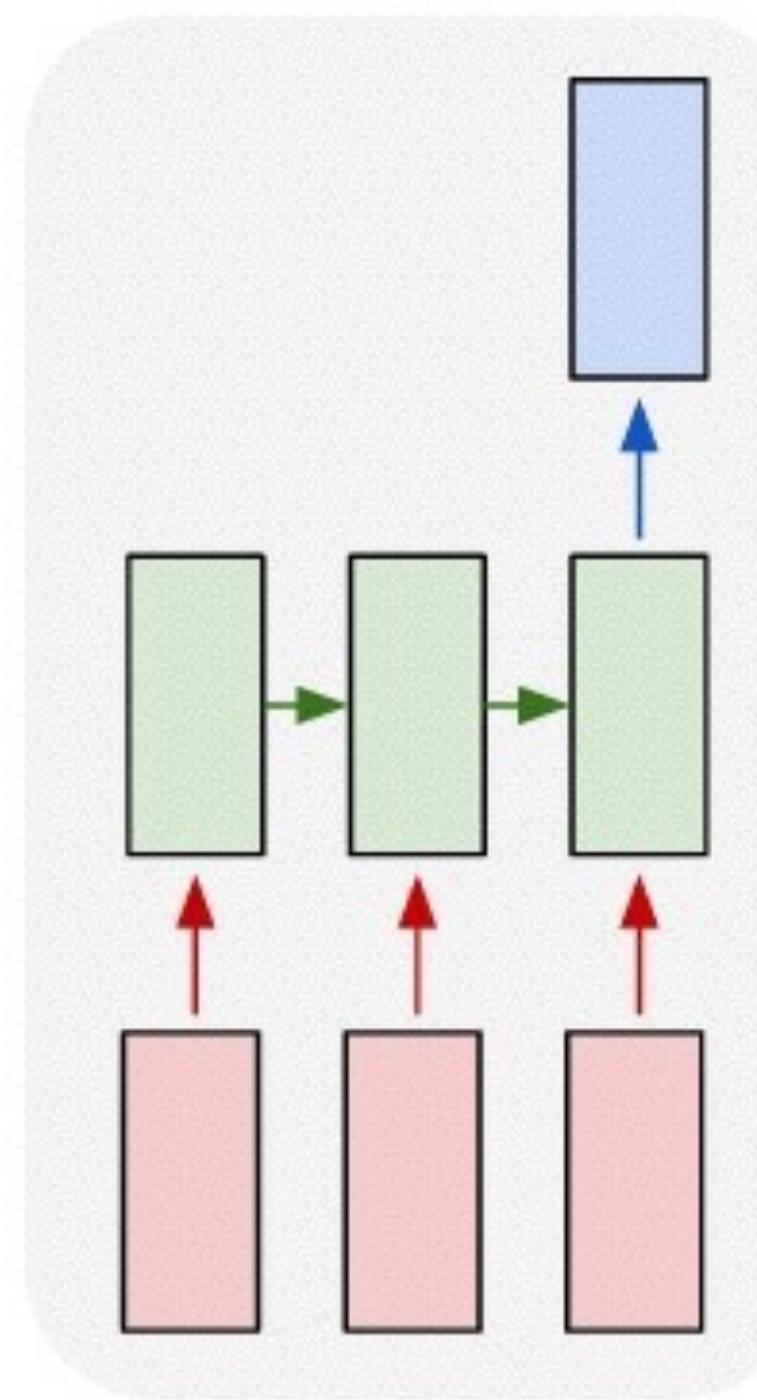
one to one



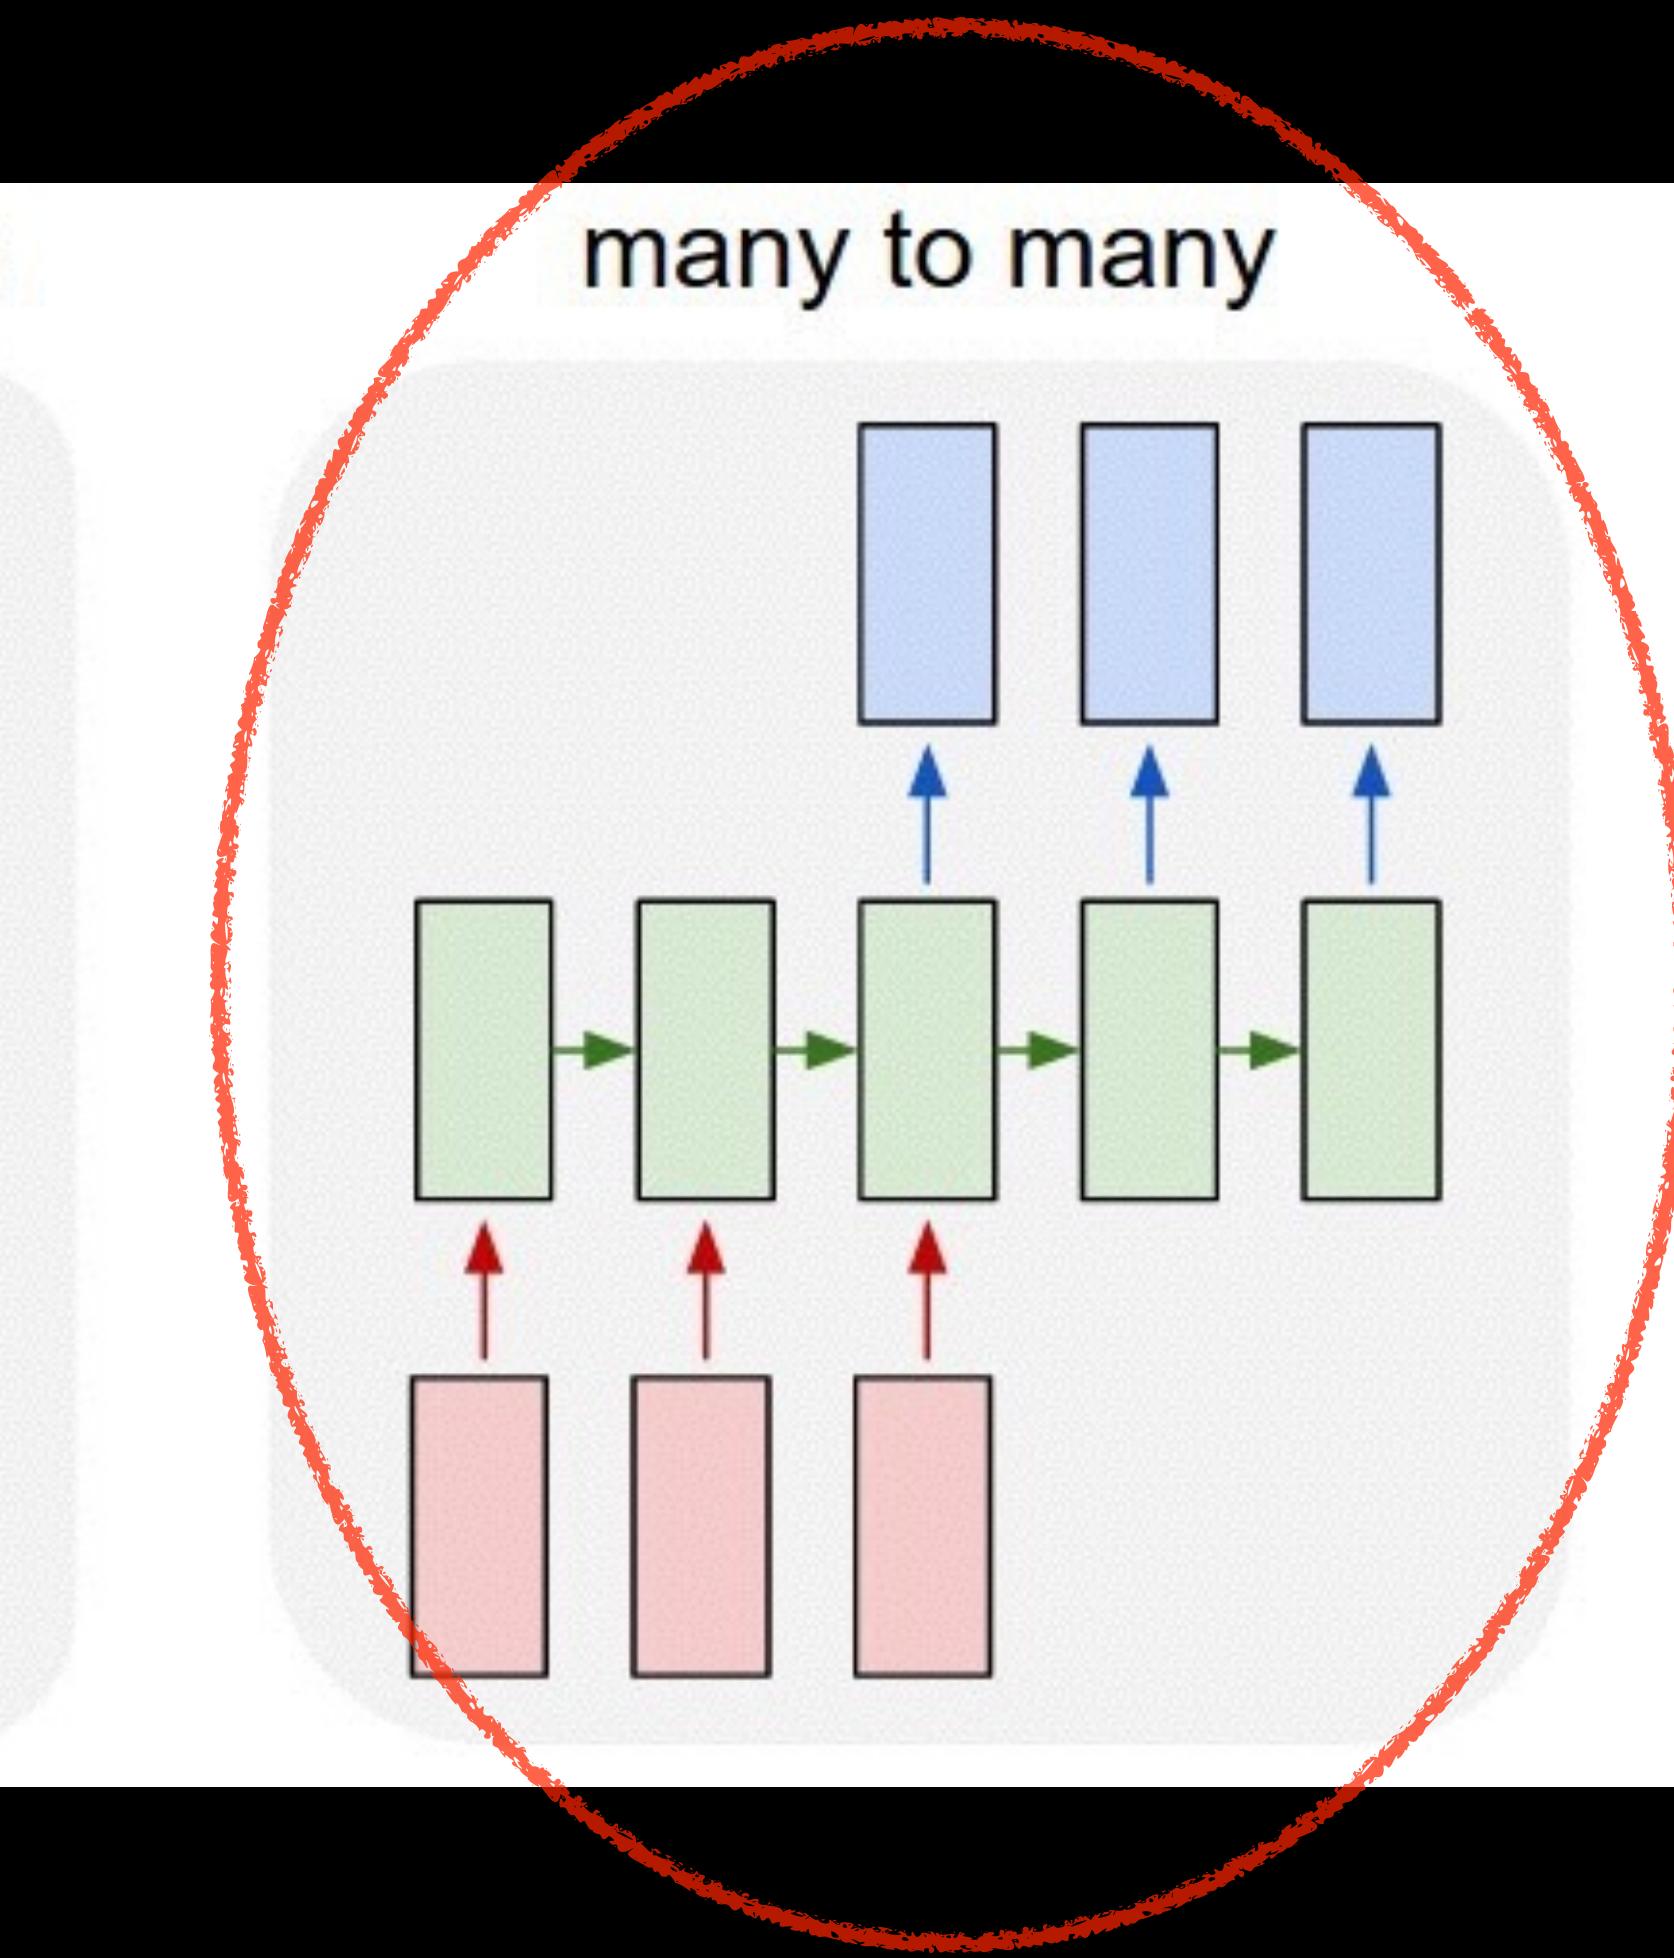
one to many



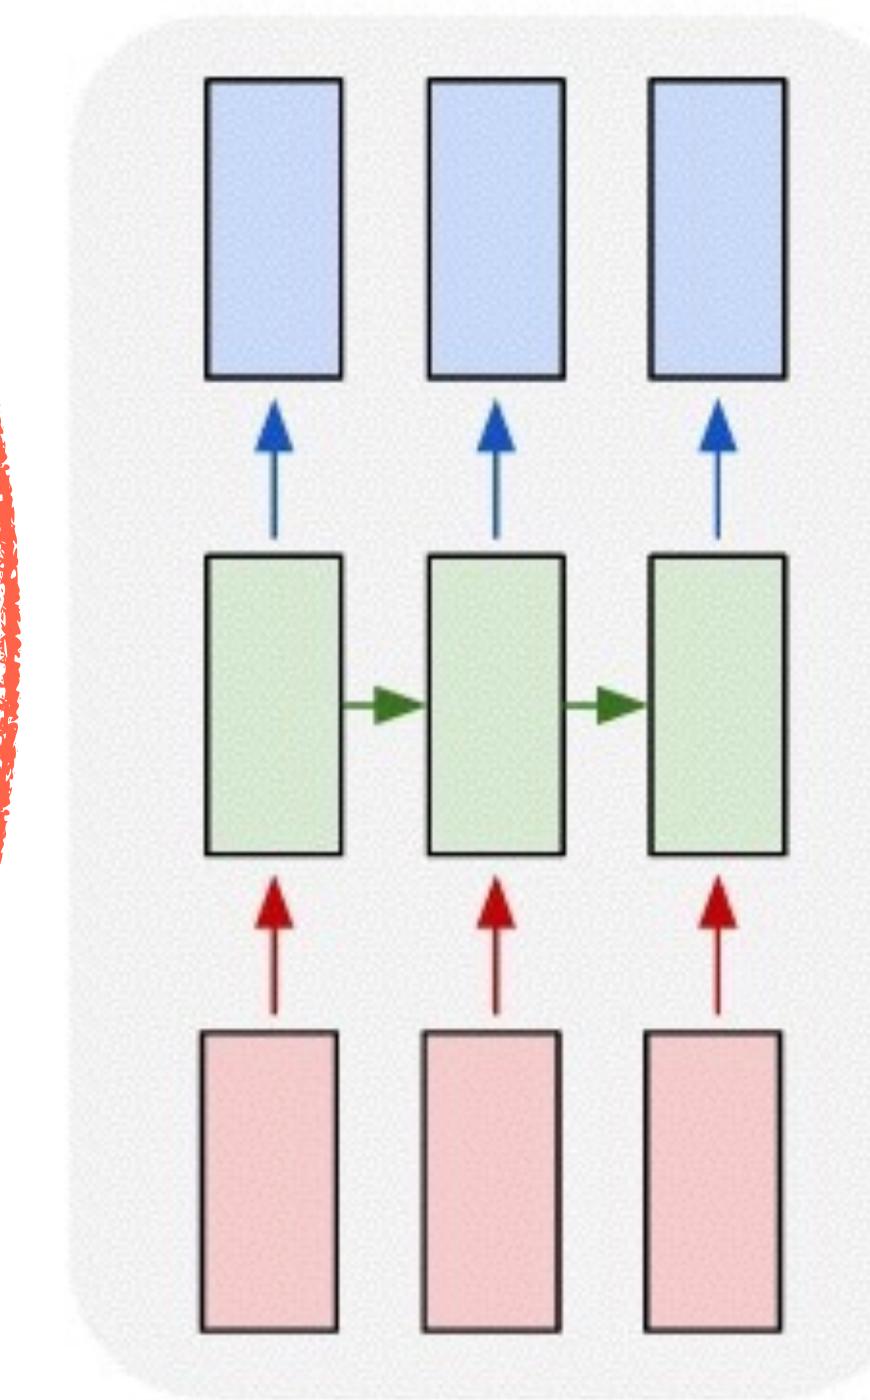
many to one



many to many

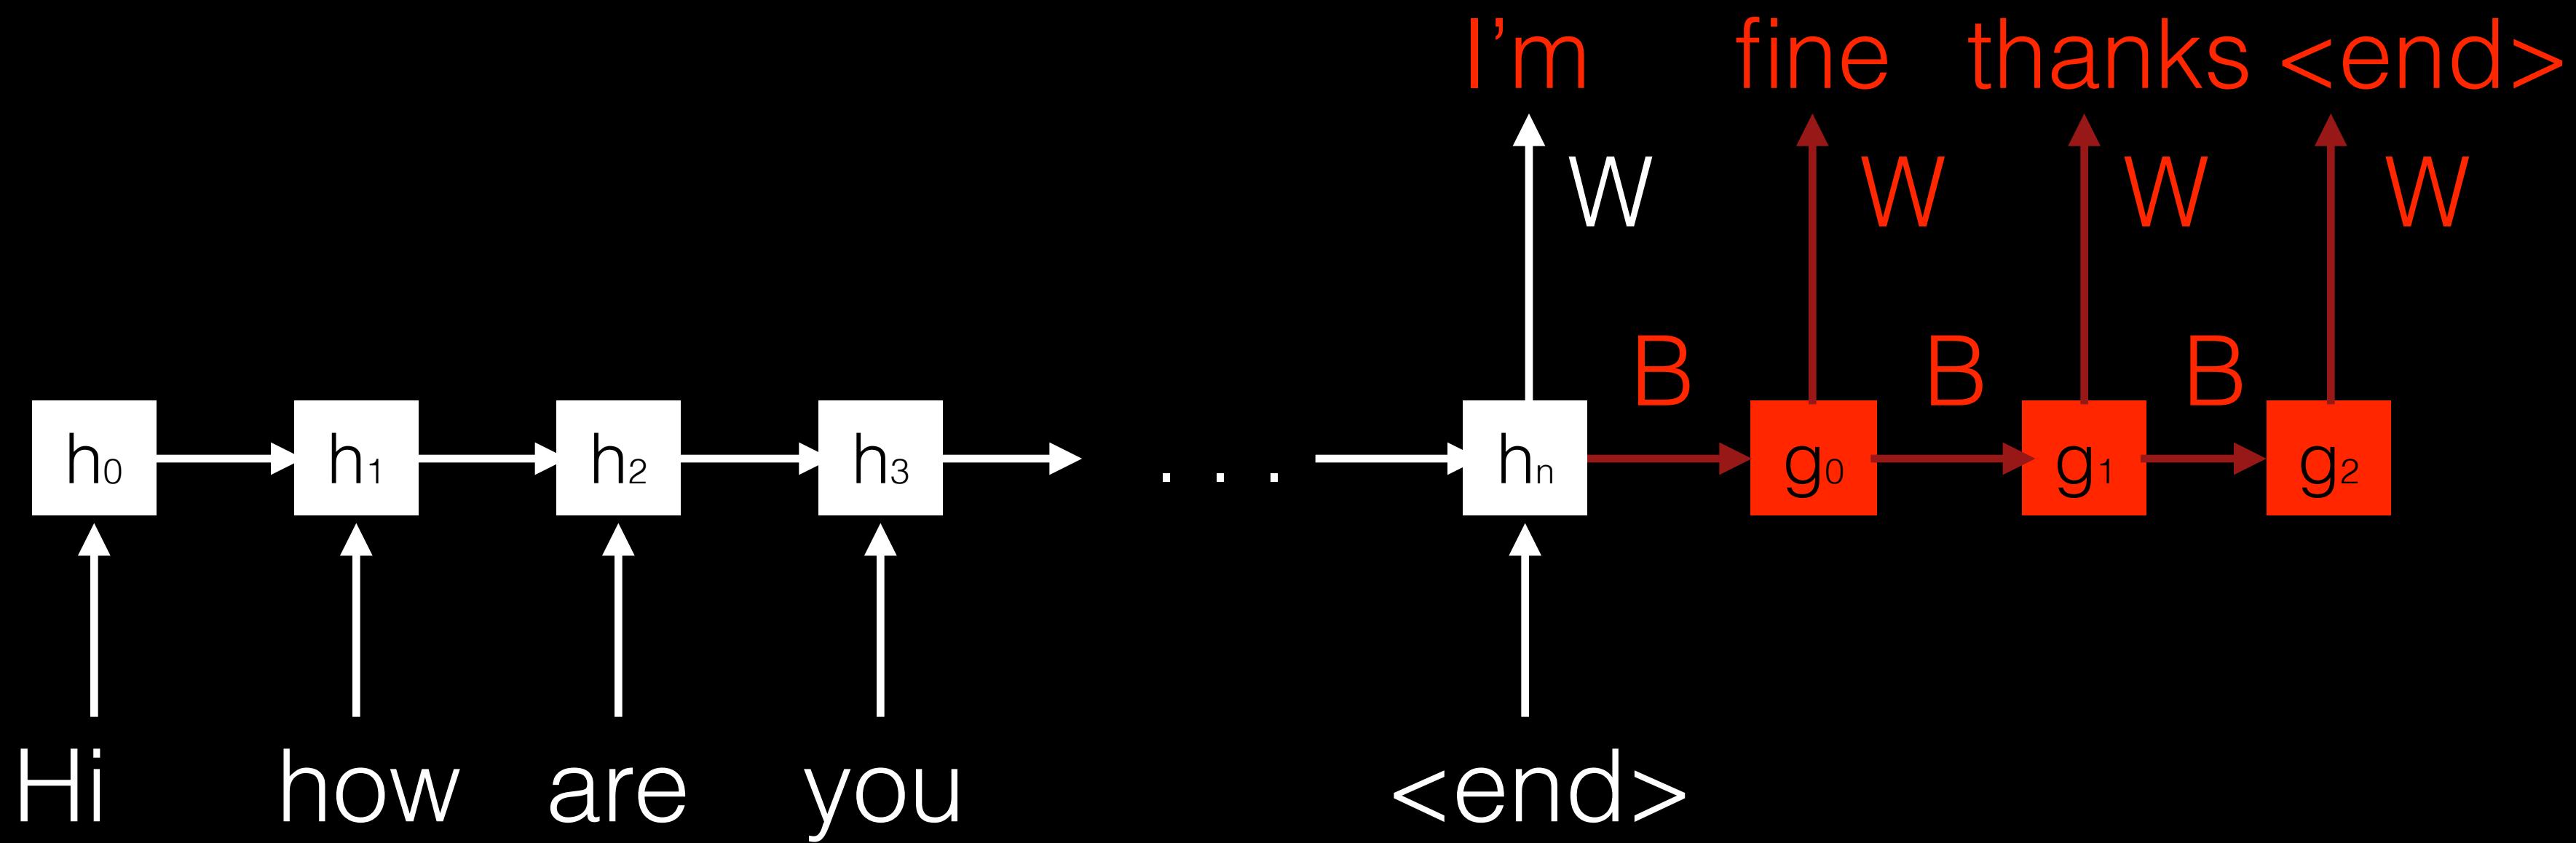


many to many

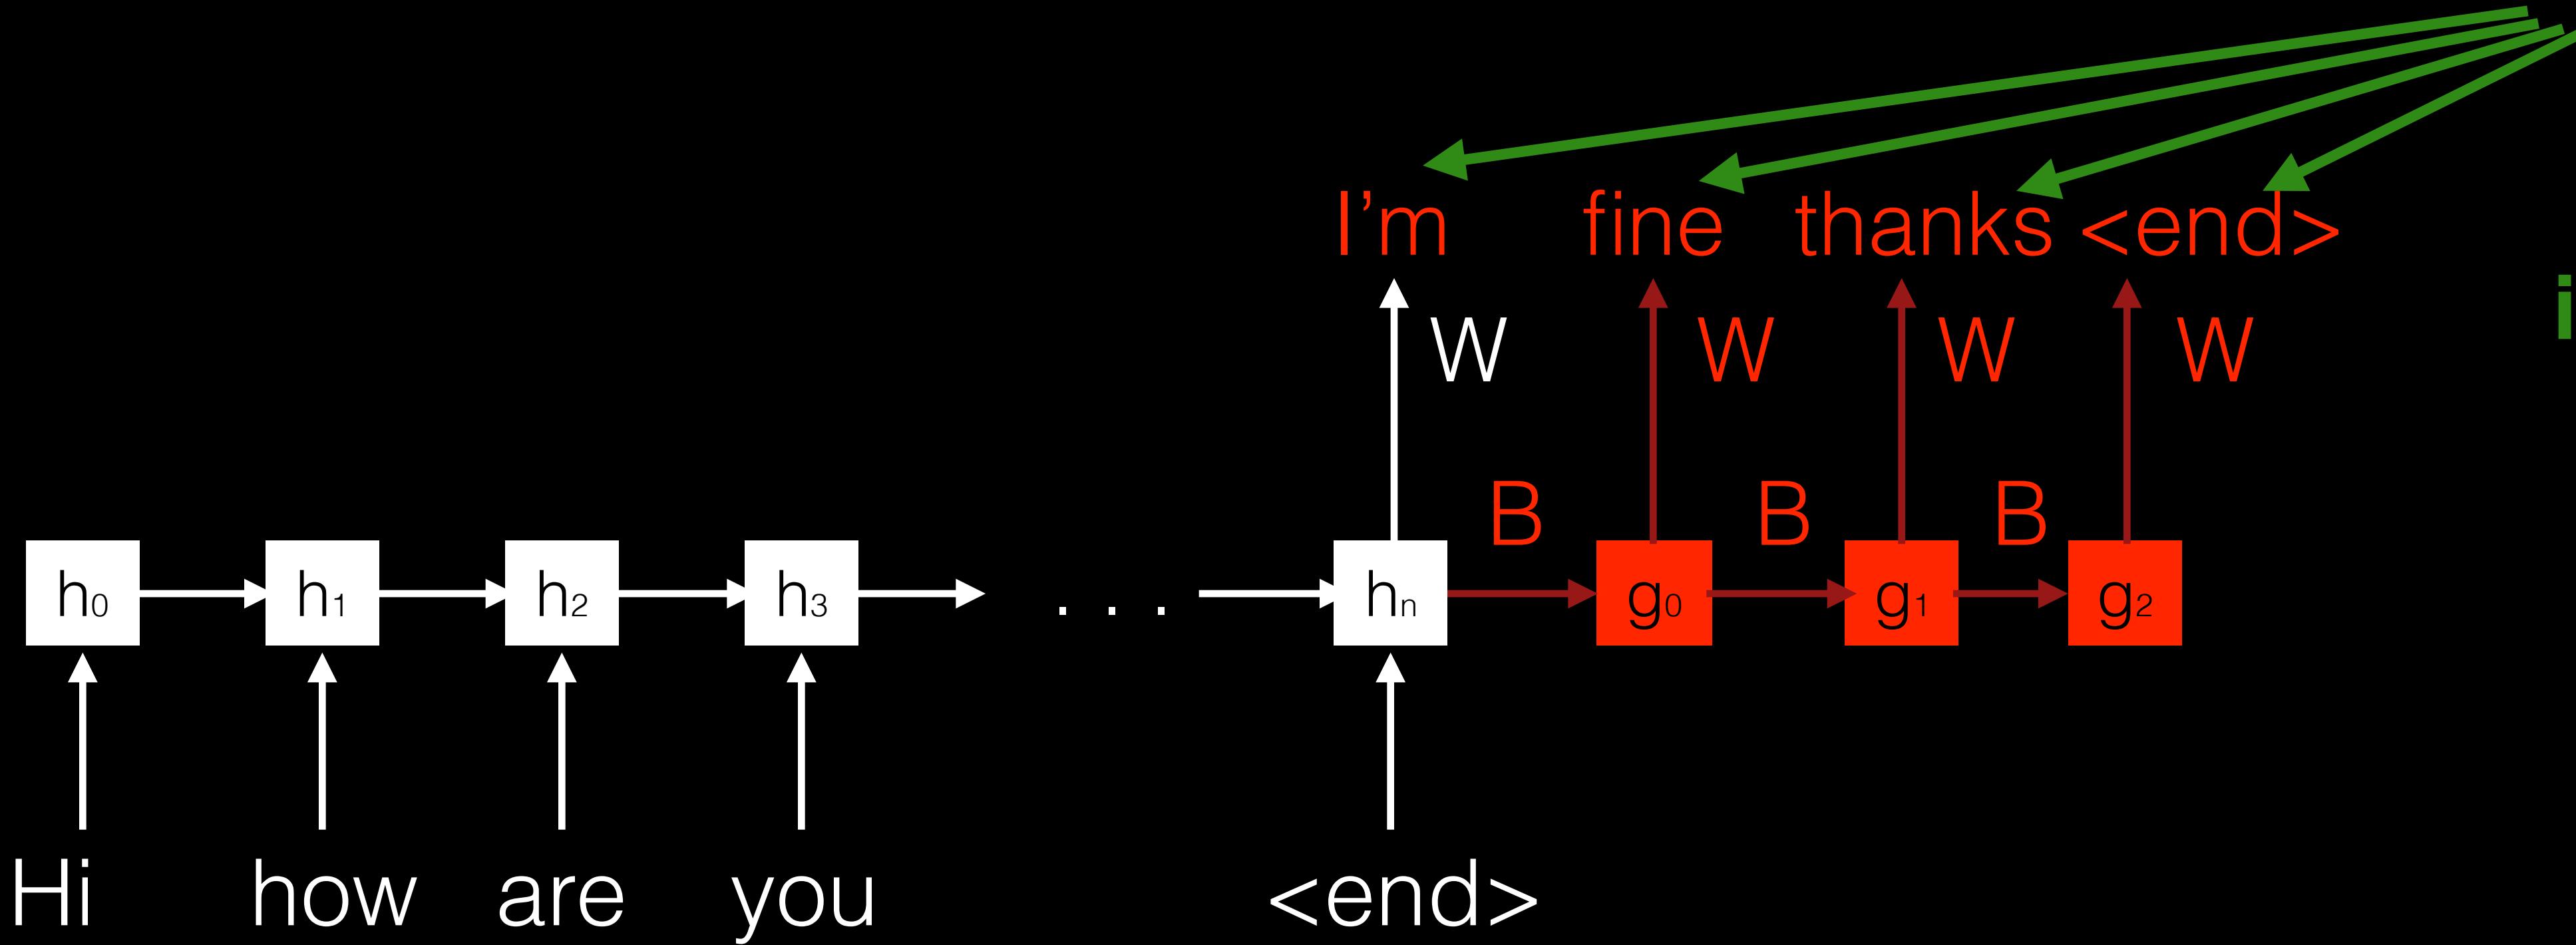


Andrej Karpathy. The Unreasonable Effectiveness of Recurrent Neural Networks

Better Formulation

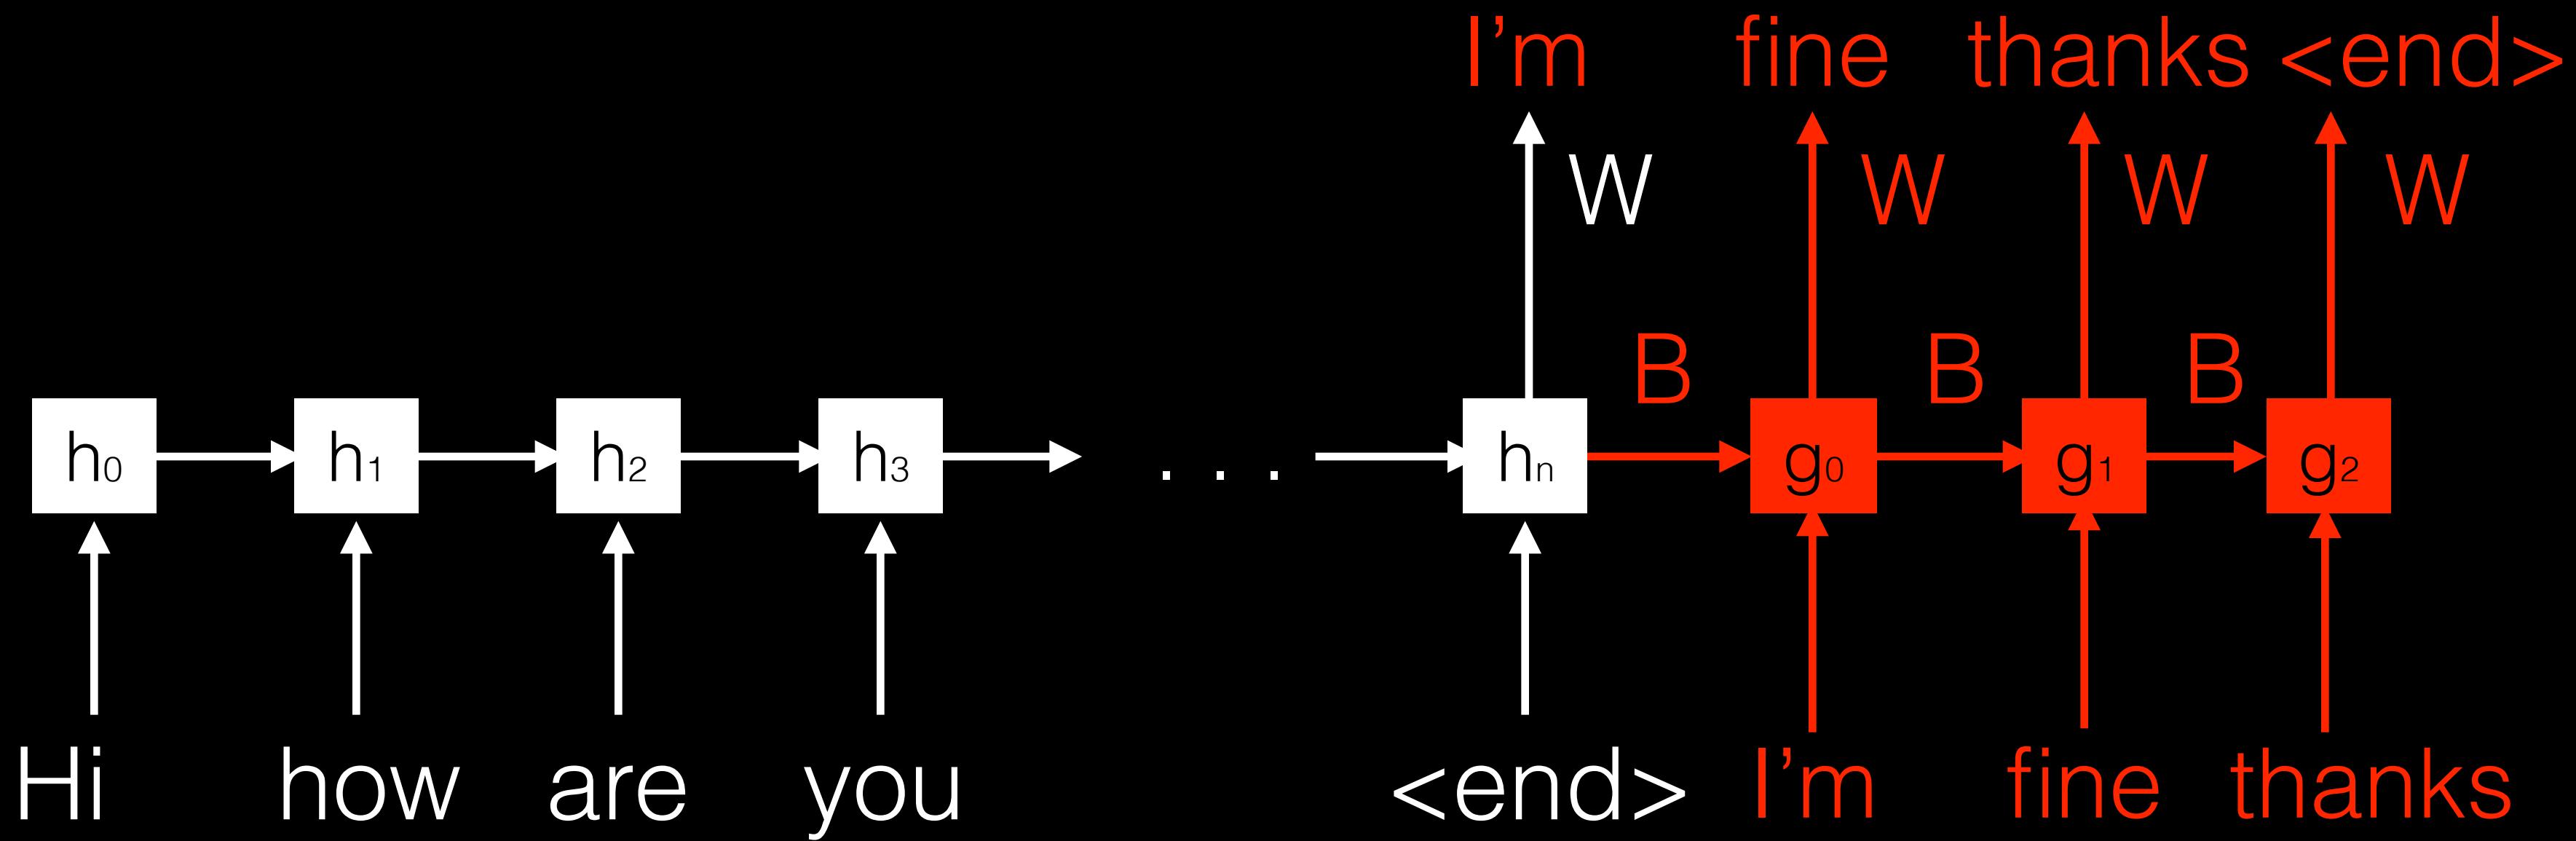


Better Formulation

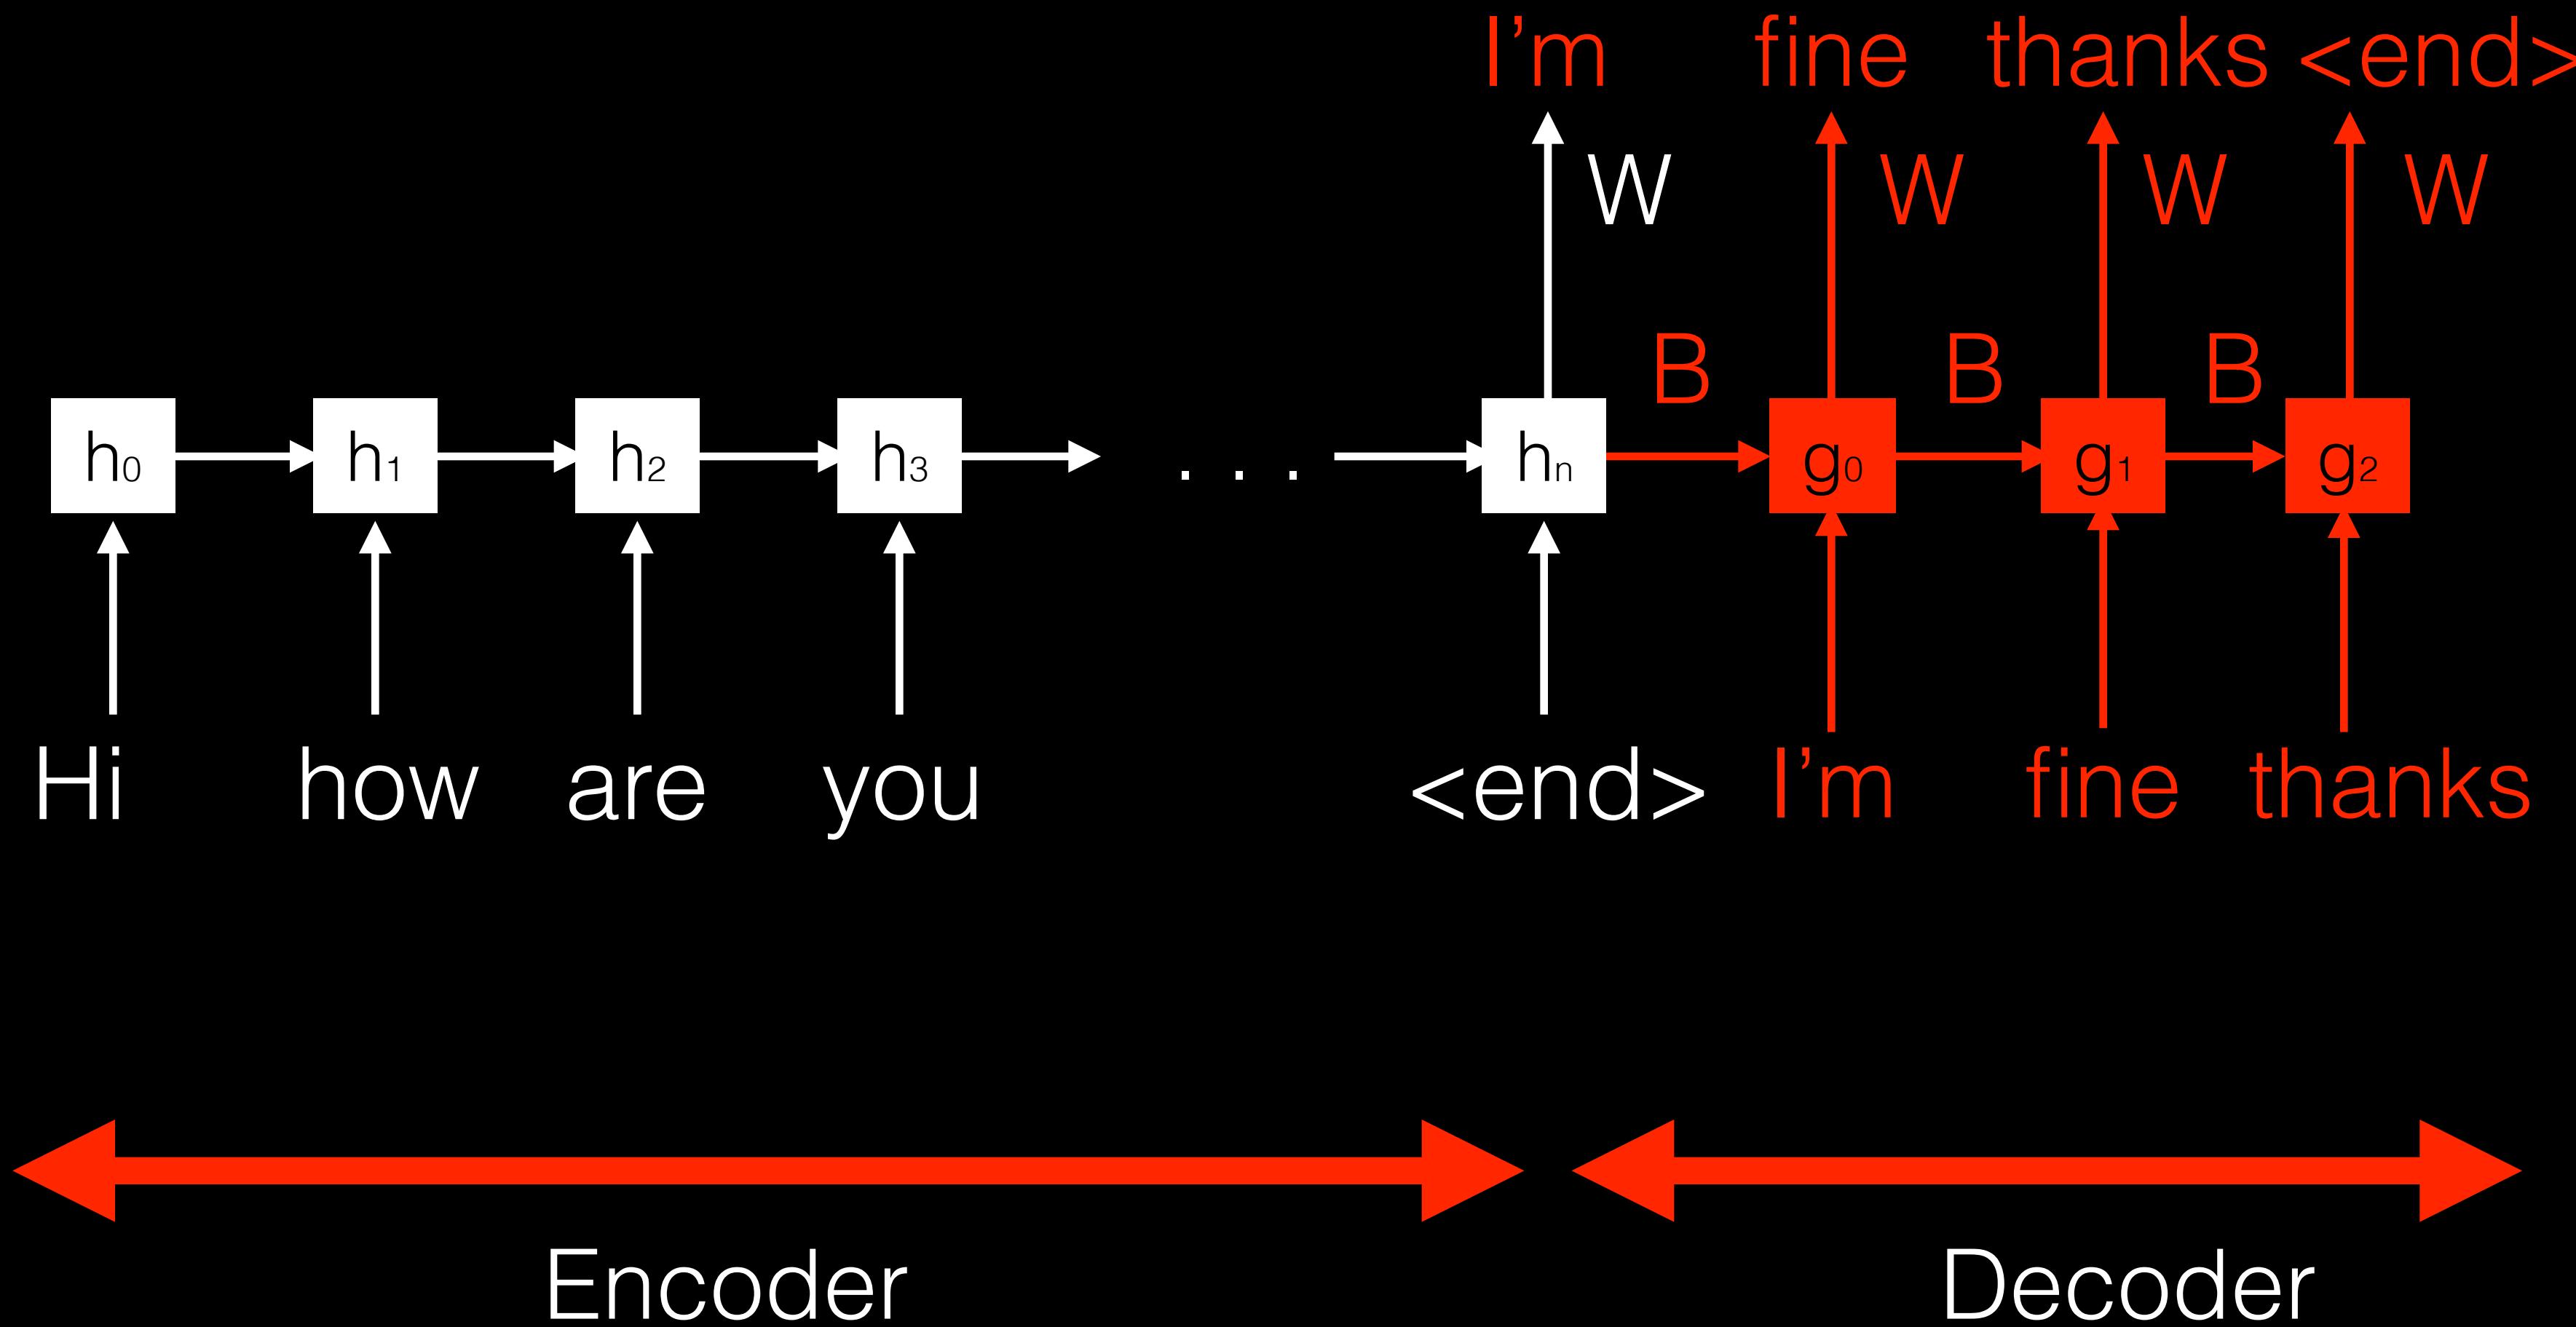


Number of choices = number of words in vocabulary

Better Formulation



Better Formulation



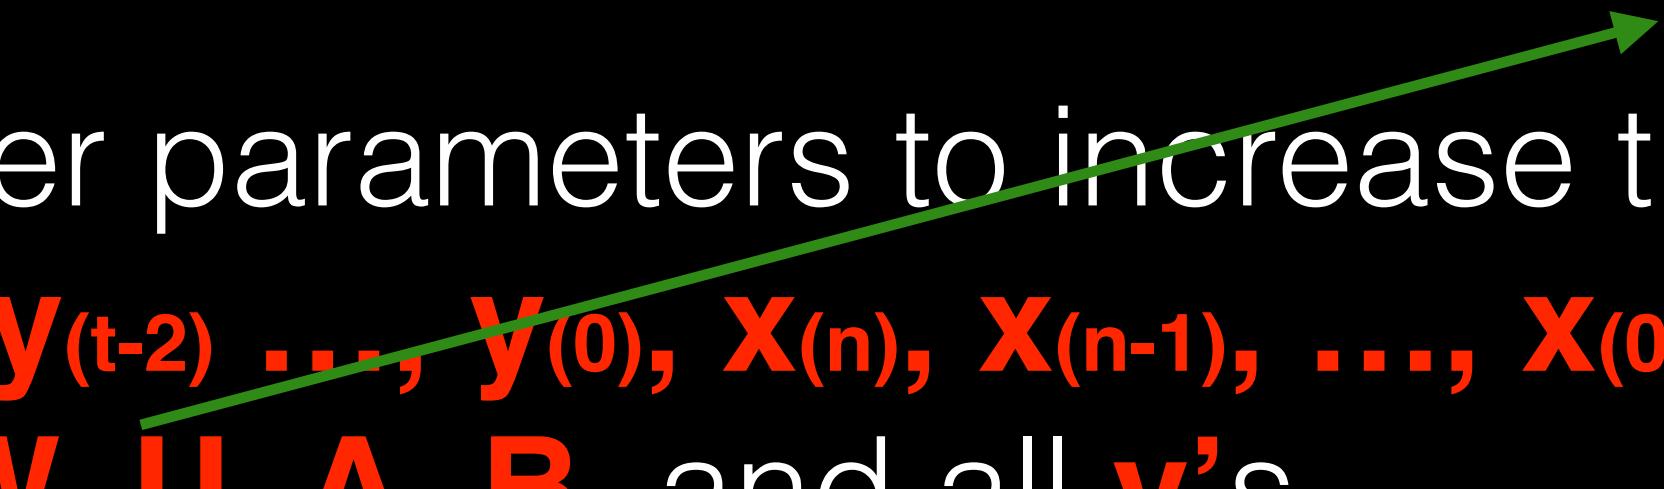
Sequence to Sequence Training with SGD

- For iteration 1, 2, 3, ..., 1000000
 - Sample an email \mathbf{x} and a reply \mathbf{y}
 - Sample a random word $\mathbf{y}_{(t)}$ in \mathbf{y}
 - Update RNN encoder and decoder parameters to increase the probability of word $\mathbf{y}_{(t)}$ given $\mathbf{y}_{(t-1)}, \mathbf{y}_{(t-2)} \dots, \mathbf{y}_{(0)}, \mathbf{x}_{(n)}, \mathbf{x}_{(n-1)}, \dots, \mathbf{x}_{(0)}$ using partial derivative with respect to $\mathbf{W}, \mathbf{U}, \mathbf{A}, \mathbf{B}$, and all \mathbf{v} 's

Sequence to Sequence Training with SGD

- For iteration 1, 2, 3, ..., 1000000
 - Sample an email \mathbf{x} and a reply \mathbf{y}
 - Sample a random word $\mathbf{y}_{(t)}$ in \mathbf{y}
 - Update RNN encoder and decoder parameters to increase the probability of word $\mathbf{y}_{(t)}$ given $\mathbf{y}_{(t-1)}, \mathbf{y}_{(t-2)}, \dots, \mathbf{y}_{(0)}, \mathbf{x}_{(n)}, \mathbf{x}_{(n-1)}, \dots, \mathbf{x}_{(0)}$ using partial derivative with respect to $\mathbf{W}, \mathbf{U}, \mathbf{A}, \mathbf{B}$, and all \mathbf{v} 's

Very hard
to derive!
Use
autodiff :)



Sequence to Sequence Prediction

- For any incoming email \mathbf{x}
 - Given \mathbf{x} , find word $\mathbf{y}_{(0)}$ with highest probability using RNN
 - Given $\mathbf{y}_{(0)}$ and \mathbf{x} , find word $\mathbf{y}_{(1)}$ with highest probability using RNN
 - ...
 - Stop when see <end>

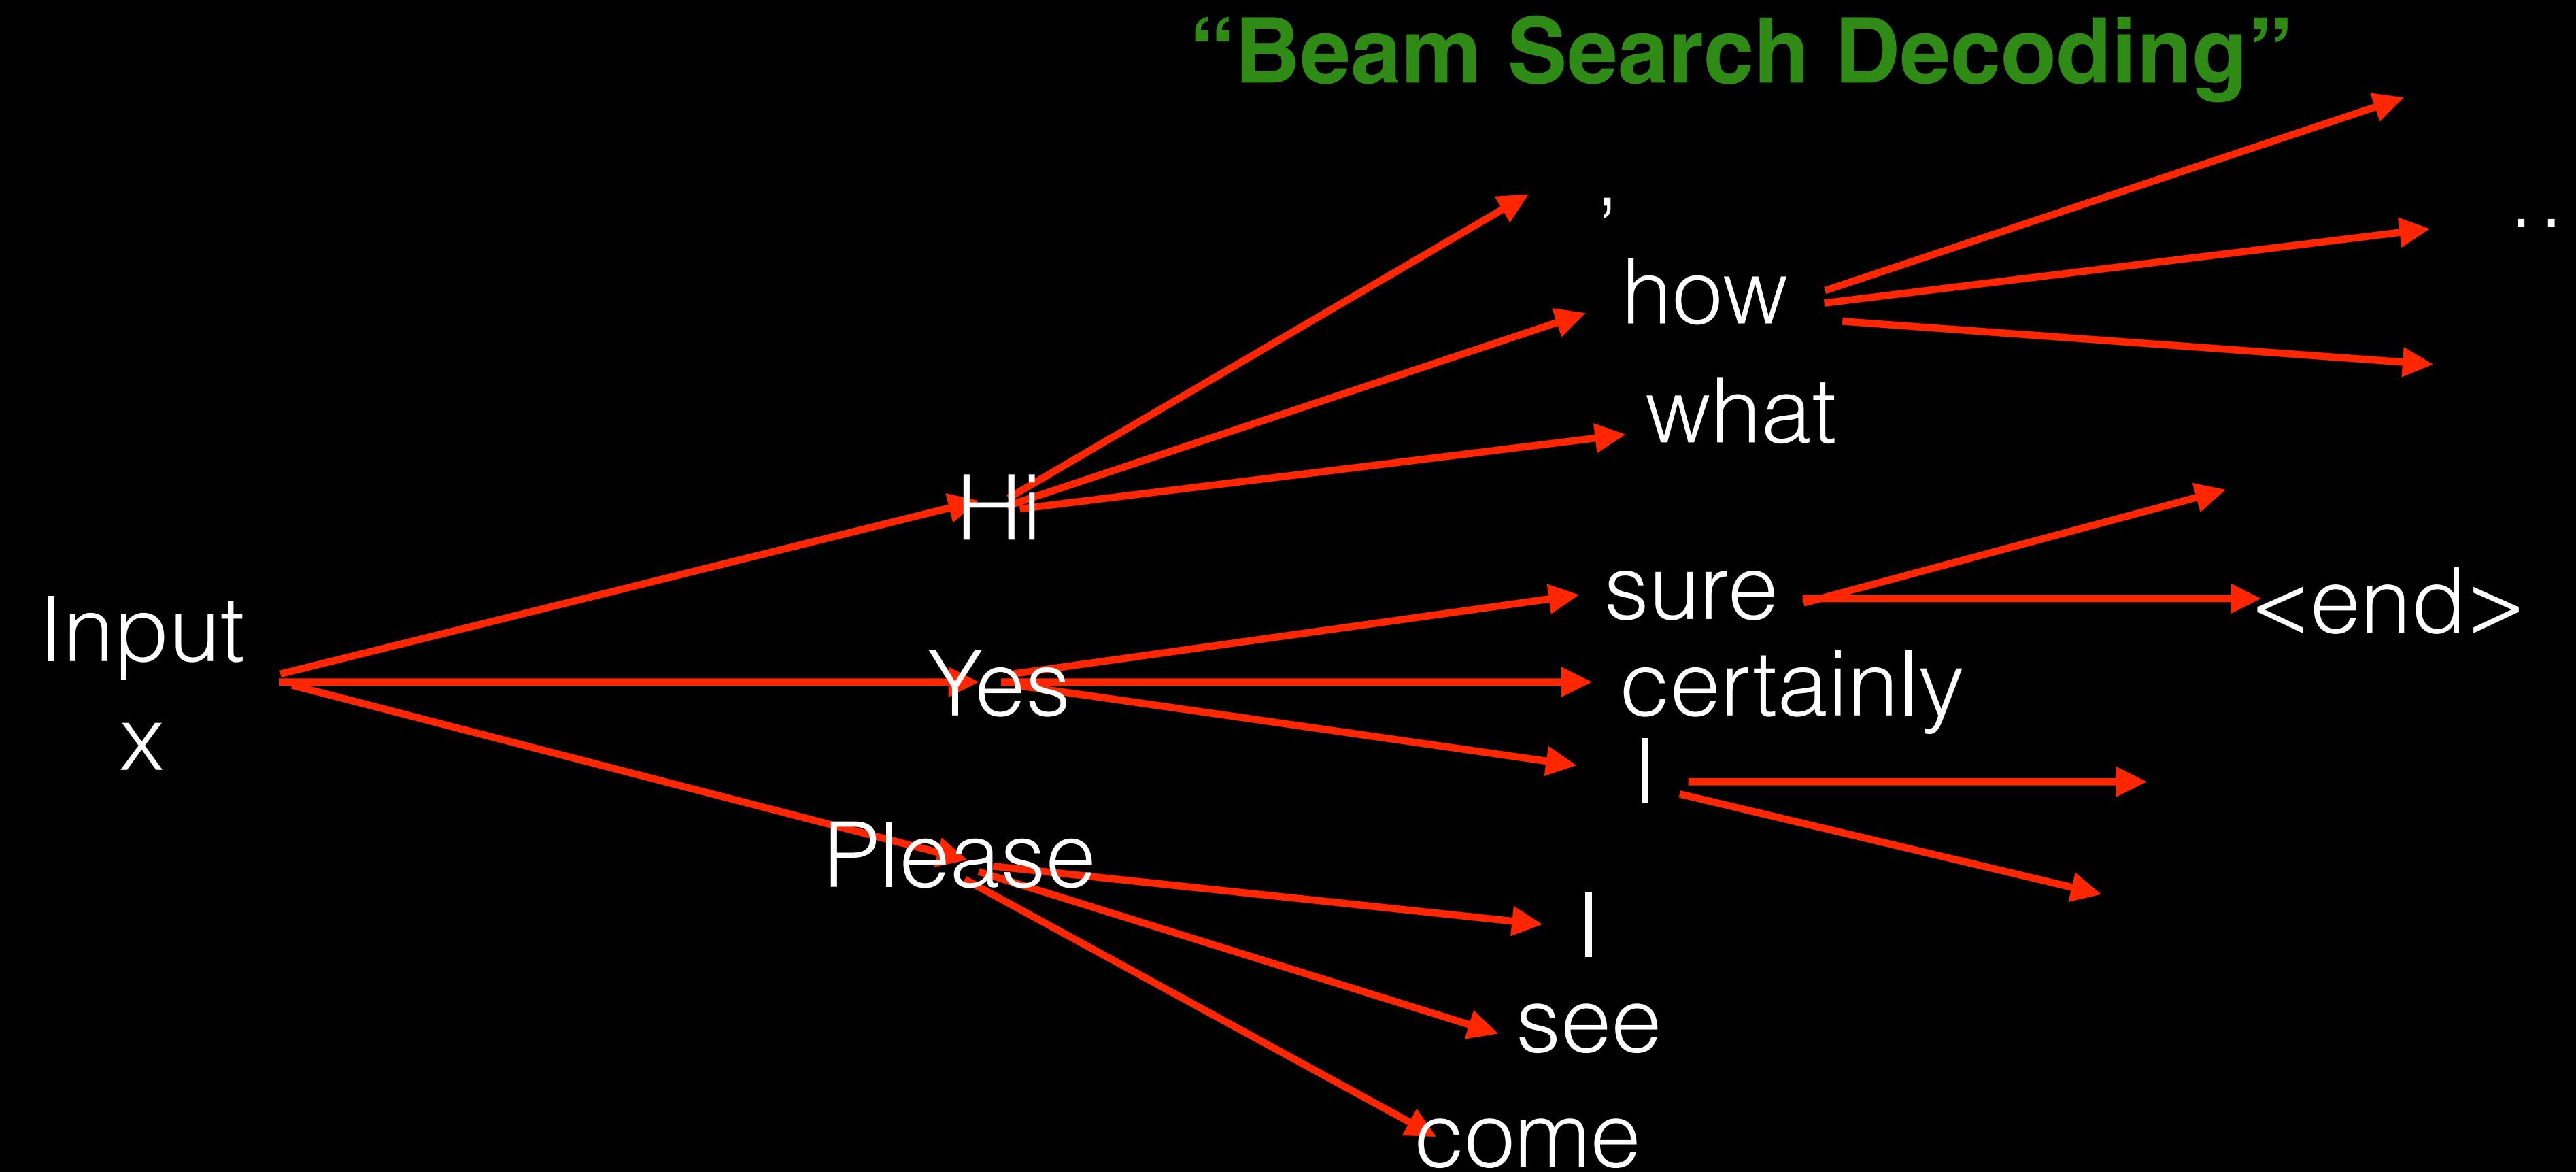
“Greedy Decoding”

Sequence to Sequence Prediction

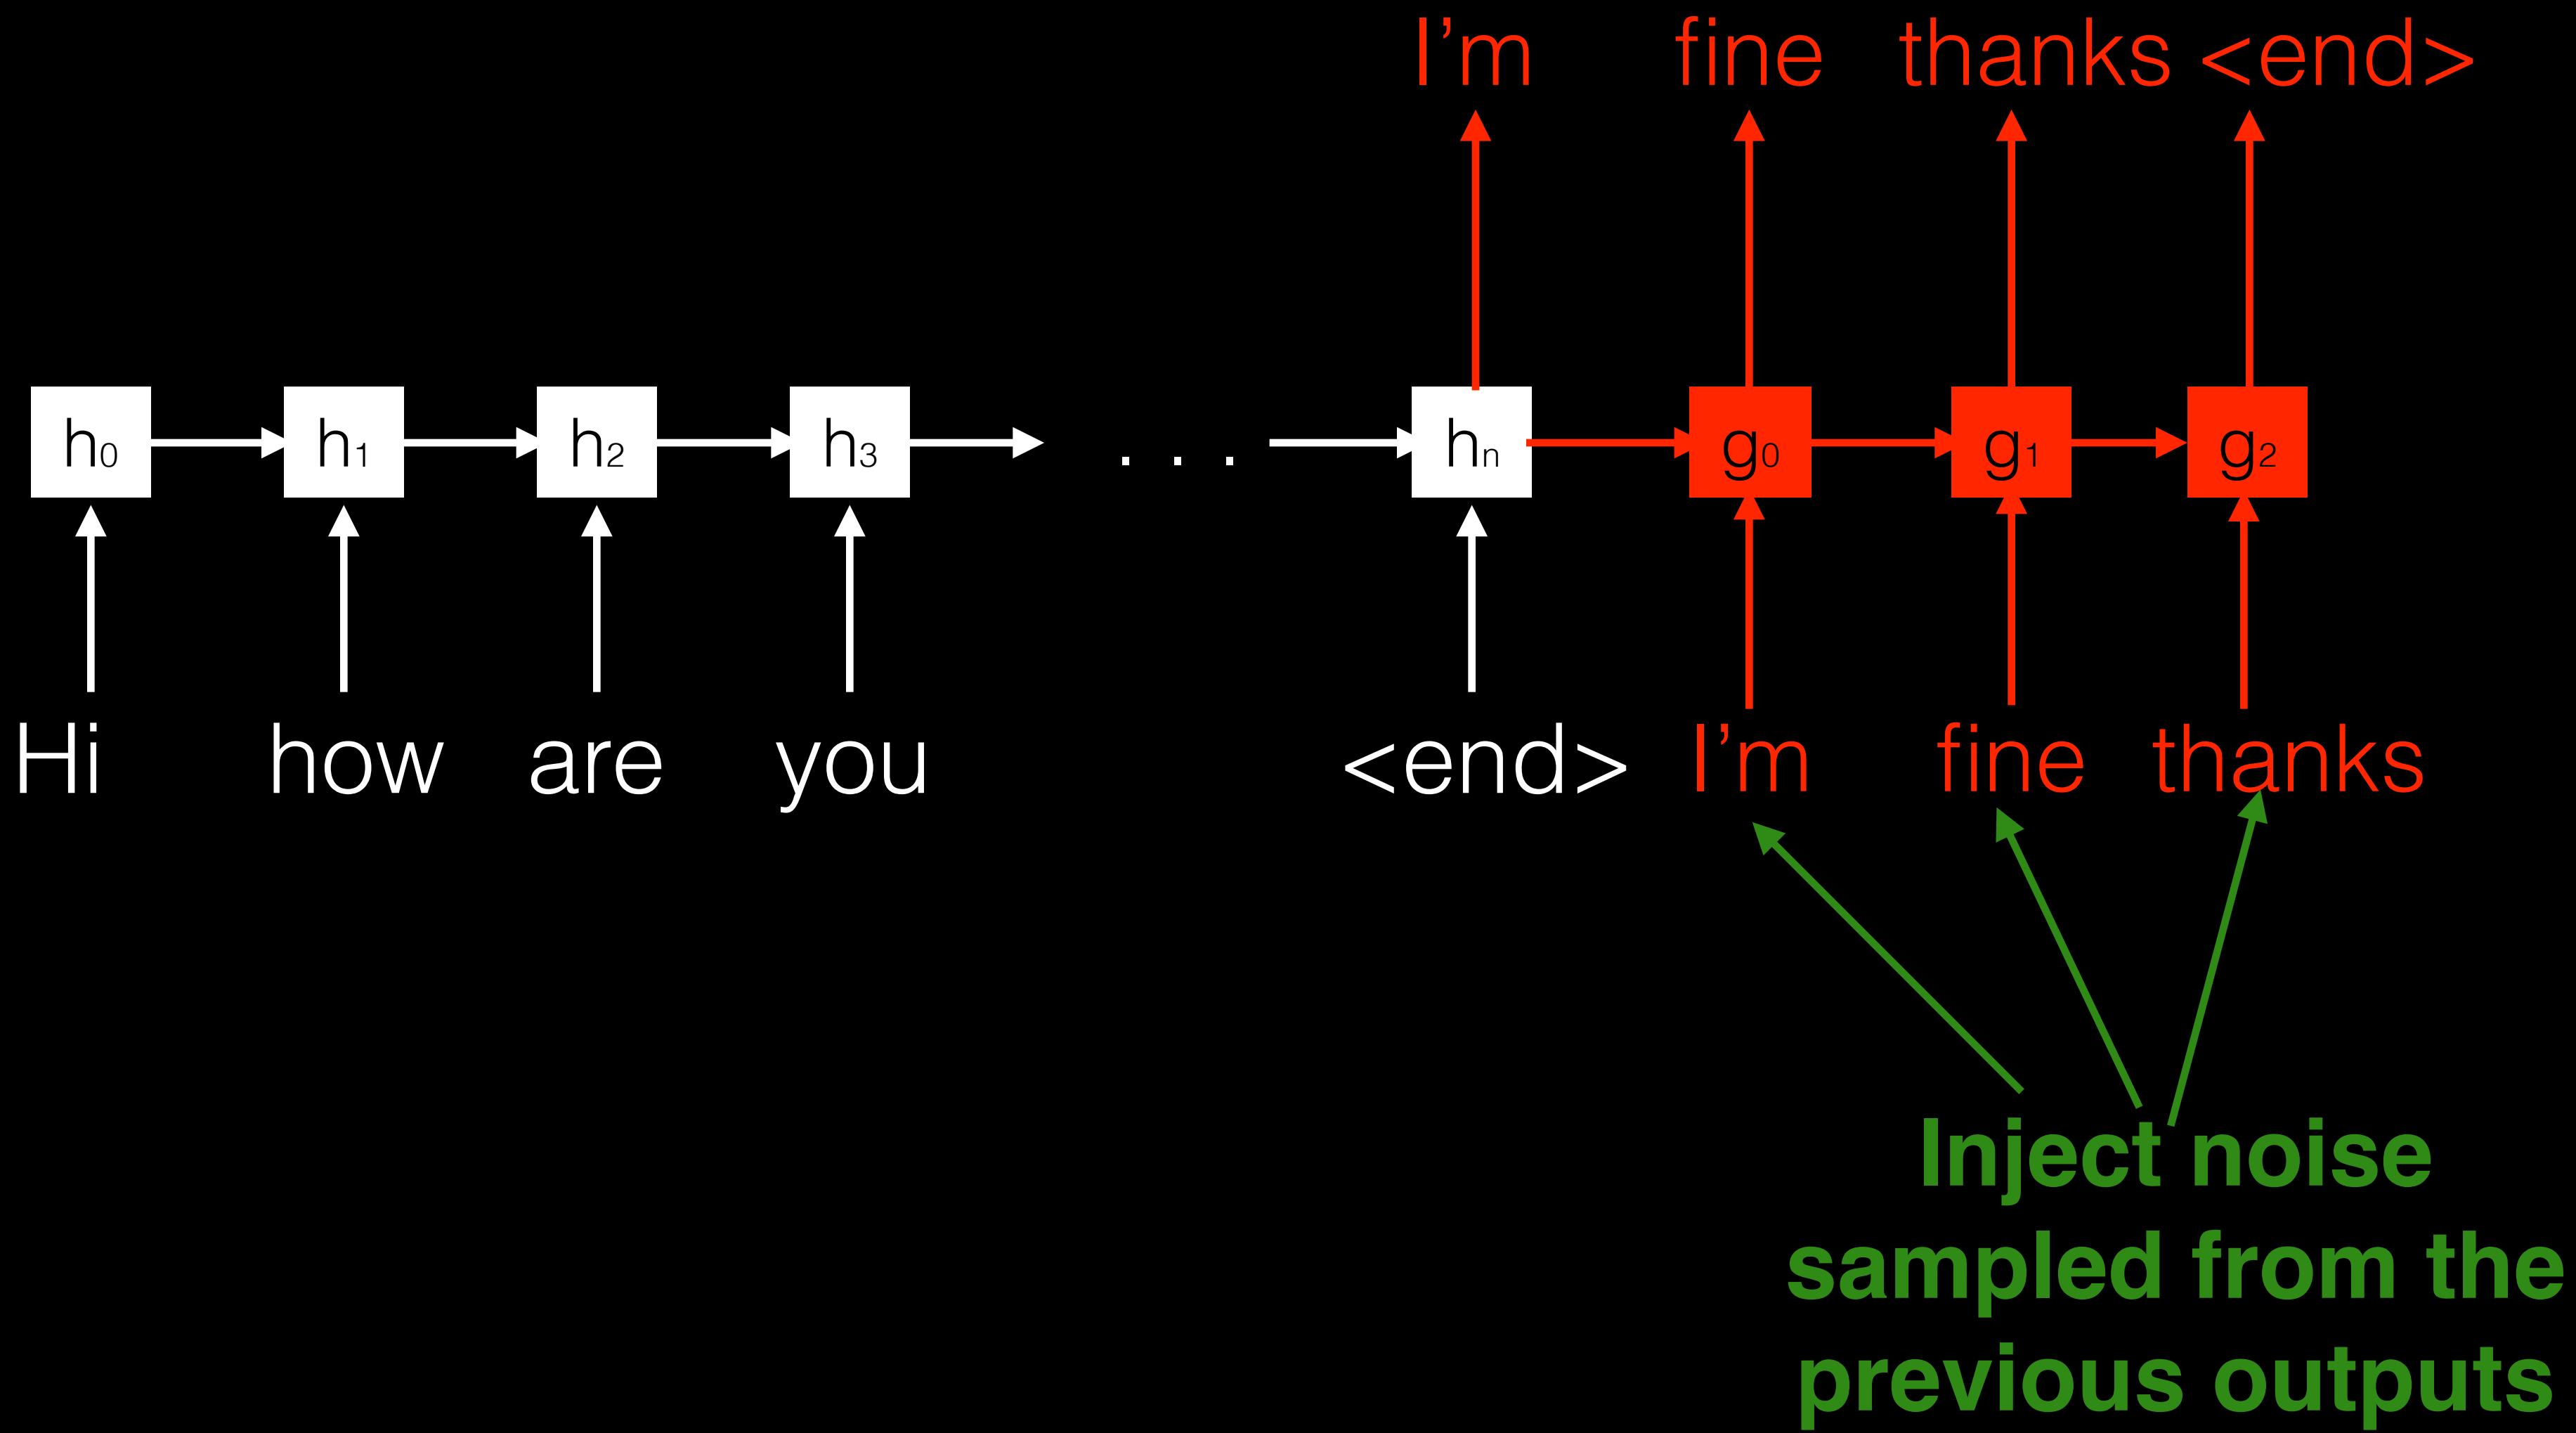
- For any incoming email \mathbf{x}
 - Given \mathbf{x} , find \mathbf{k} candidates for $\mathbf{y}_{(0)}$ with highest probability using RNN
 - Given \mathbf{x} , for each candidate $\mathbf{y}_{(0)}$, find \mathbf{k} candidates for word $\mathbf{y}_{(1)}$ with highest probability using RNN
 - ...
 - Stop when see <end> on each beam
 - Reply = beam with highest probability

“Beam Search Decoding”

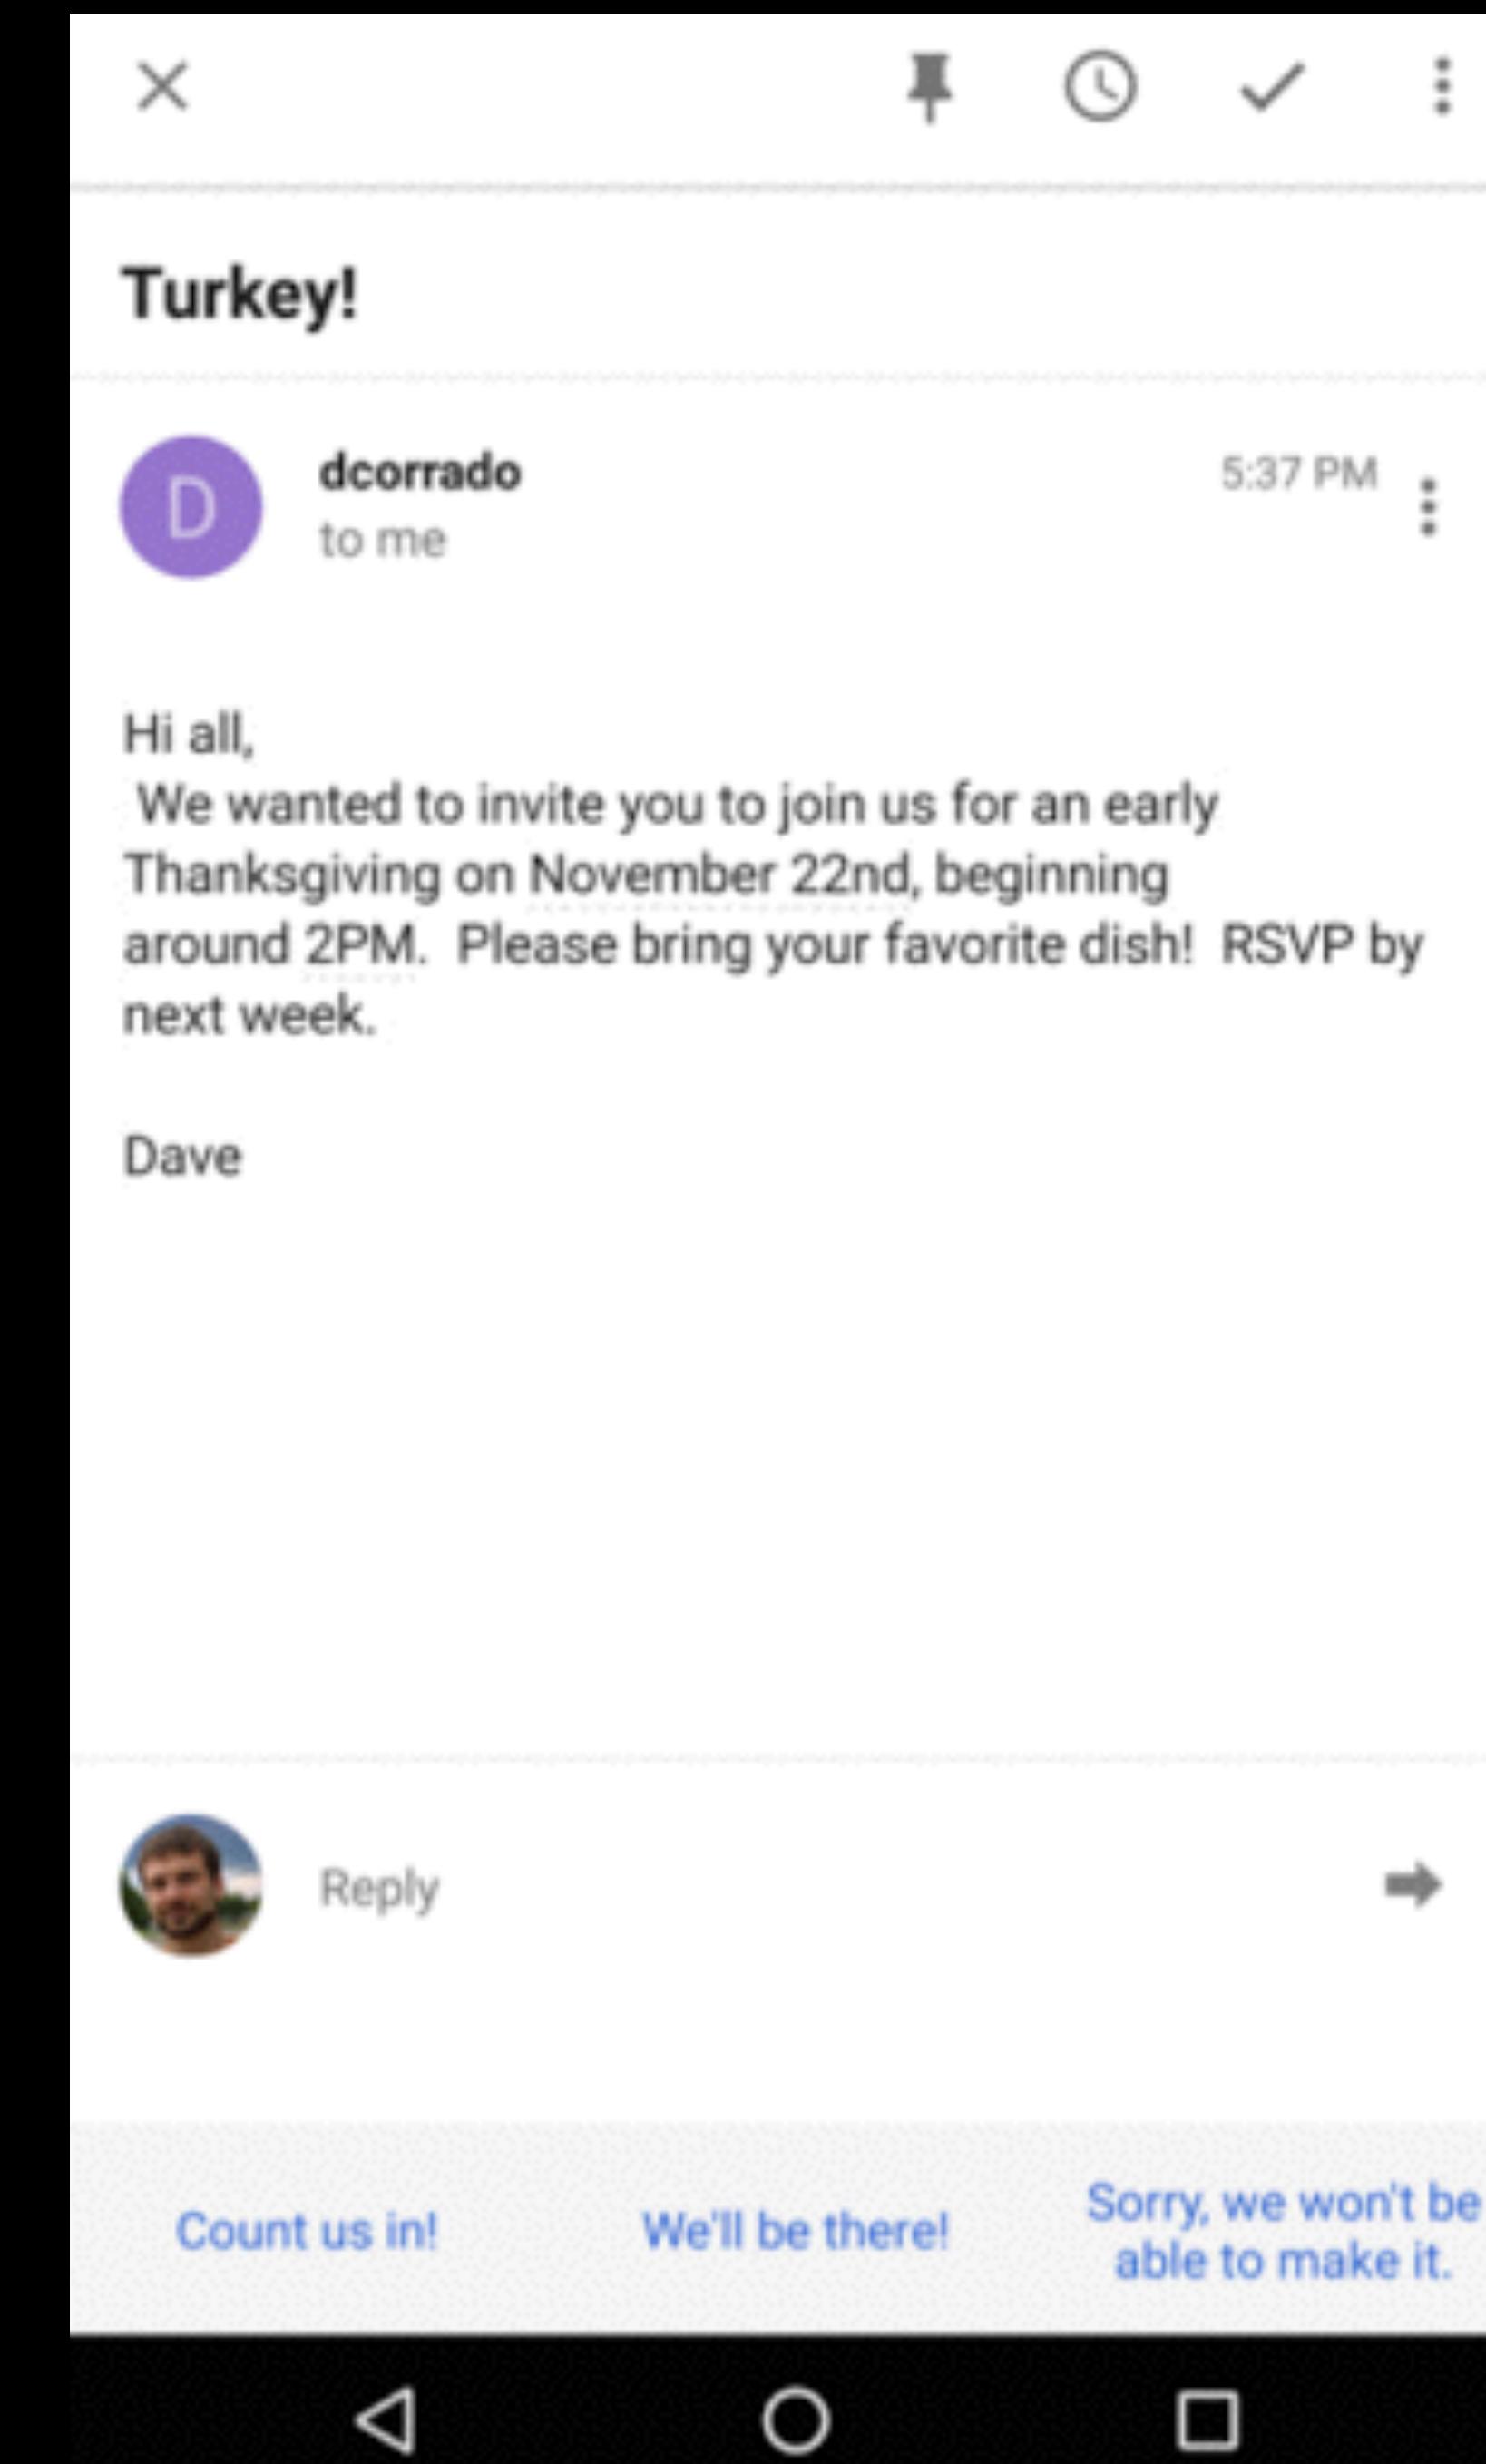
Sequence to Sequence Prediction



Scheduled Sampling



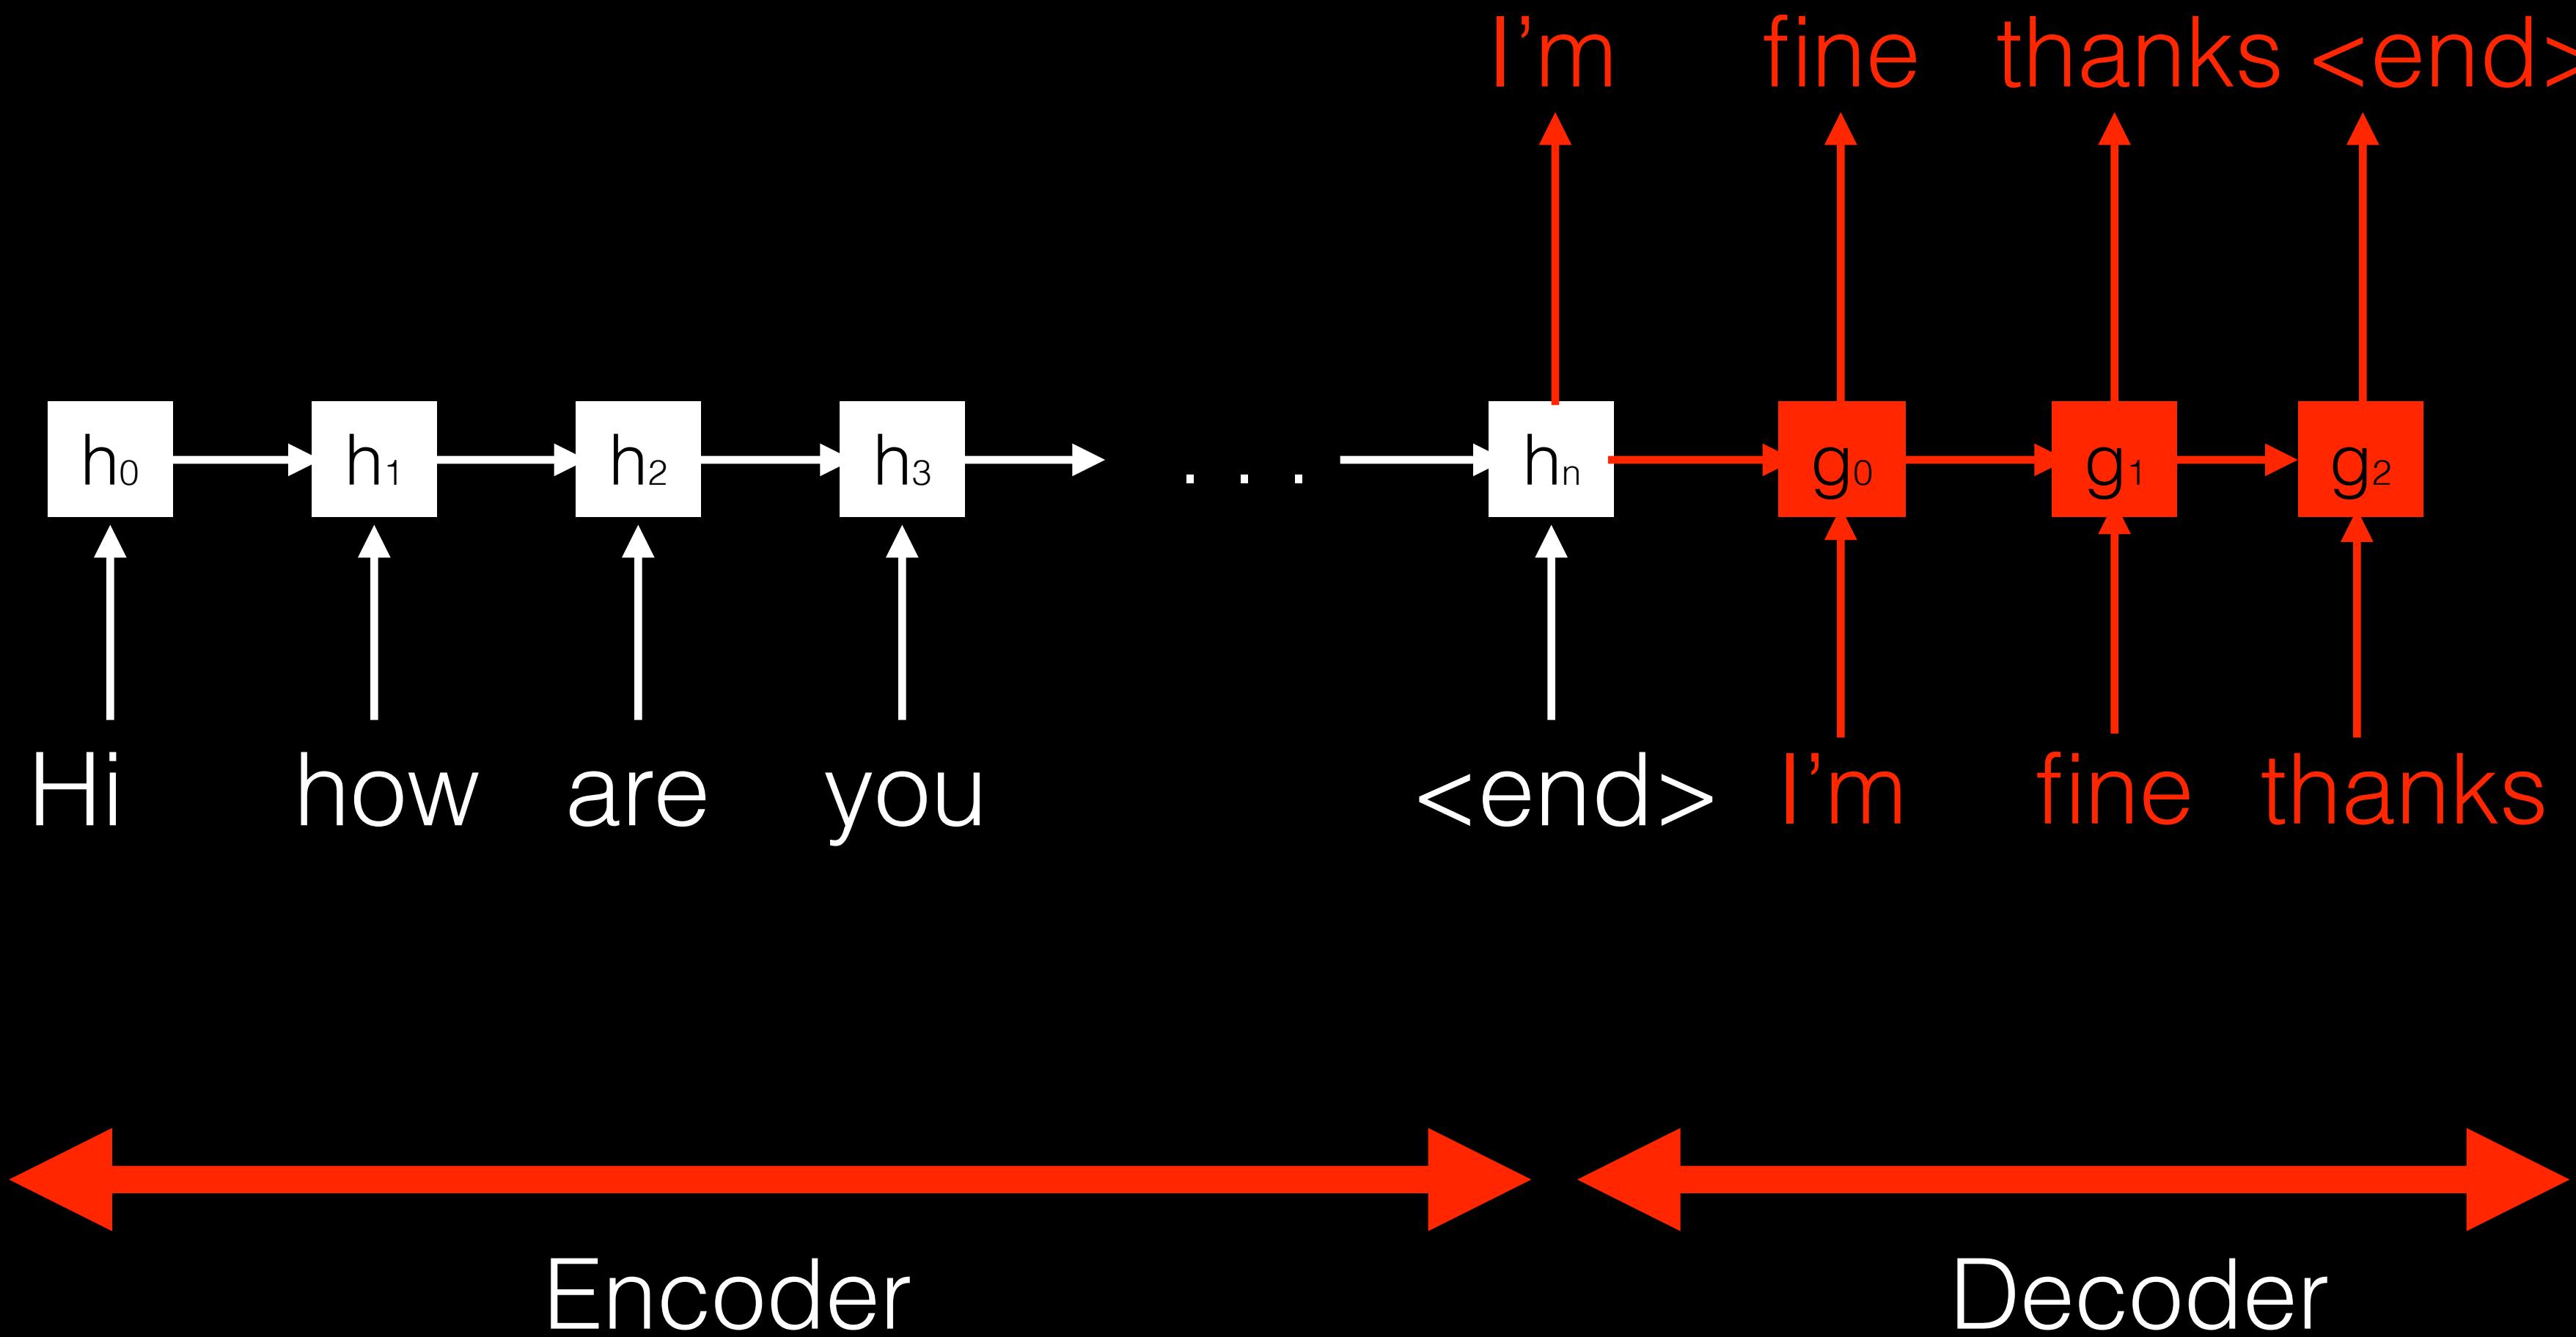
SmartReply feature in Inbox



The big picture so far

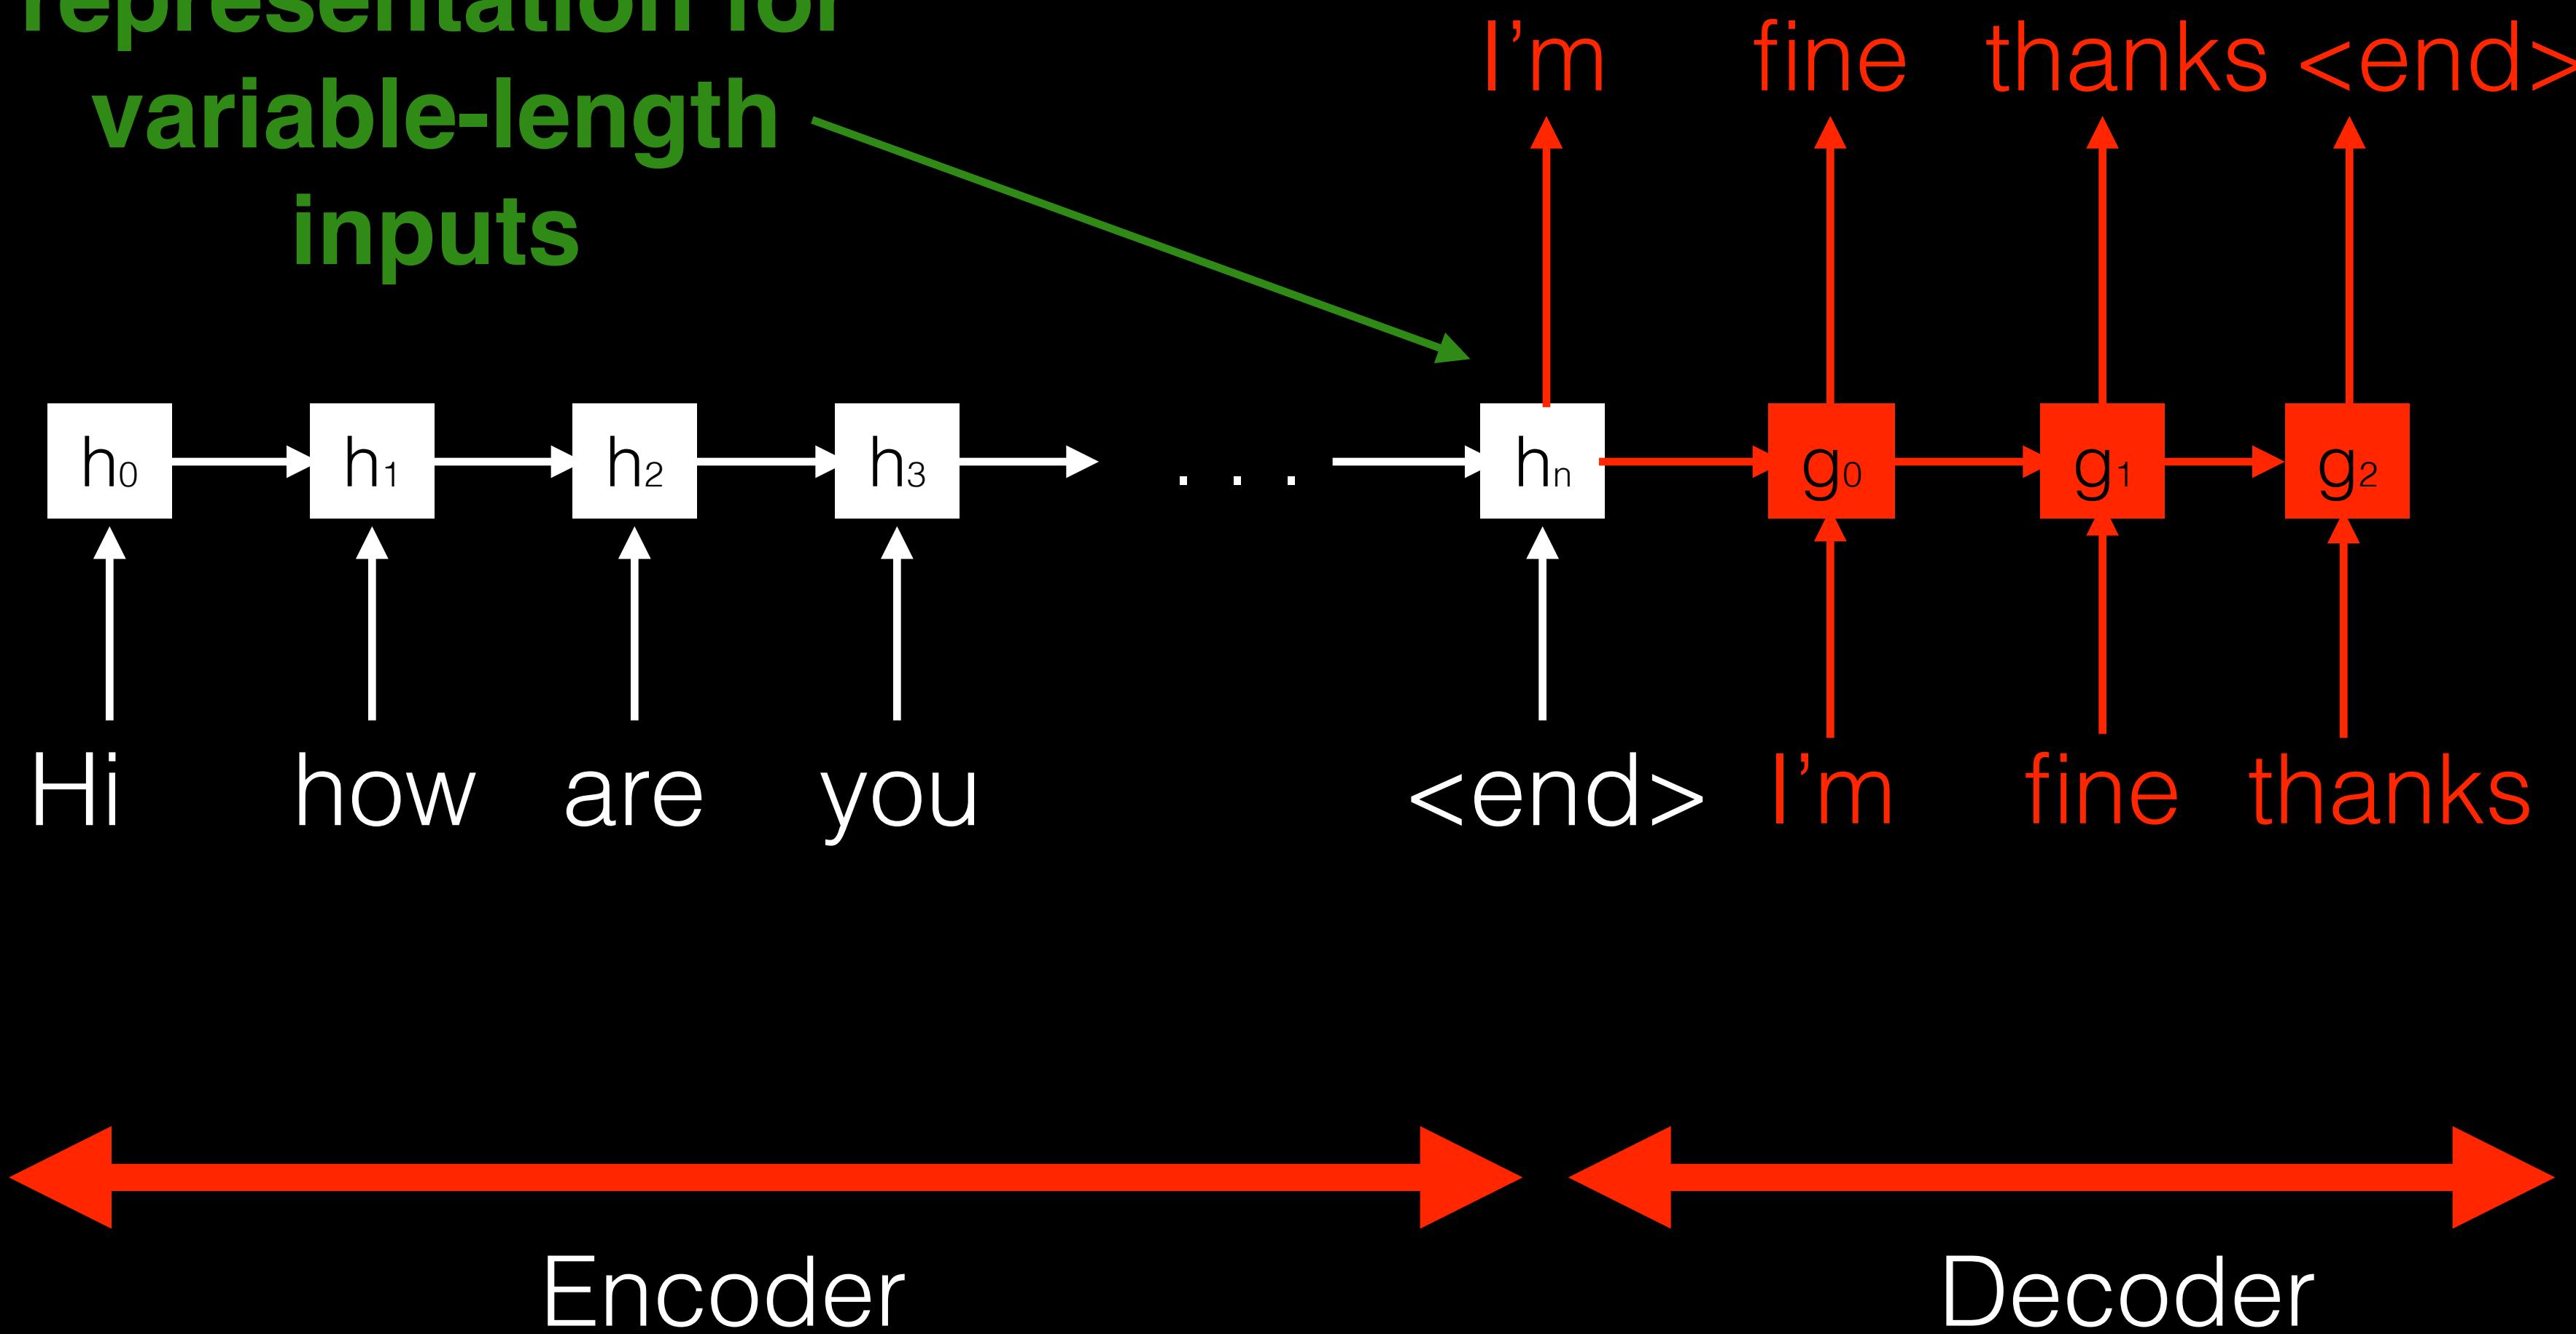
- RNN encoder and RNN decoder for sequence to sequence learning
- Use stochastic gradient descent for training
- Beam search decoding

Attention Mechanism

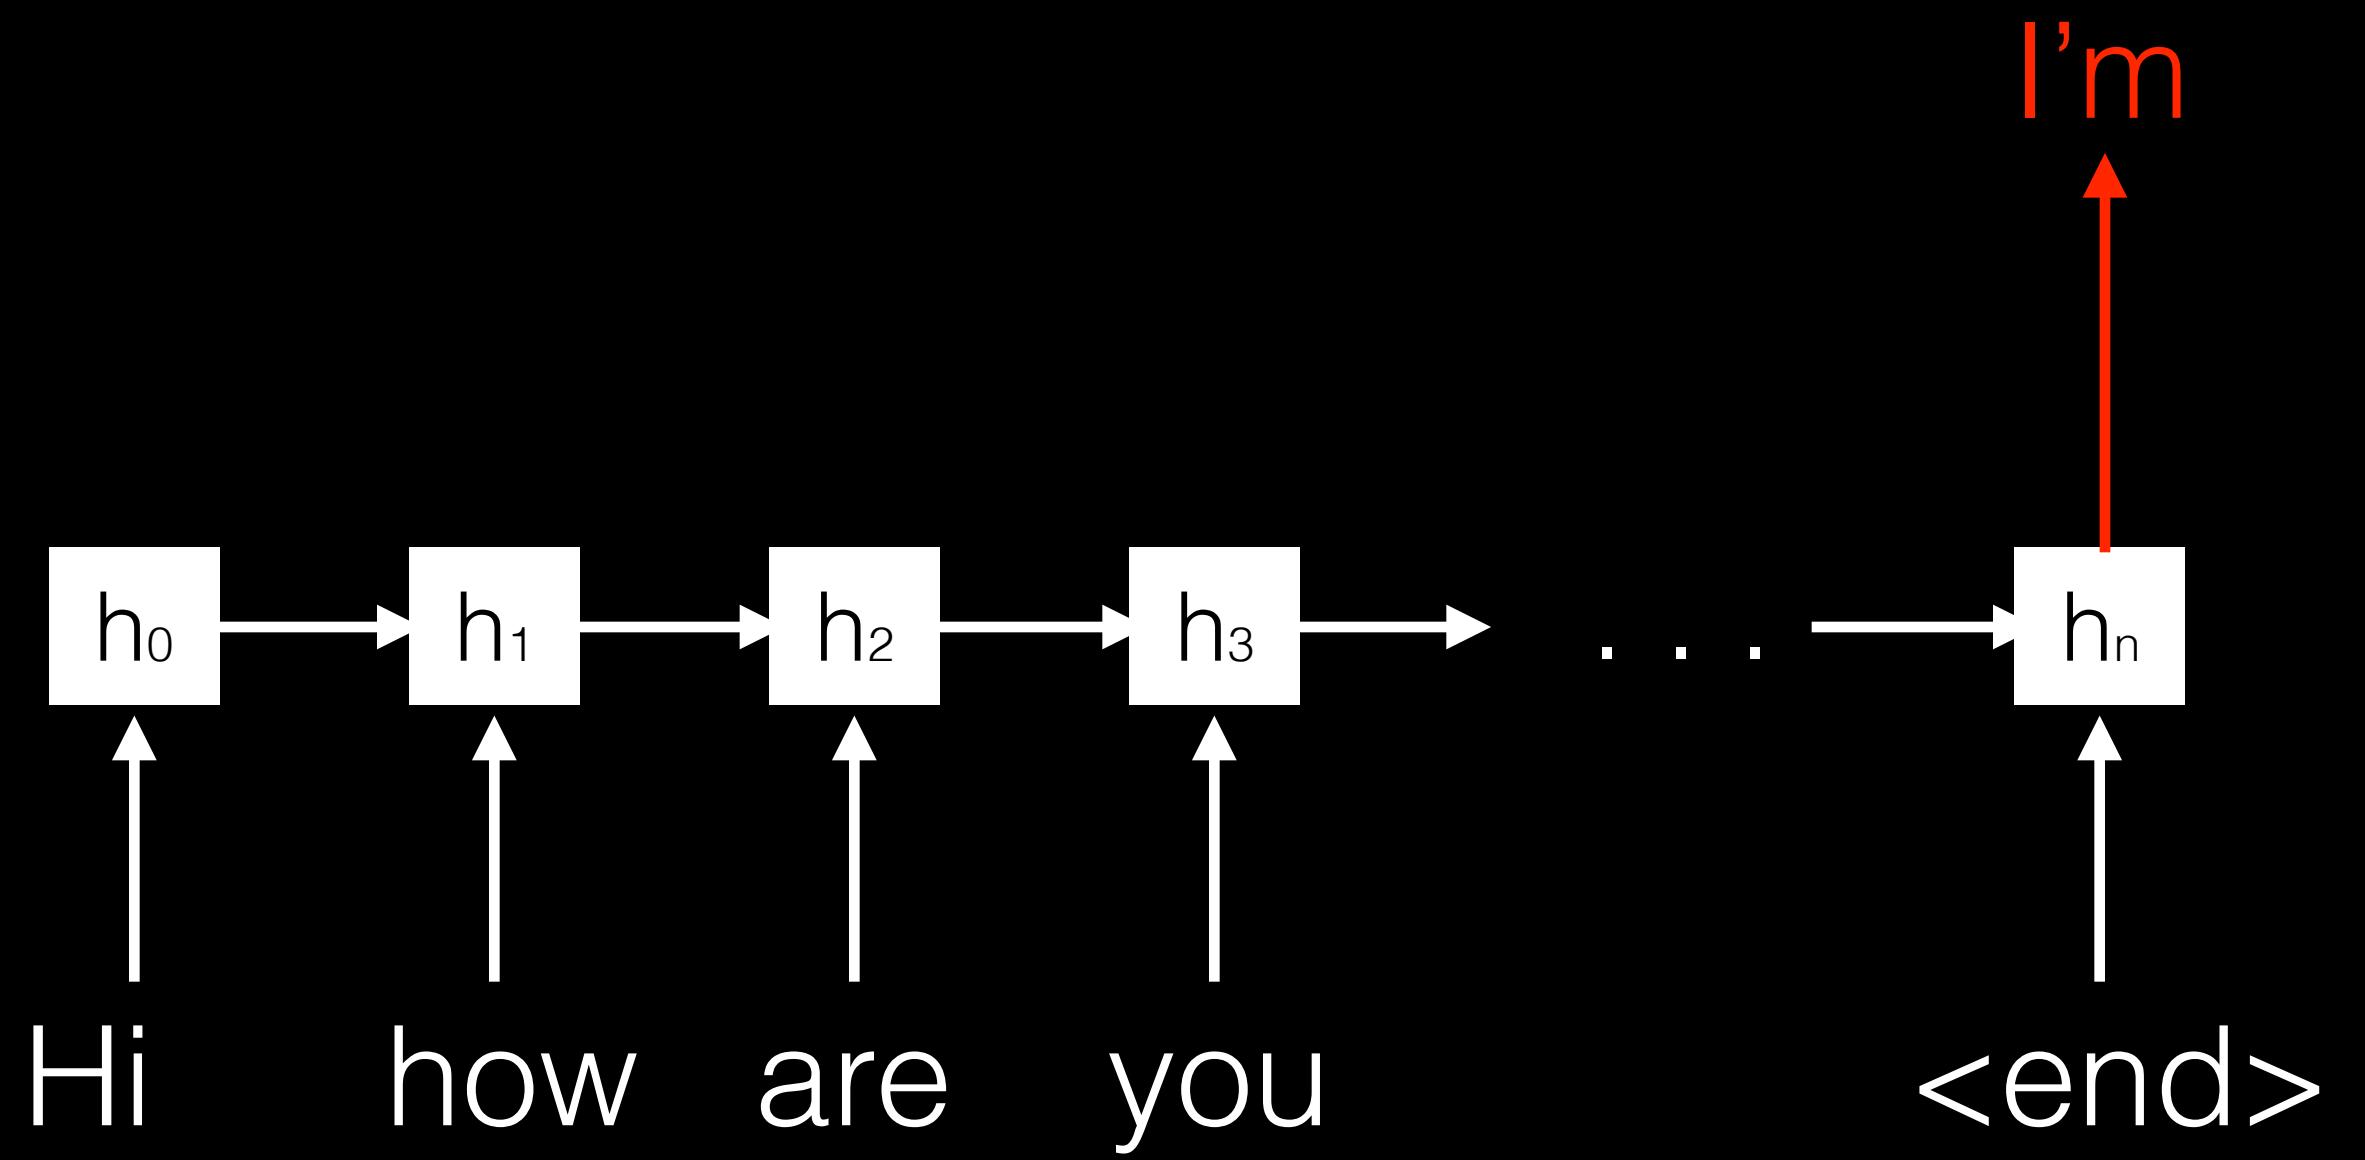


Attention Mechanism

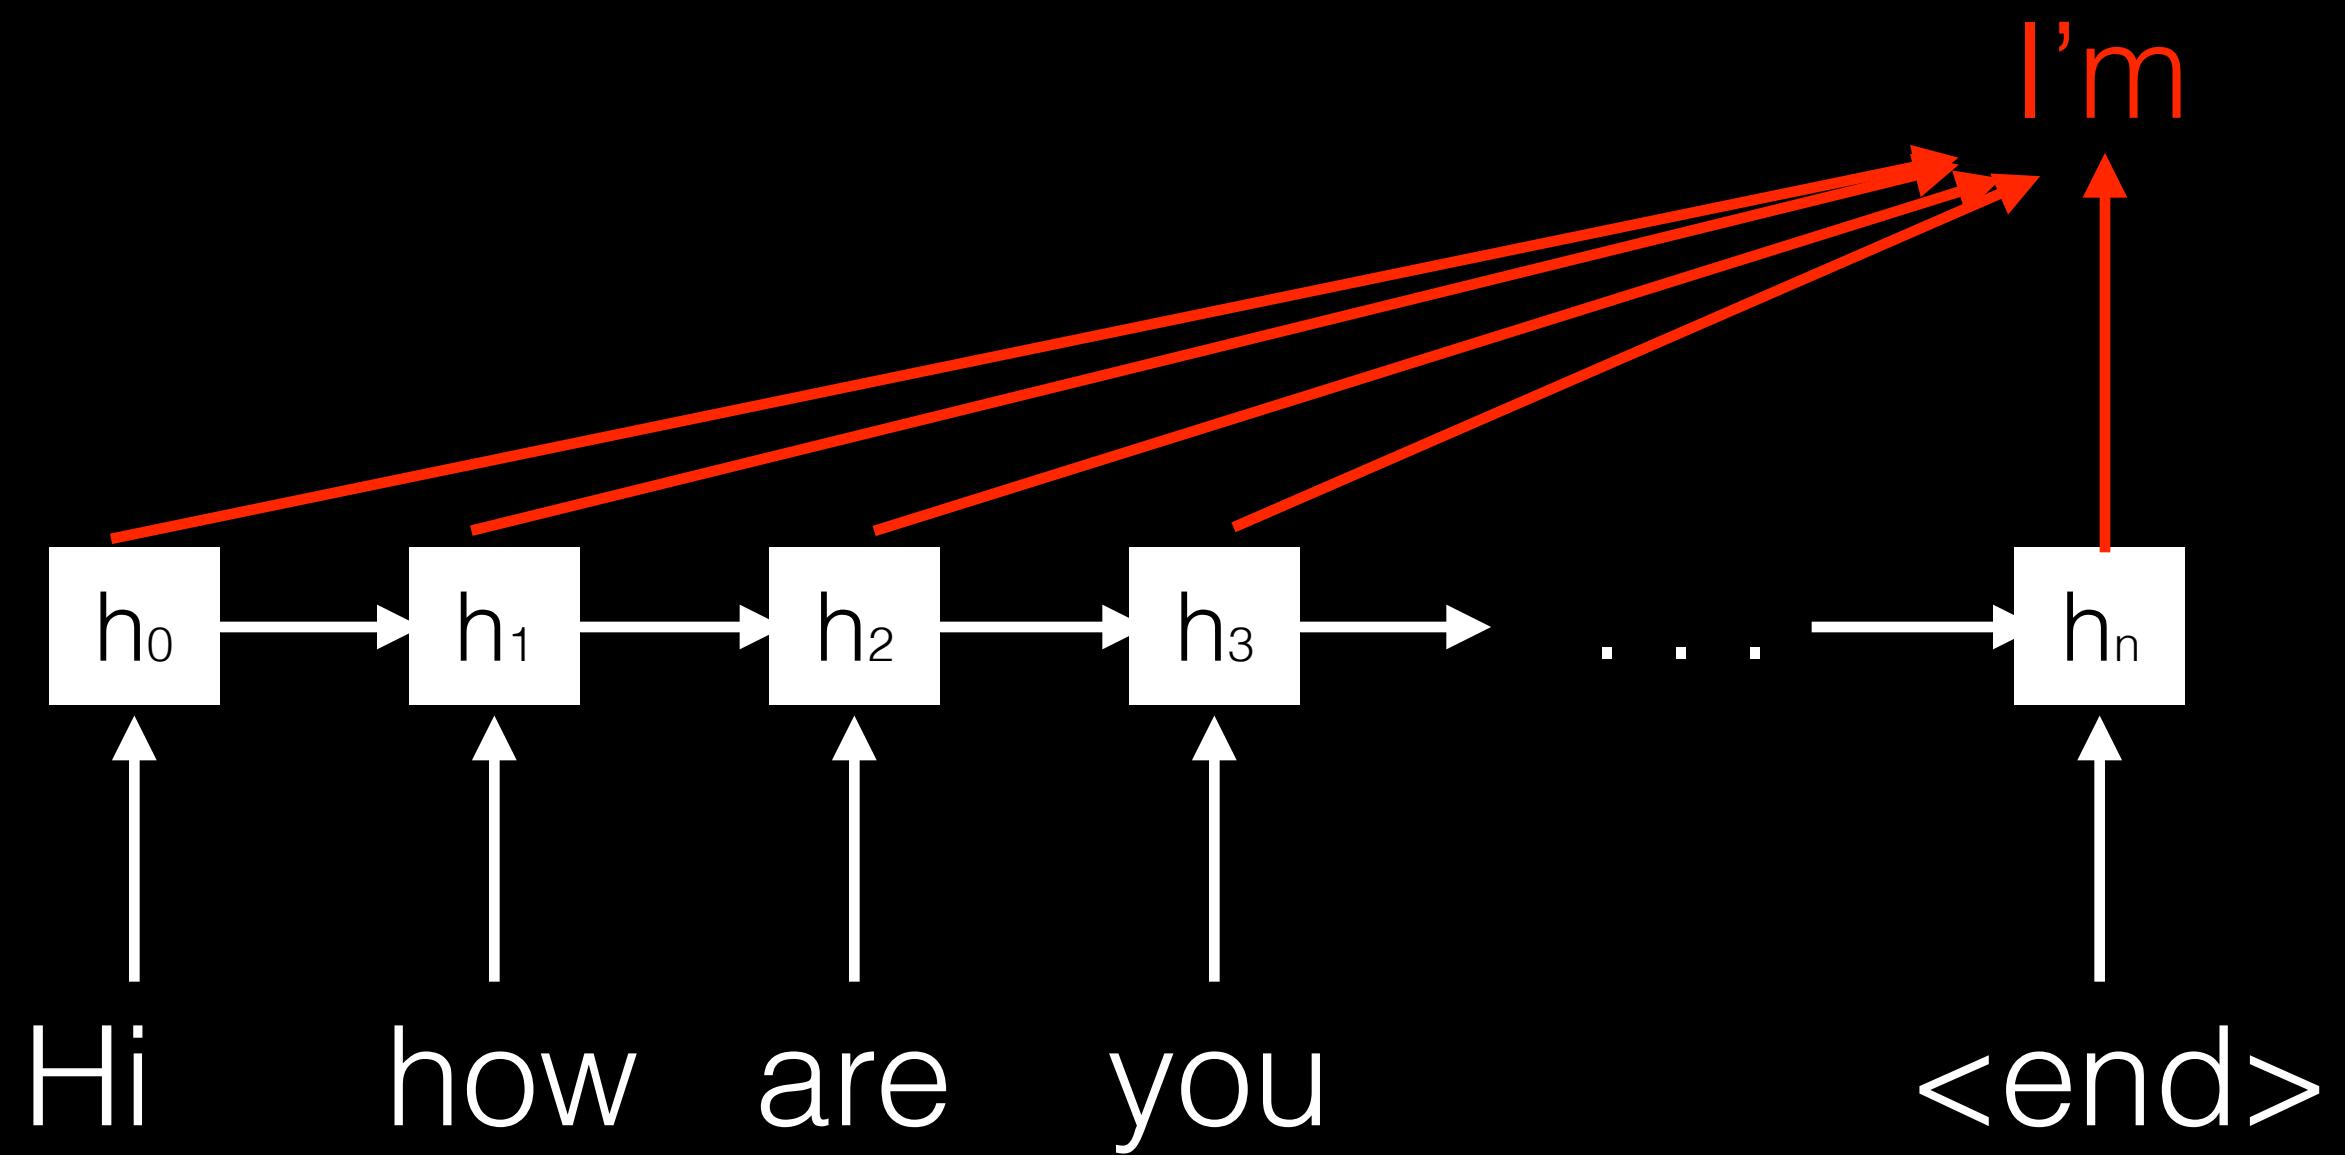
**Fixed-length
representation for
variable-length
inputs**



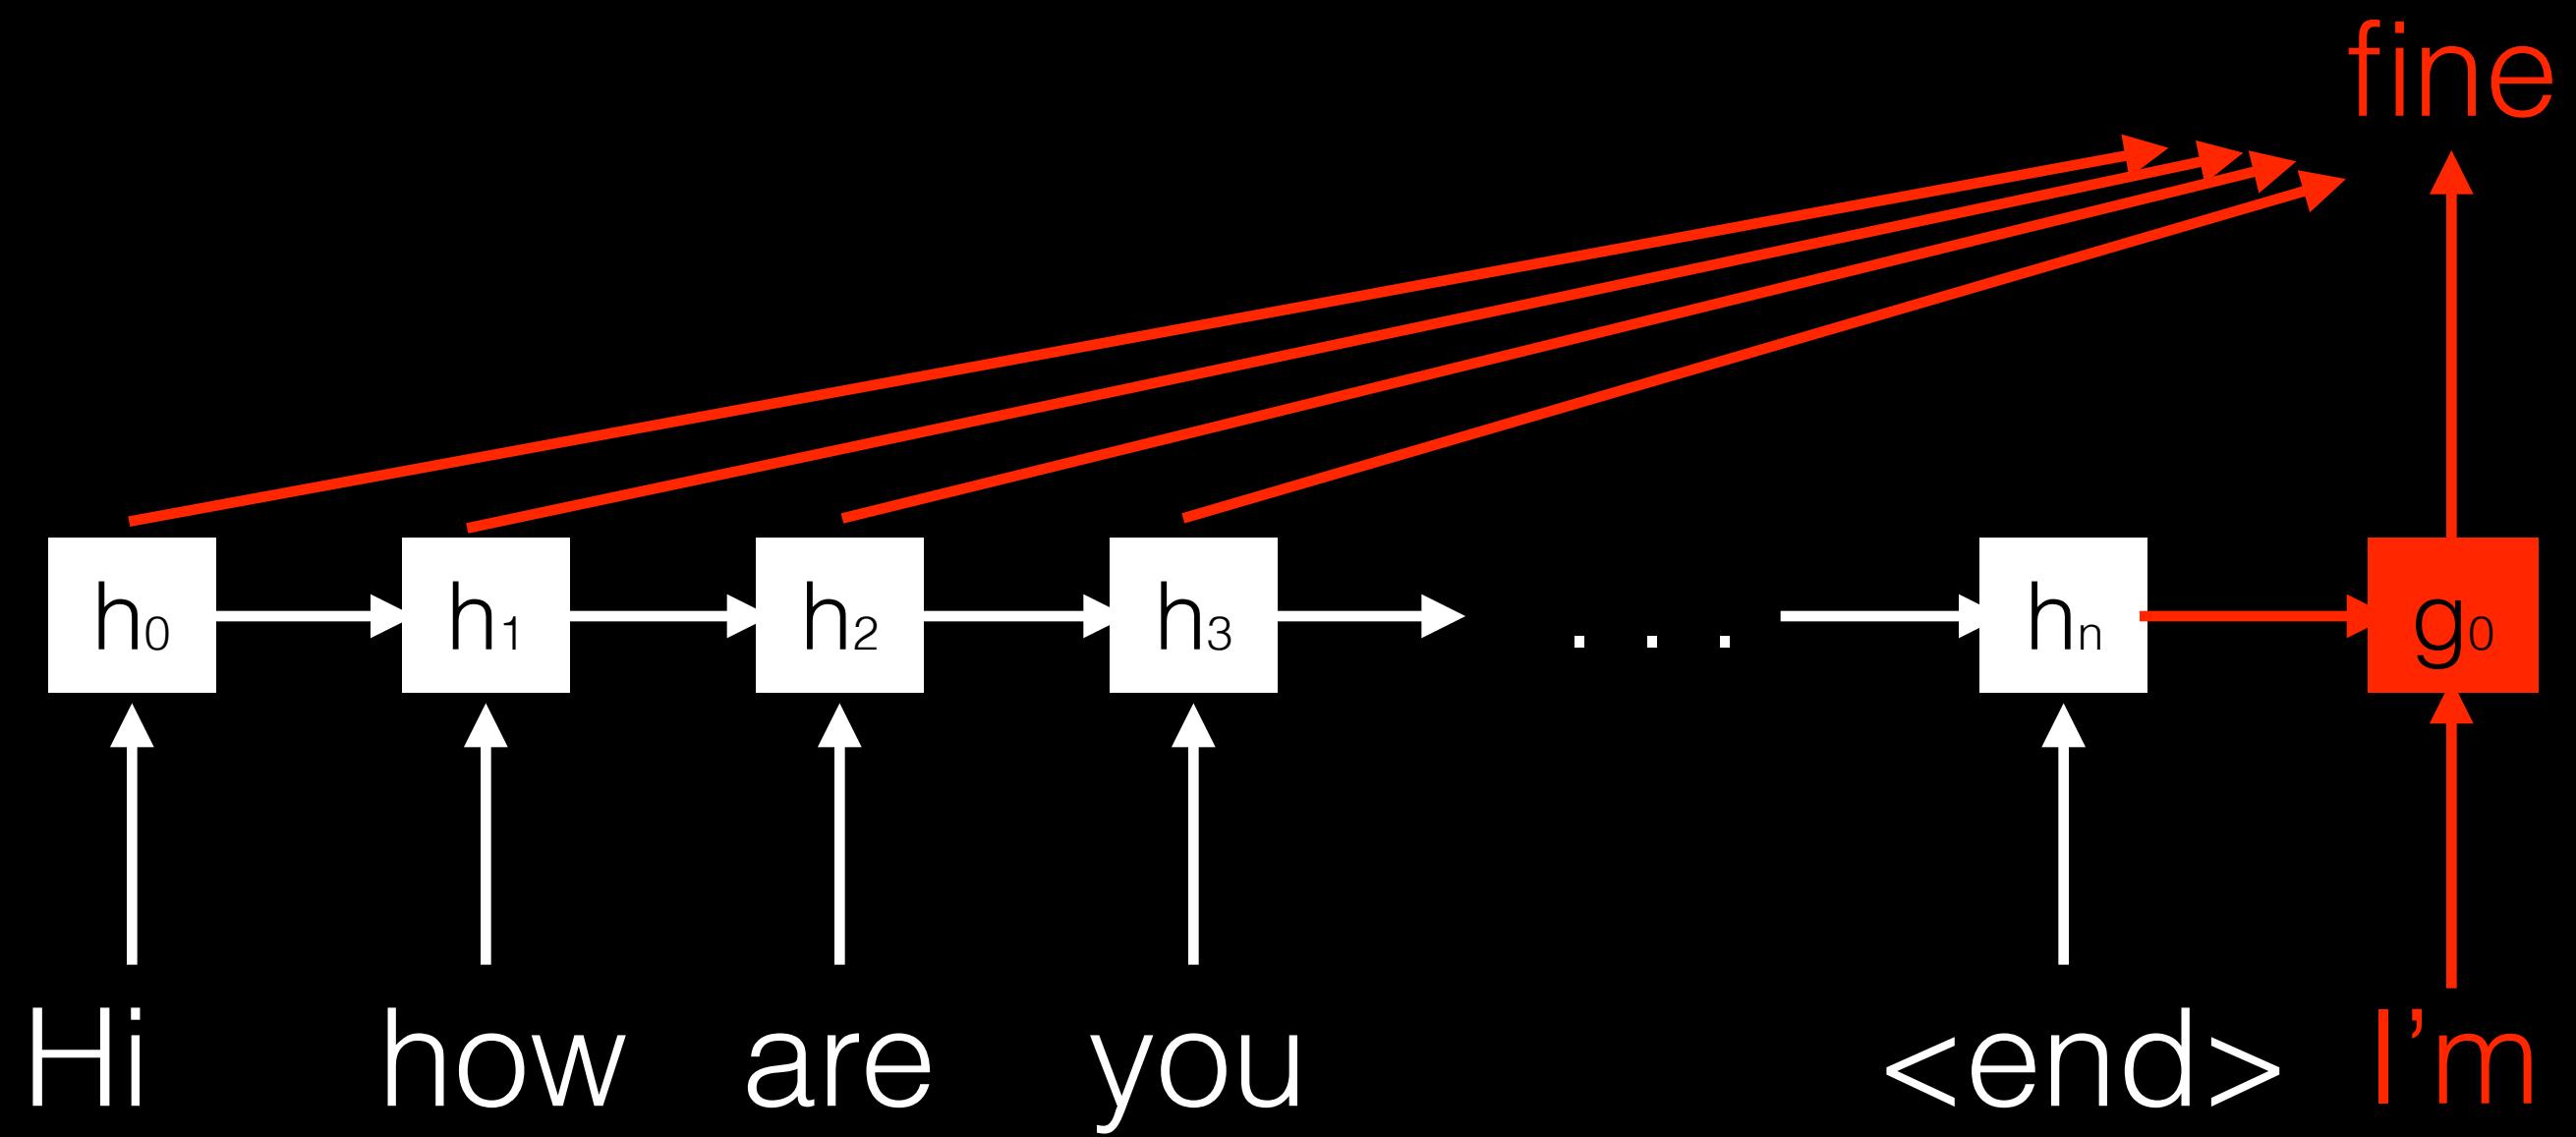
Attention Mechanism



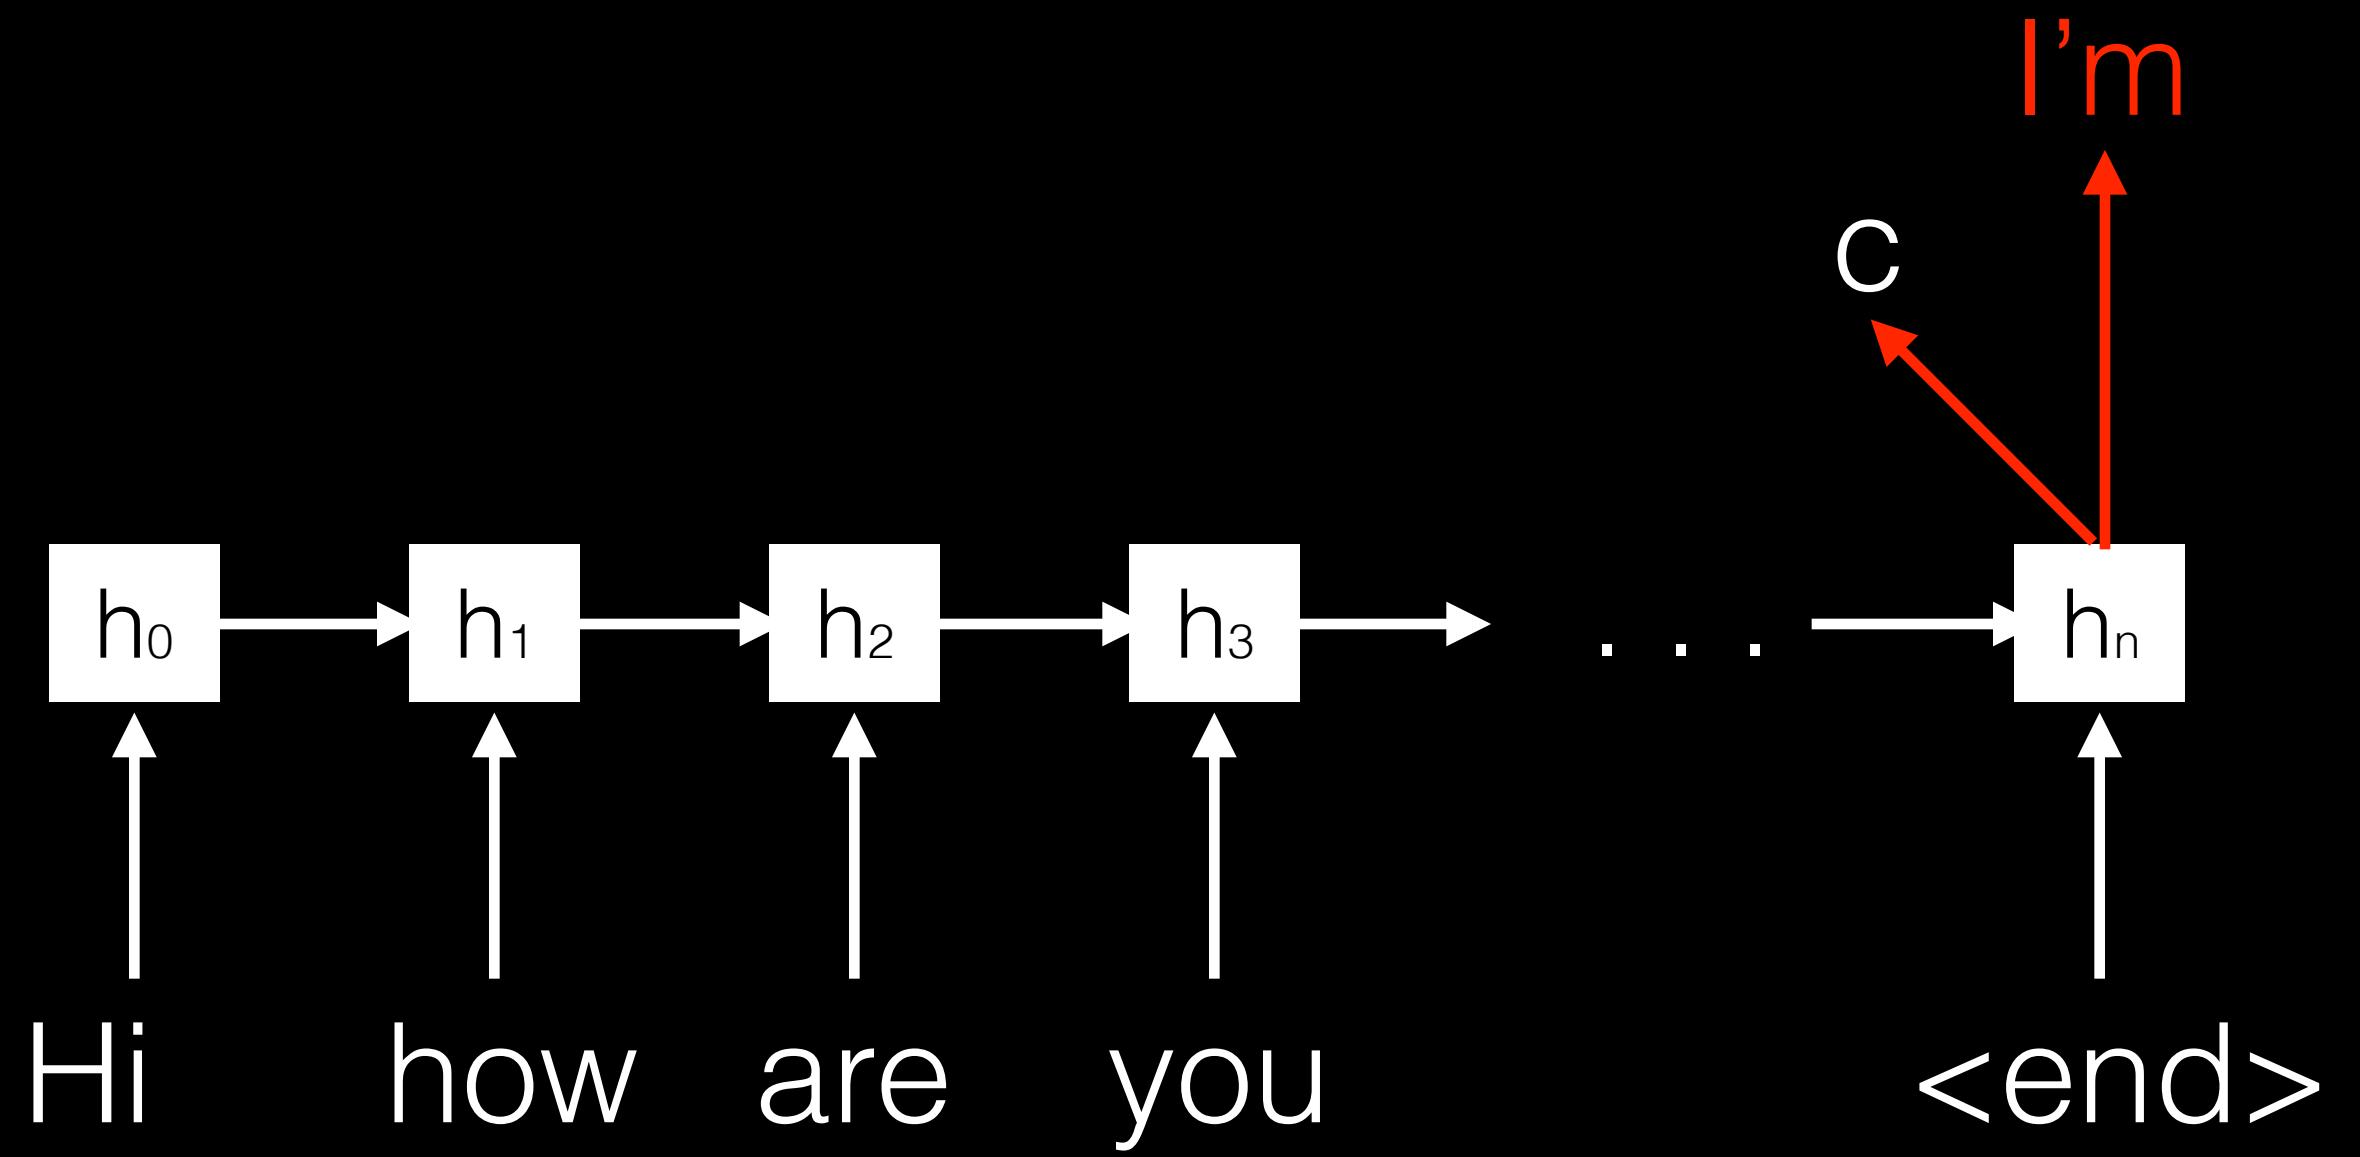
Attention Mechanism



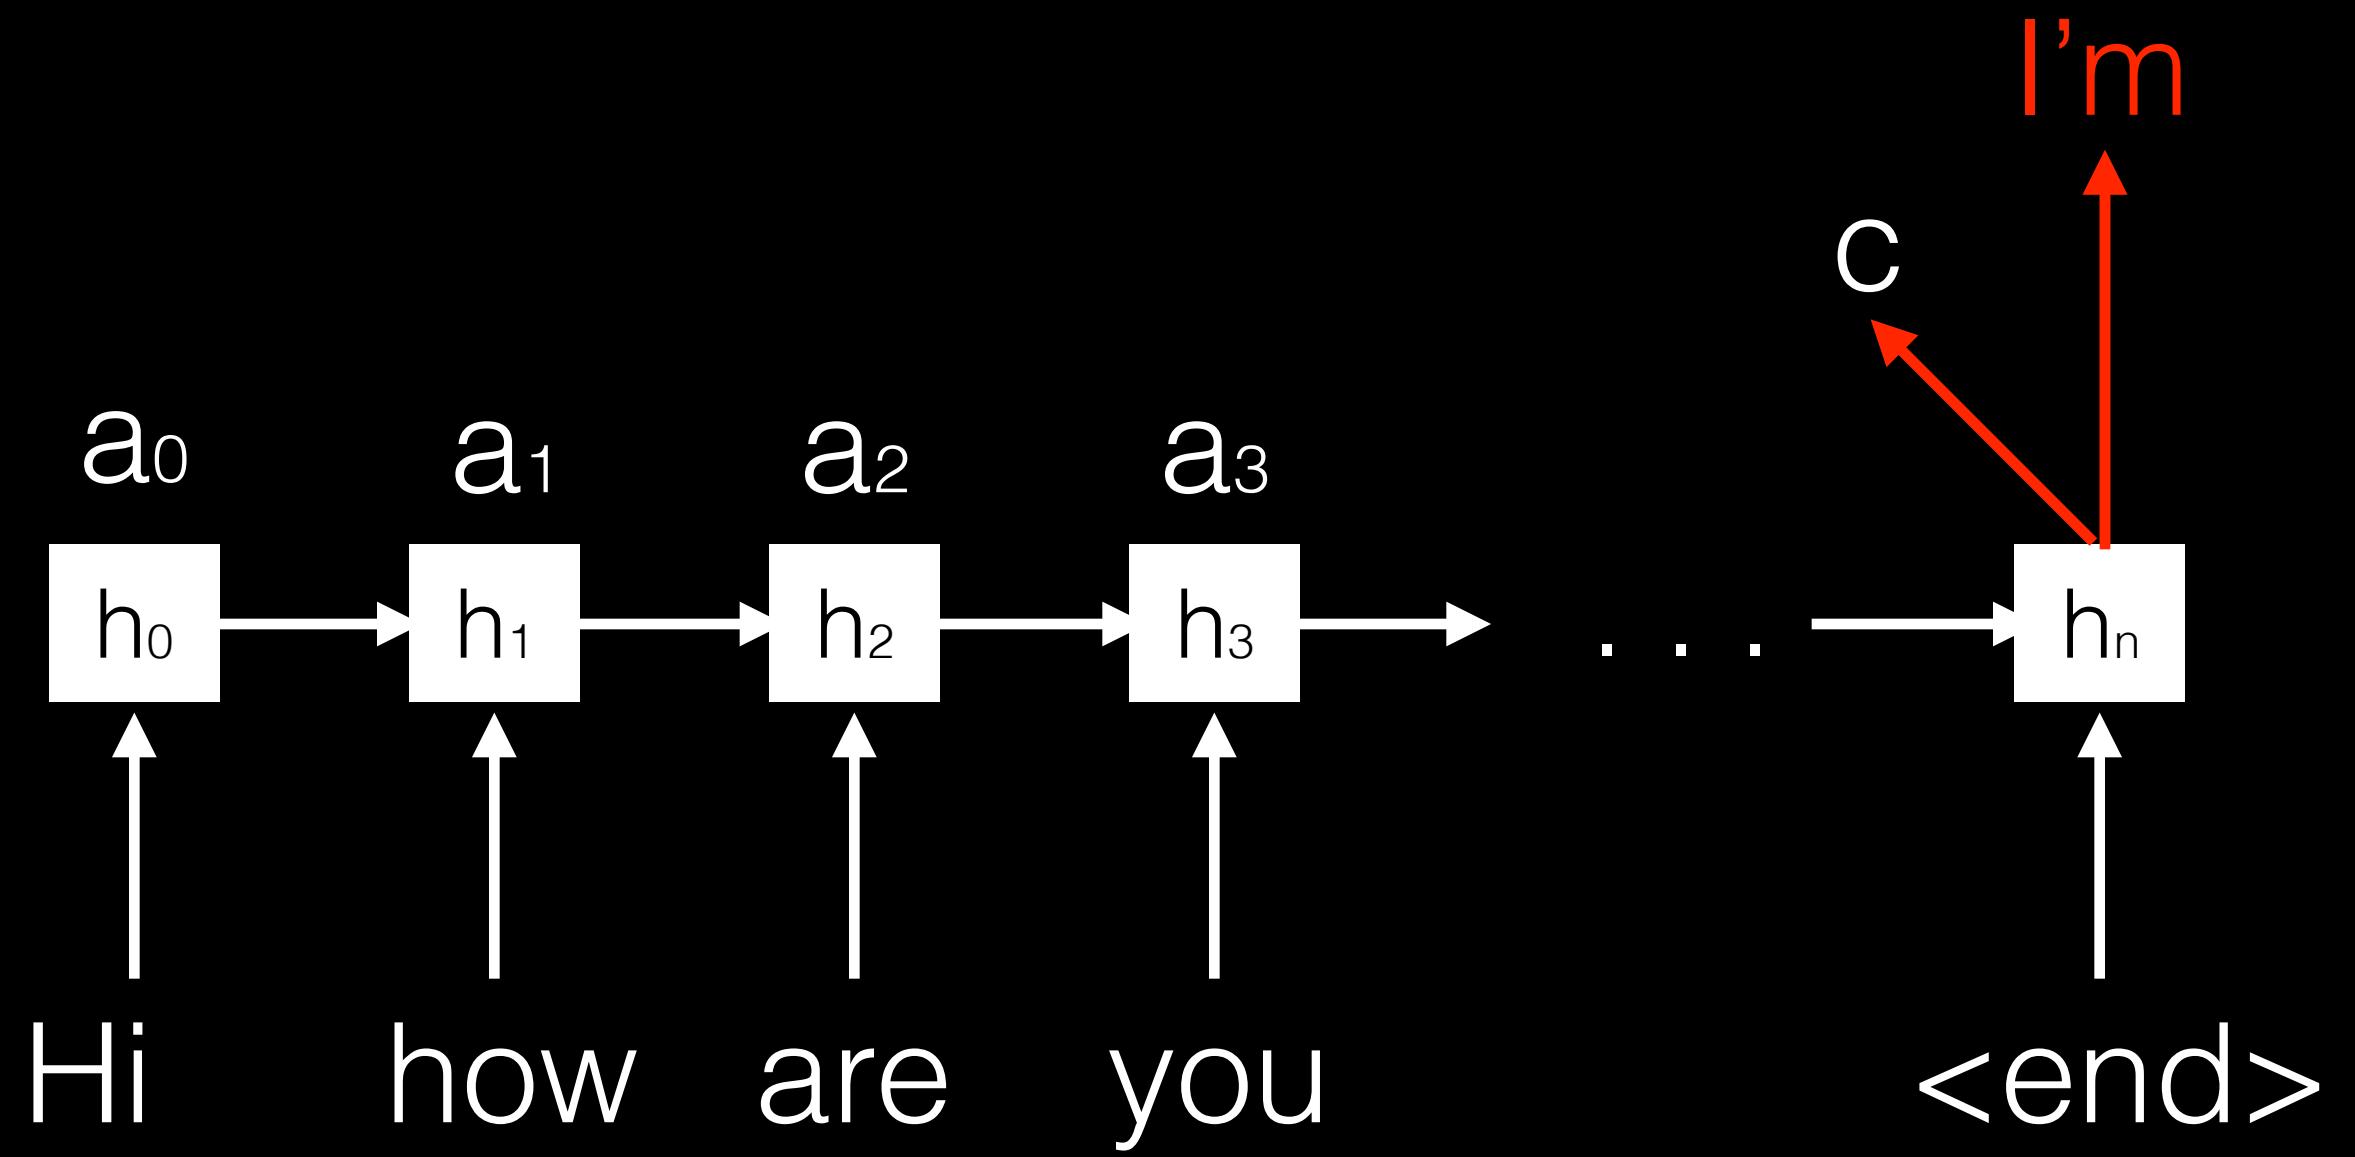
Attention Mechanism



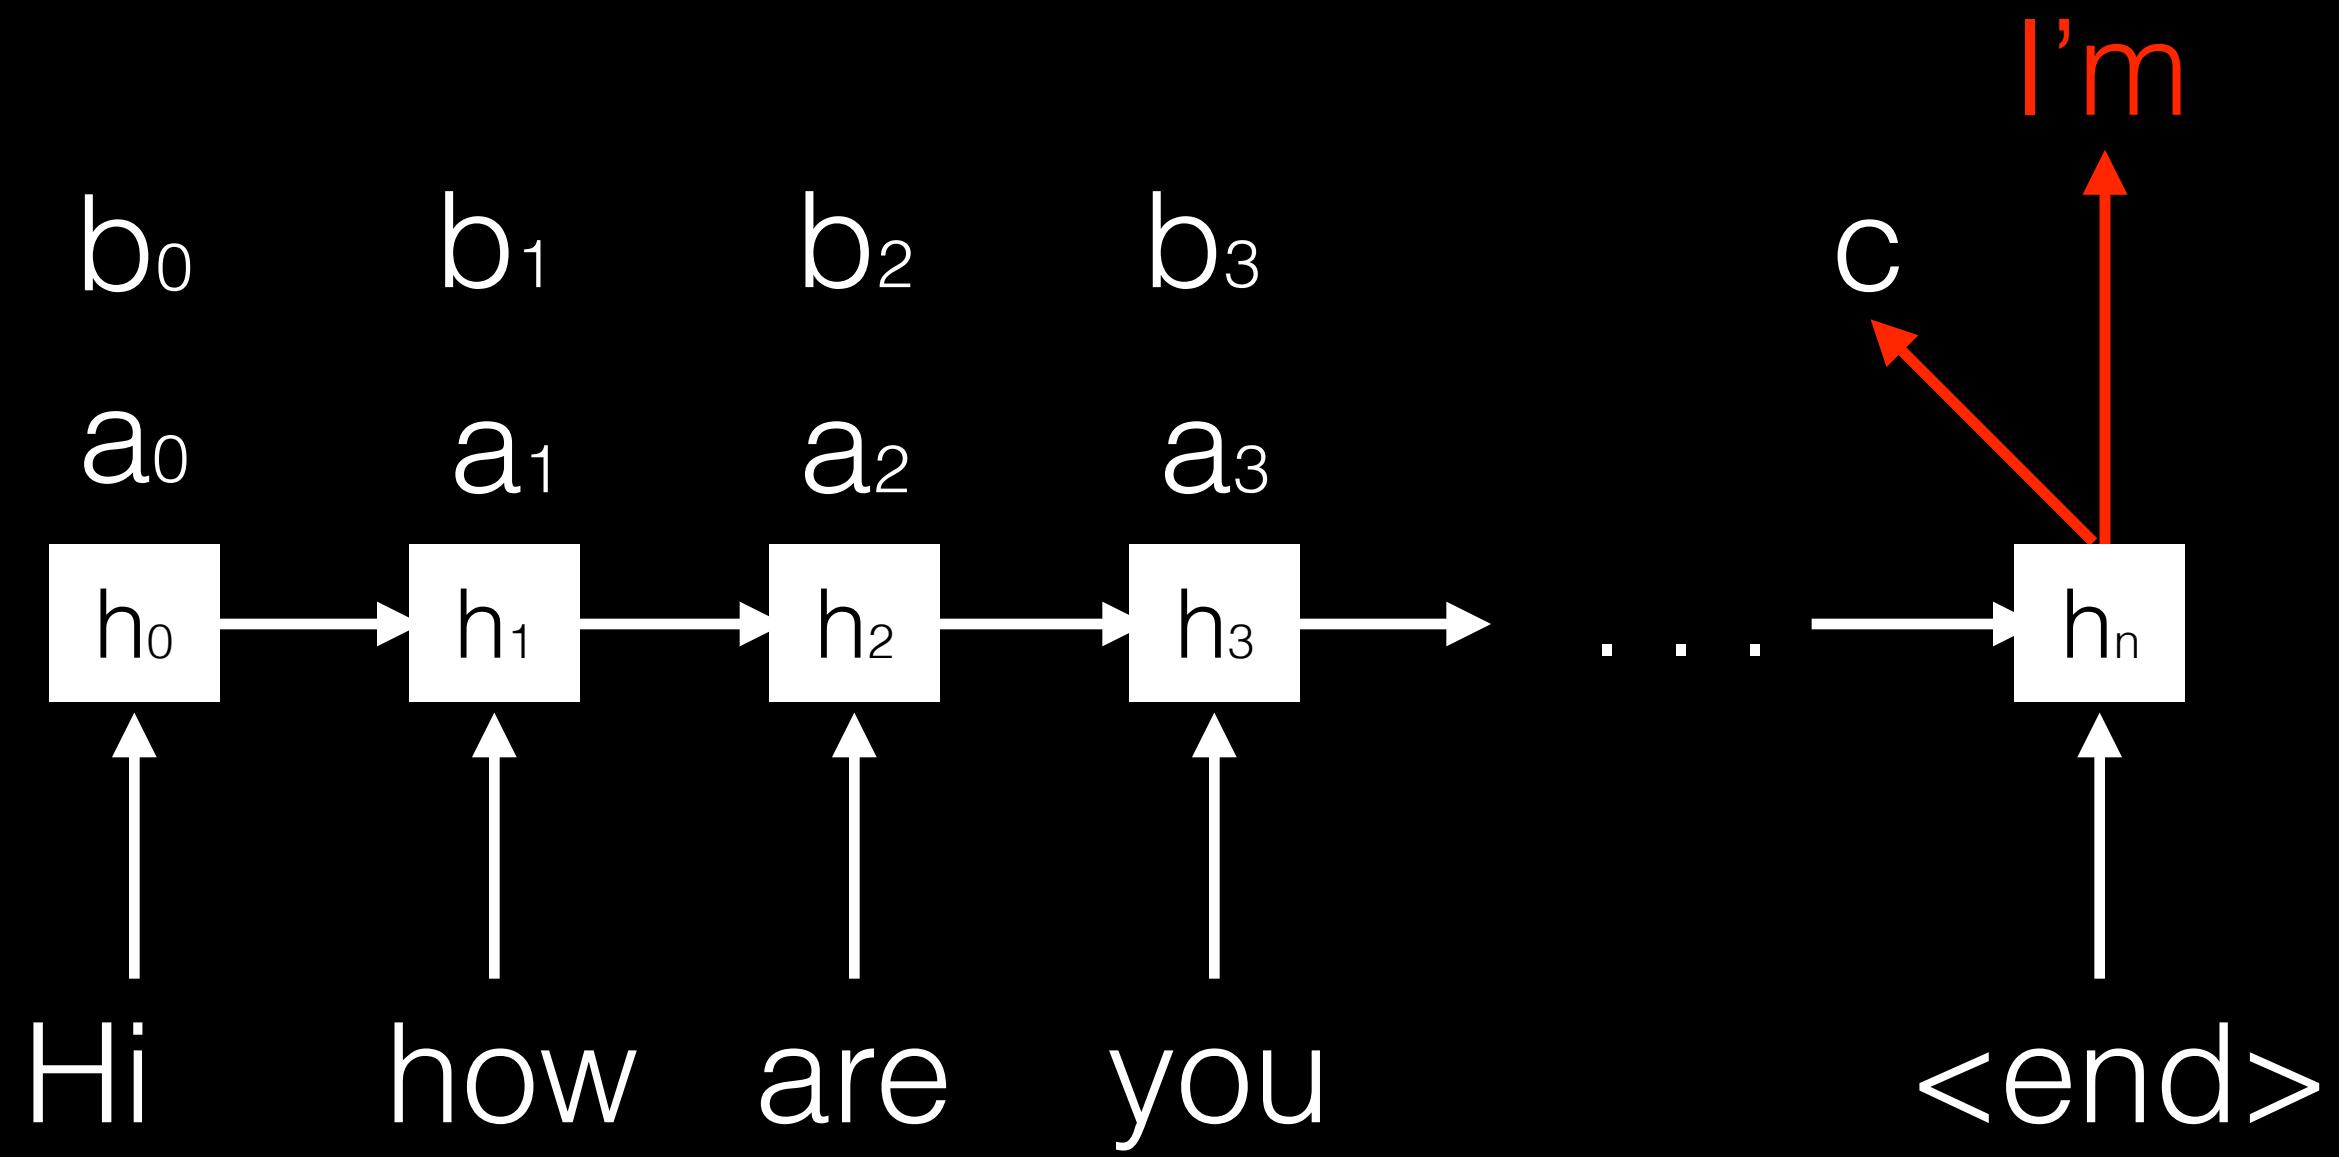
Attention Mechanism



Attention Mechanism

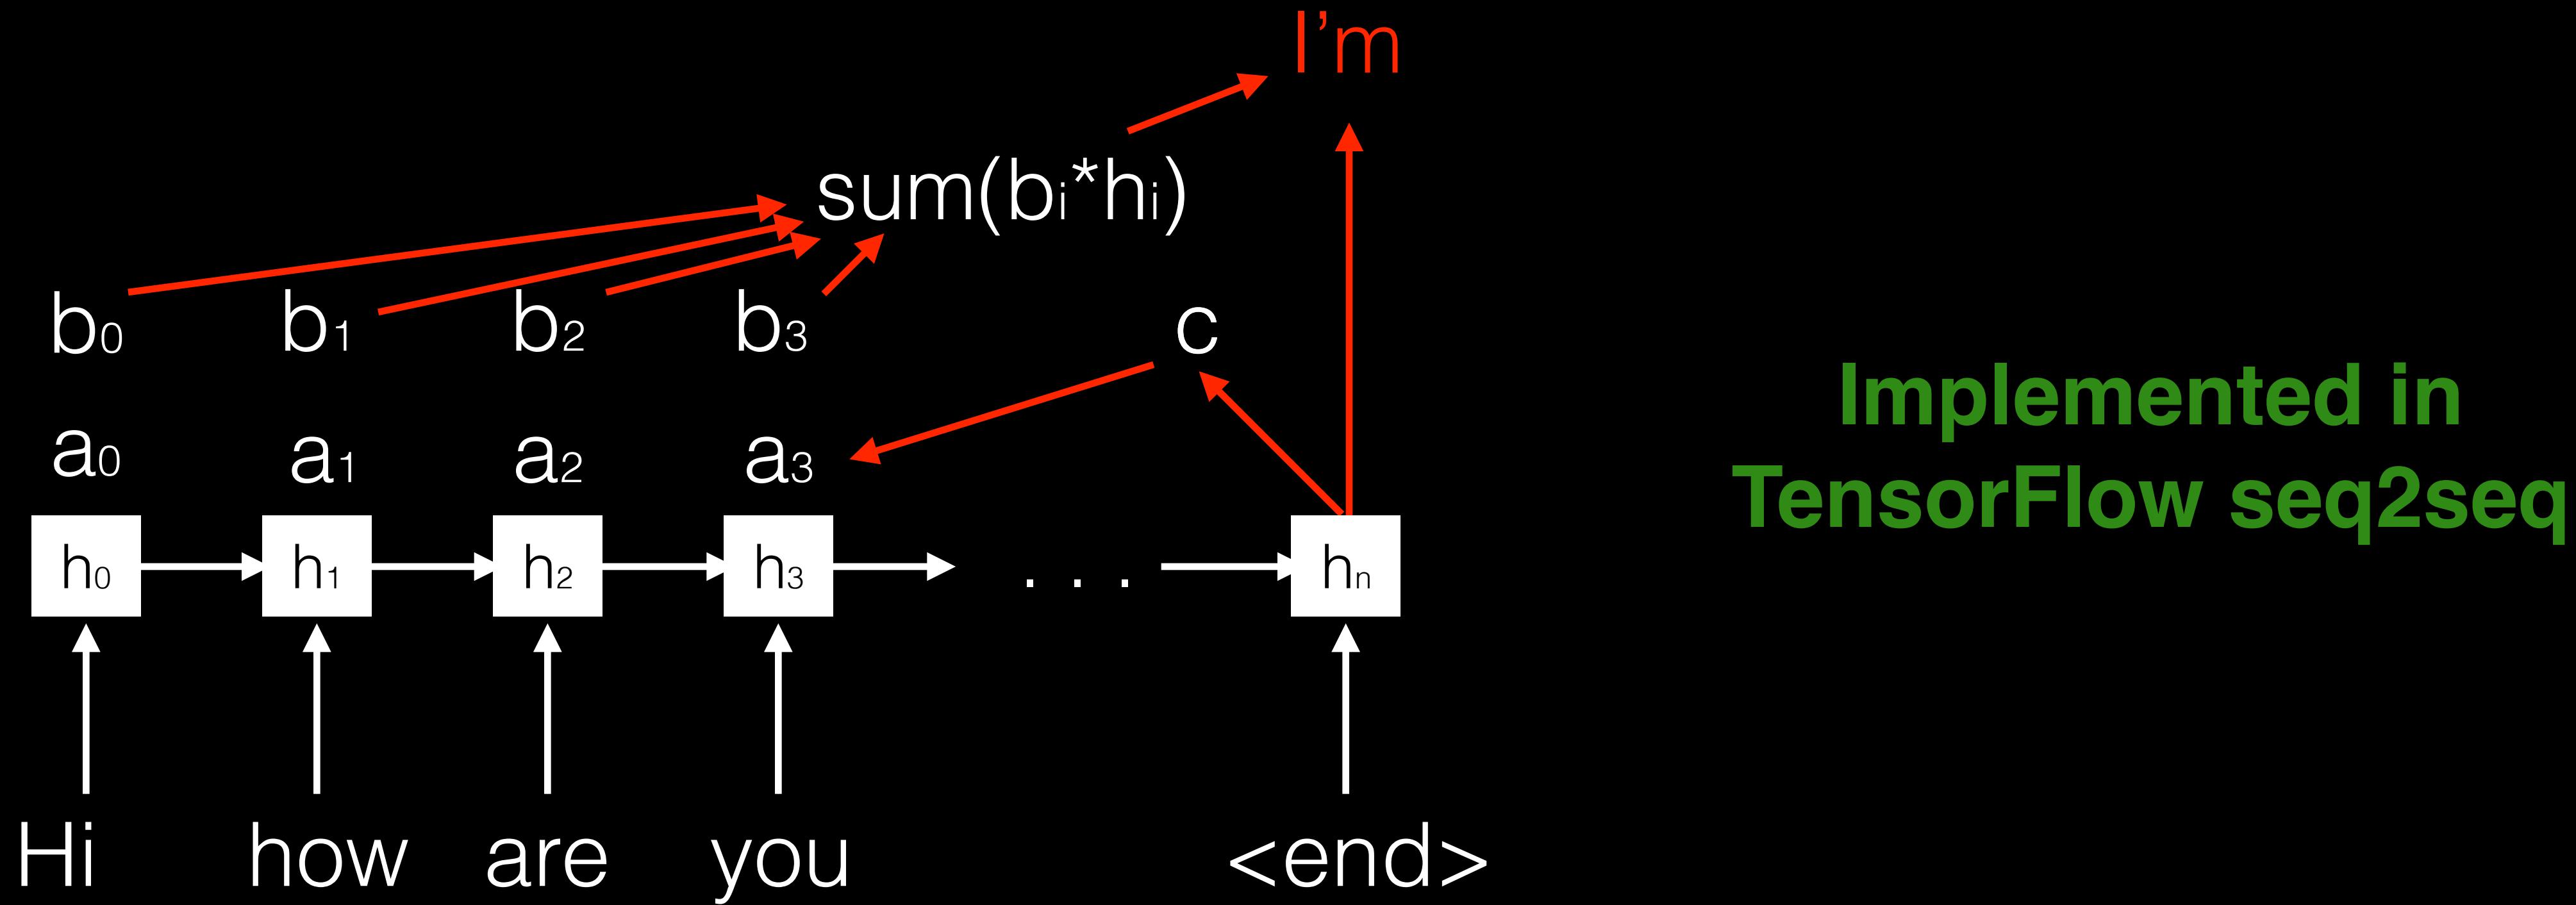


Attention Mechanism

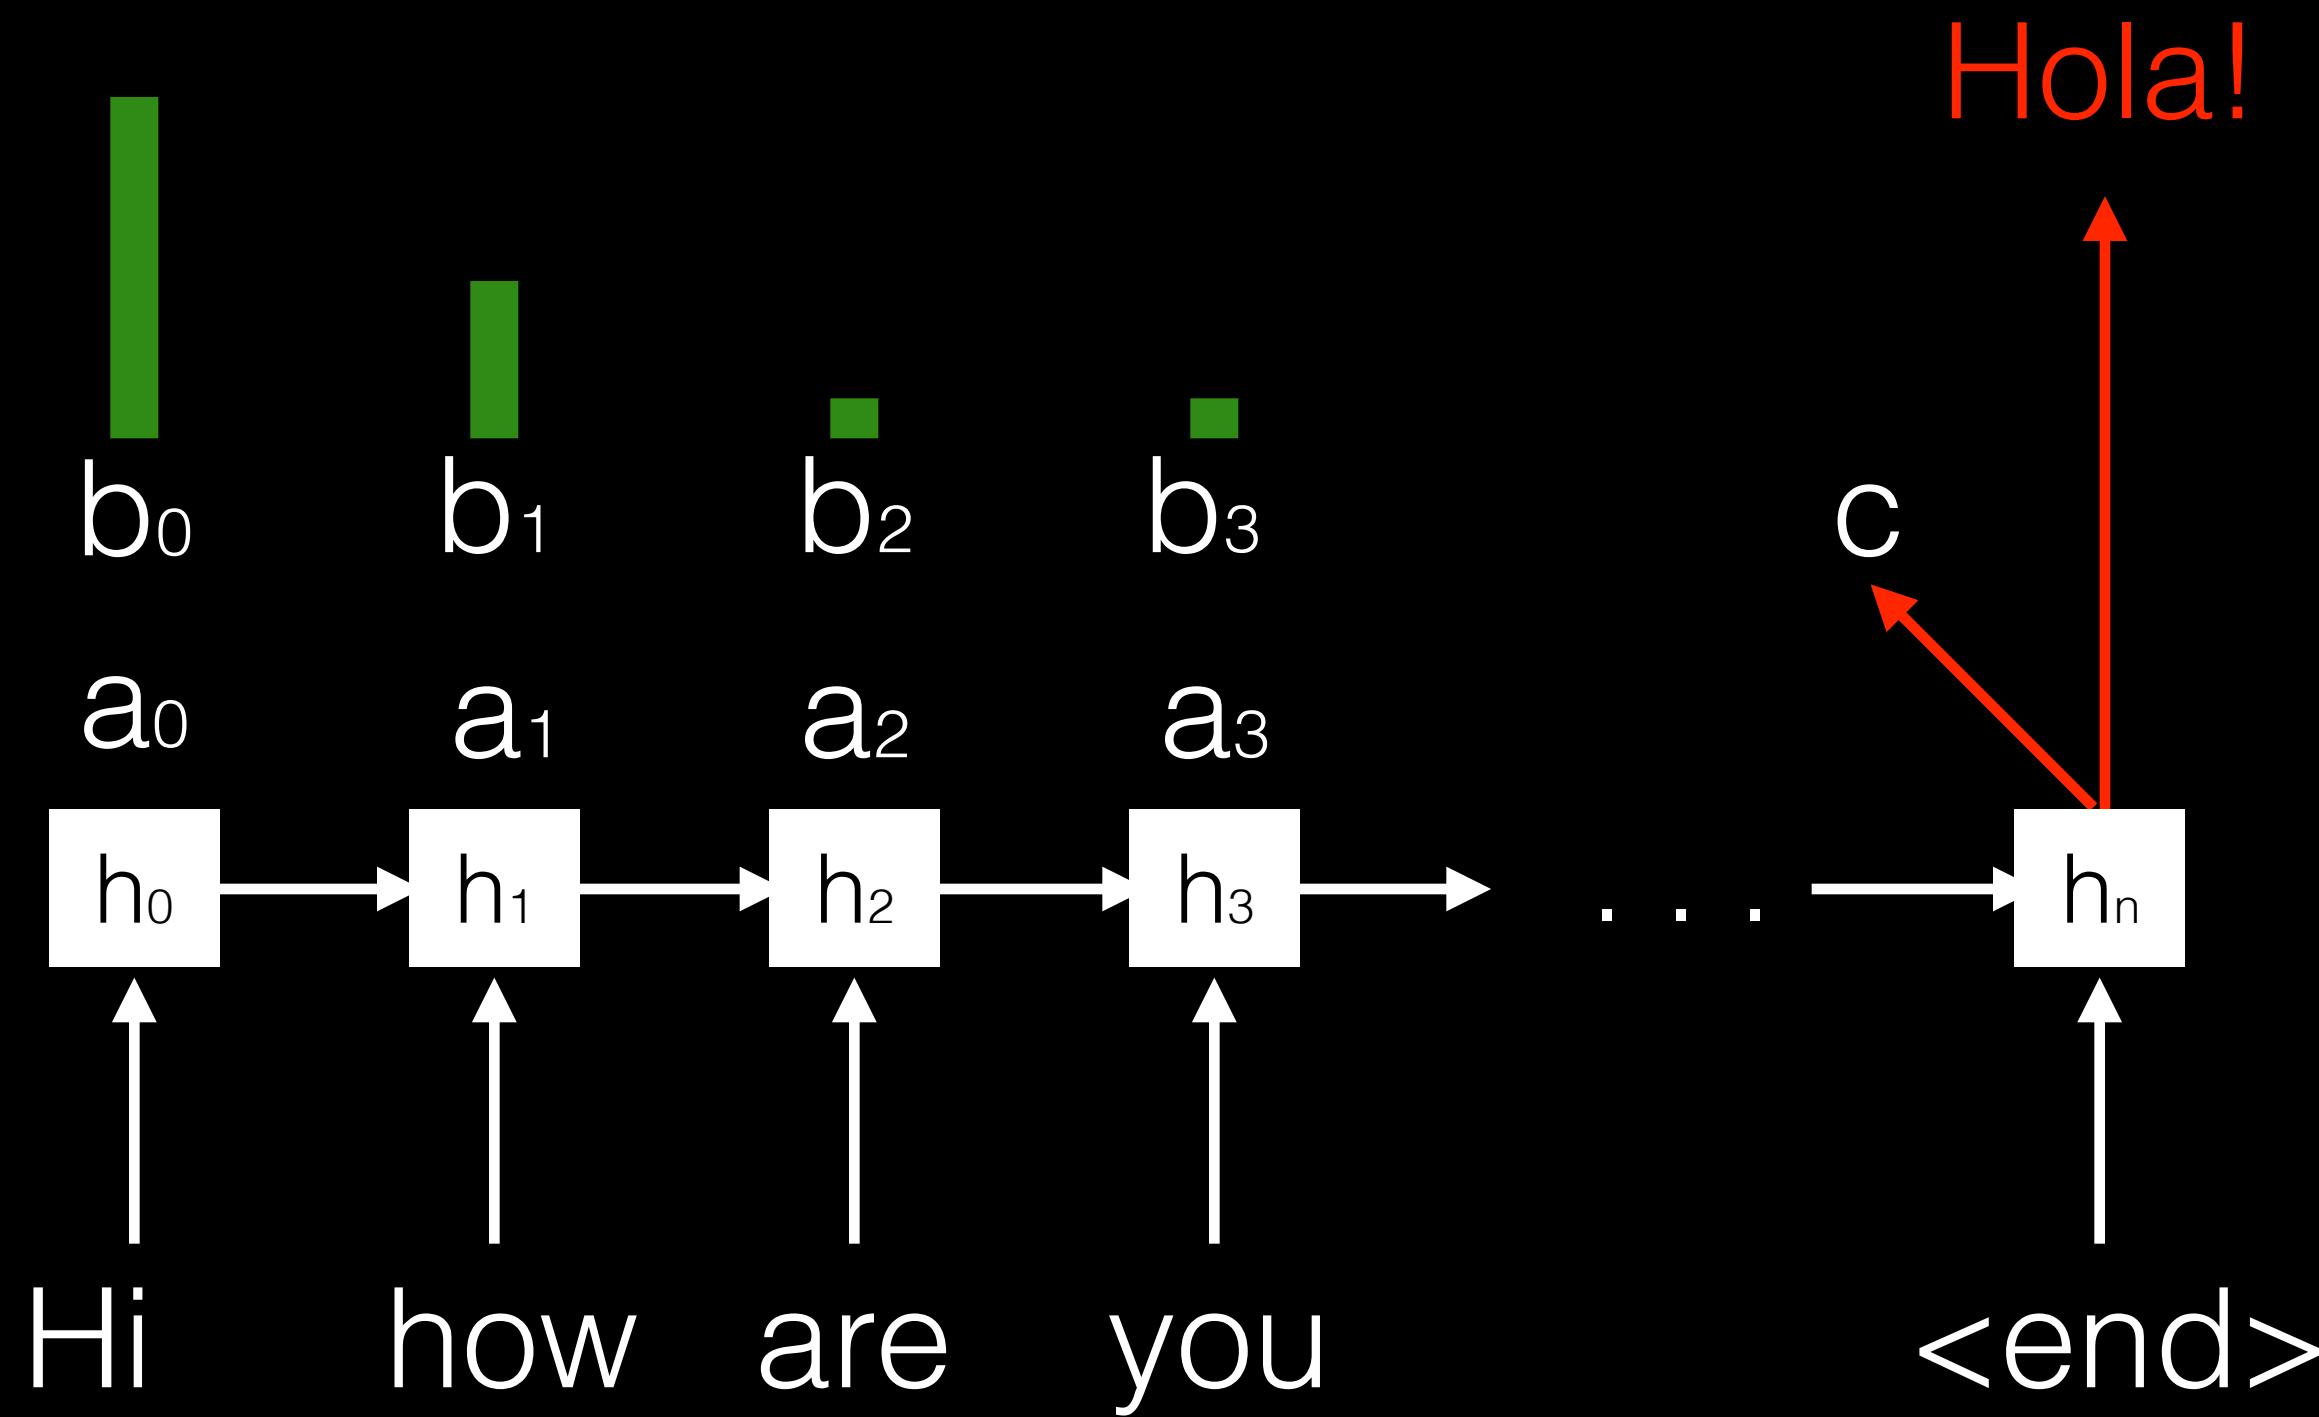


$$b_i = \frac{\exp(a_i)}{\exp(a_1) + \exp(a_2) + \dots + \exp(a_n)}$$

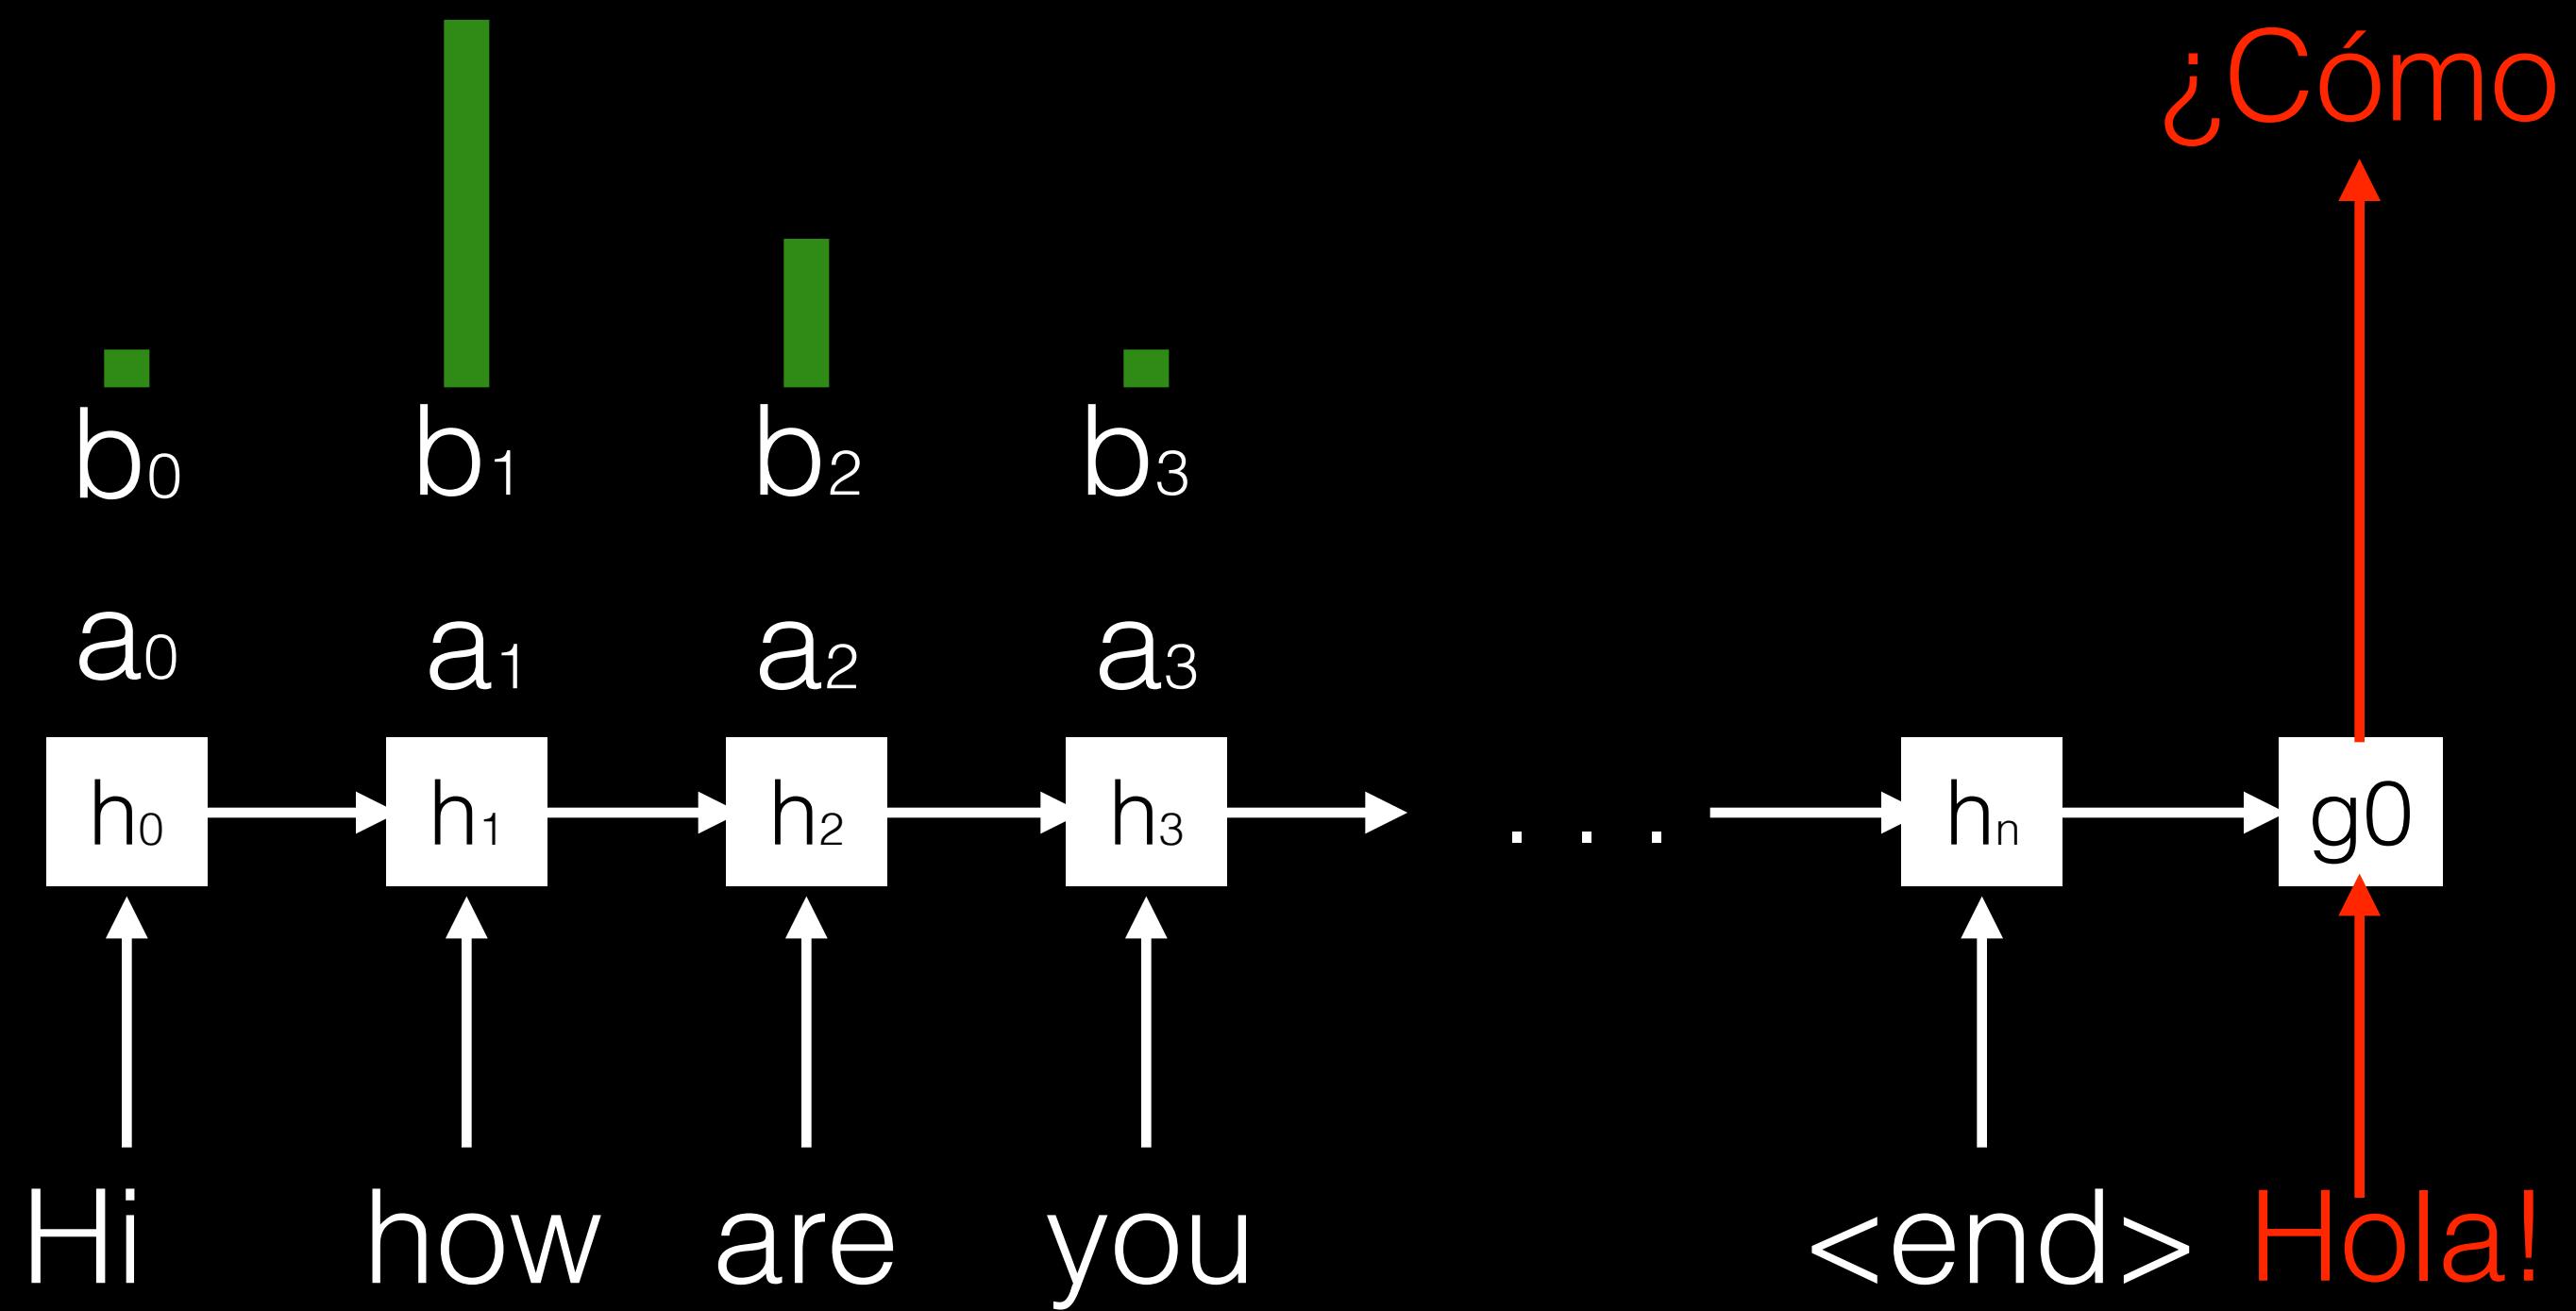
Attention Mechanism



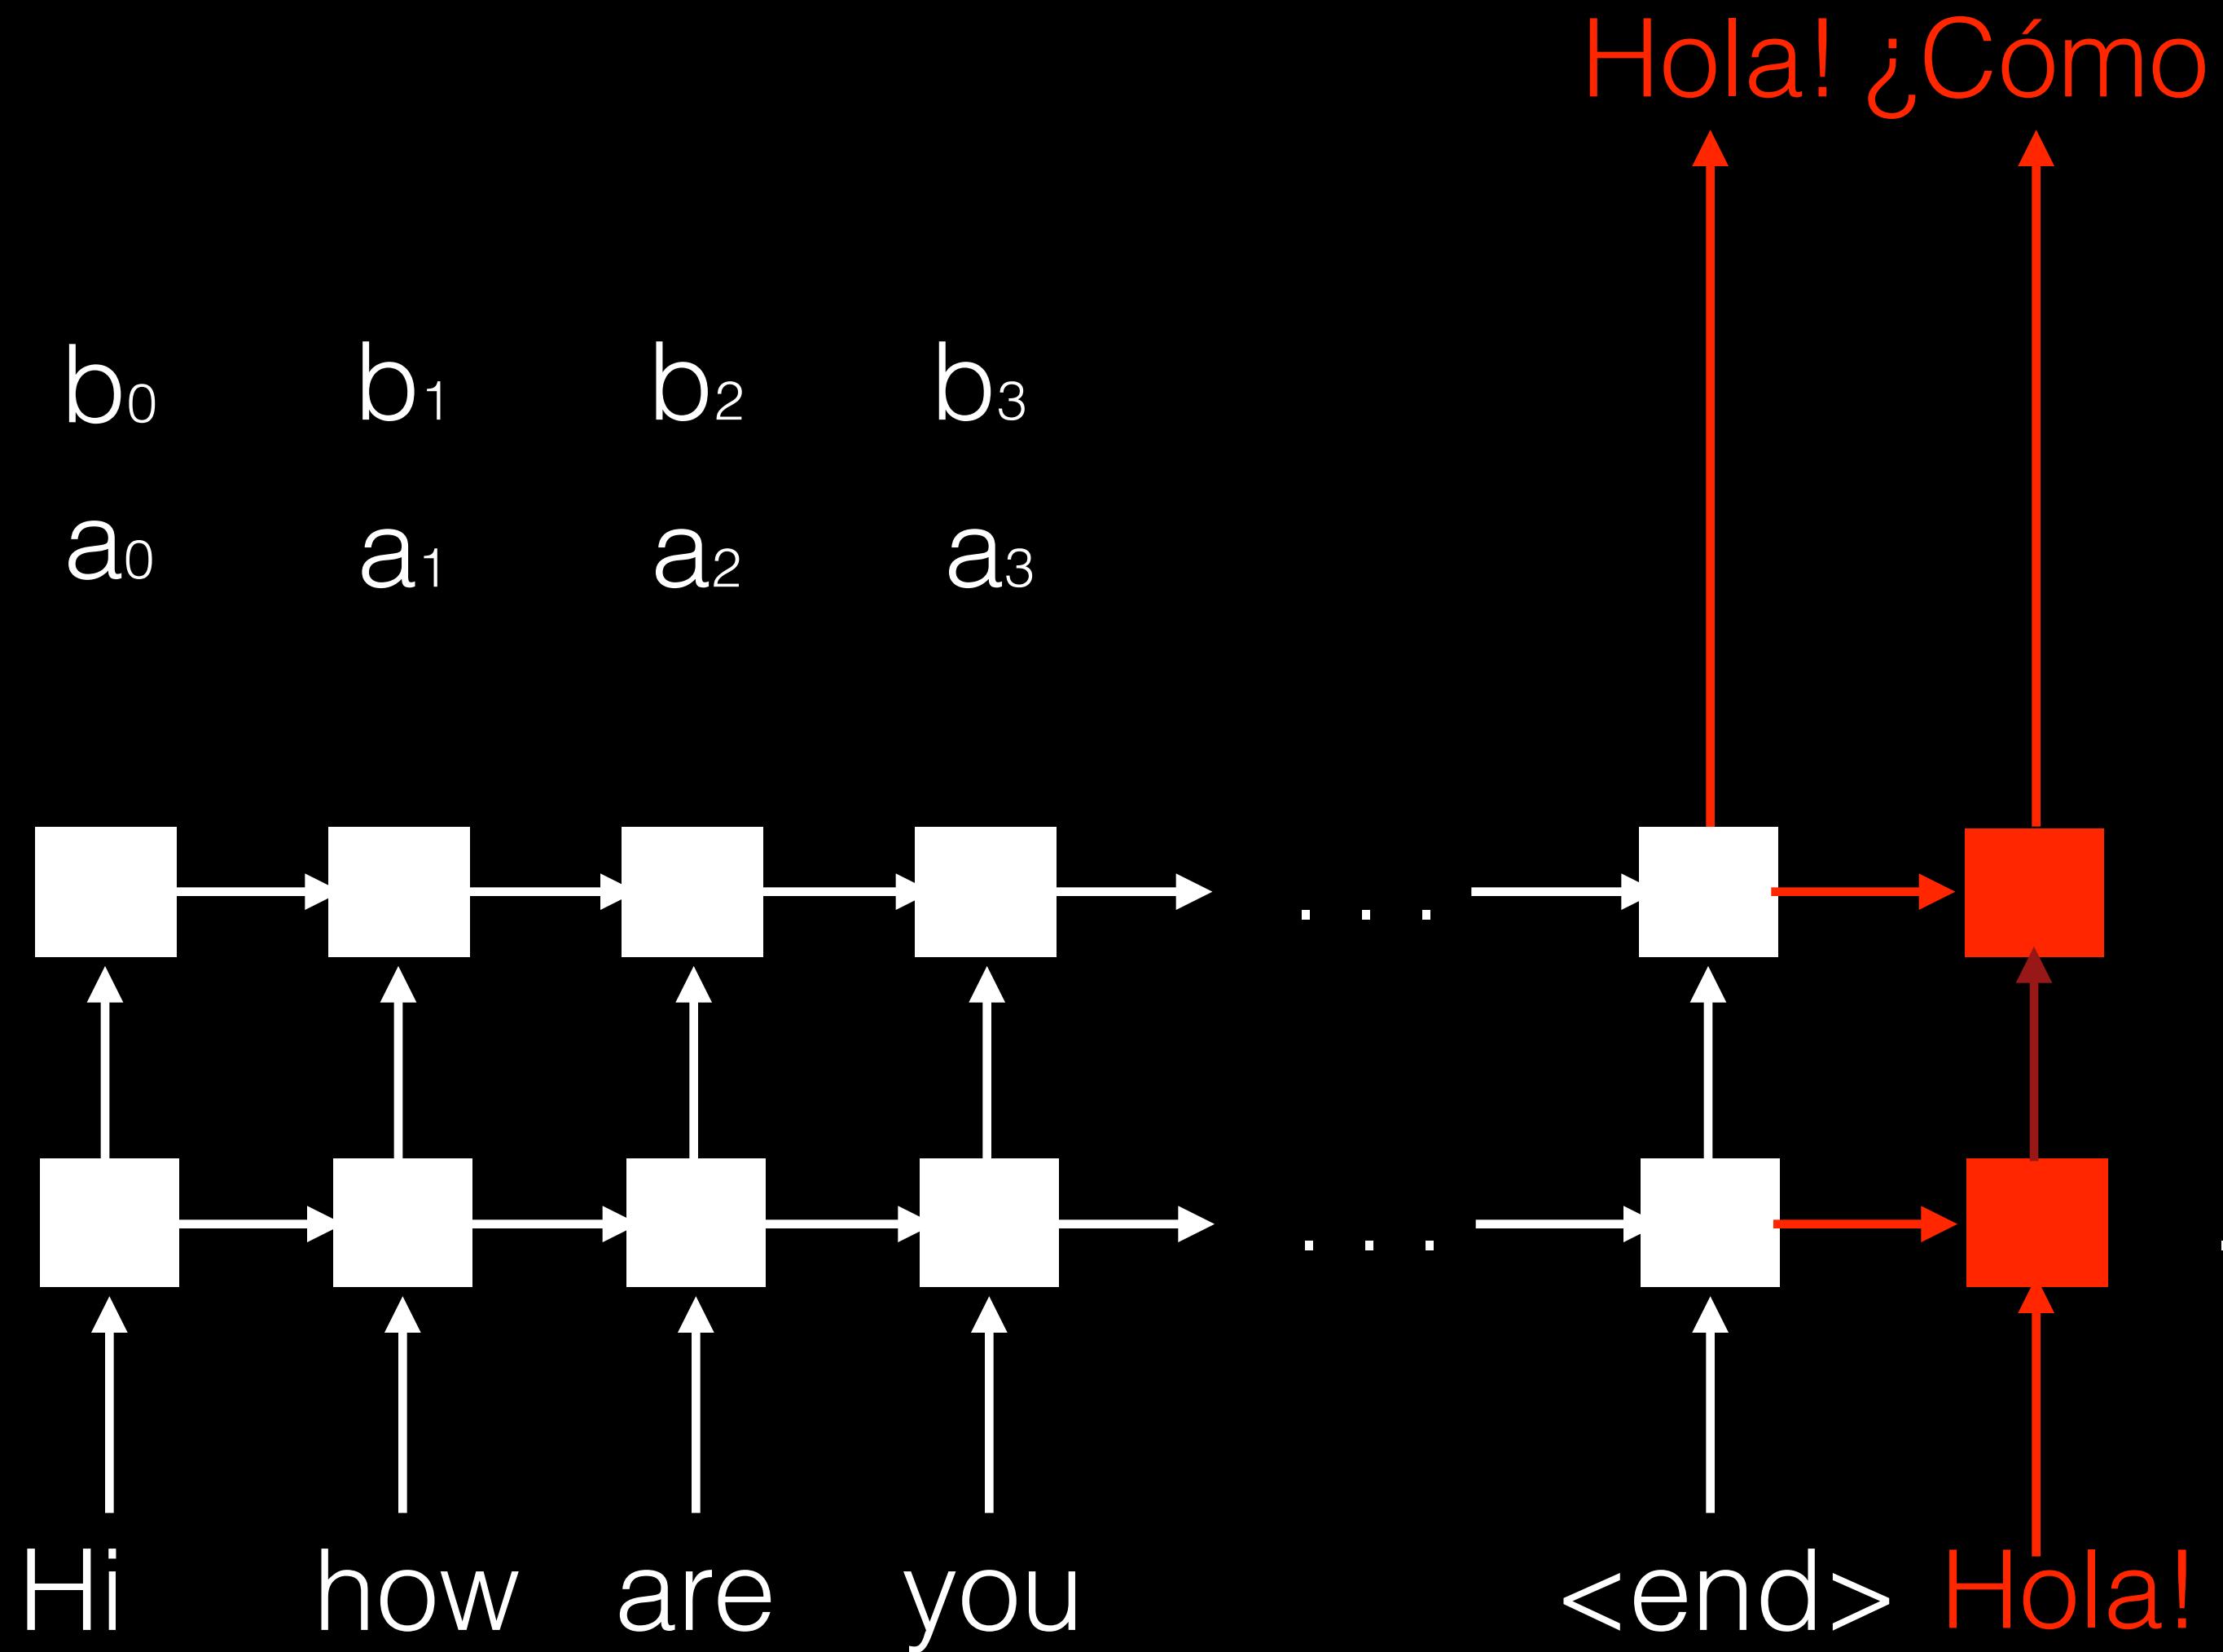
Model Understandability with Attention Mechanism



Model Understandability with Attention Mechanism



Deeper Networks work Better



Sequence to Sequence With Attention

- Currently the state-of-art in many translation tasks
 - Tip 1: Use word segments or word/character hybrid instead of just words
 - Tip 2: Gradient Clipping to prevent explosion
 - Tip 3: Use Long Short Term Memory

LSTMCell vs. RNNCell

RNNCell:

```
h = tanh(theta * [inputs, h])
```

LSTMCell:

```
Z = theta * [inputs, h]
```

```
i, j, f, o = split(1, 4, Z) # split to four blocks
```

```
new_c = c * sigmoid(f) + sigmoid(i) * tanh(j) # integral of c
```

```
new_h = tanh(new_c) * sigmoid(o)
```

Applications

- Other applications:
 - Summarization, Image Captioning,
 - Speech Transcription, Q&A

Applications

- Other applications:
 - Summarization, Image Captioning,
 - Speech Transcription, Q&A

seq2seq for Speech

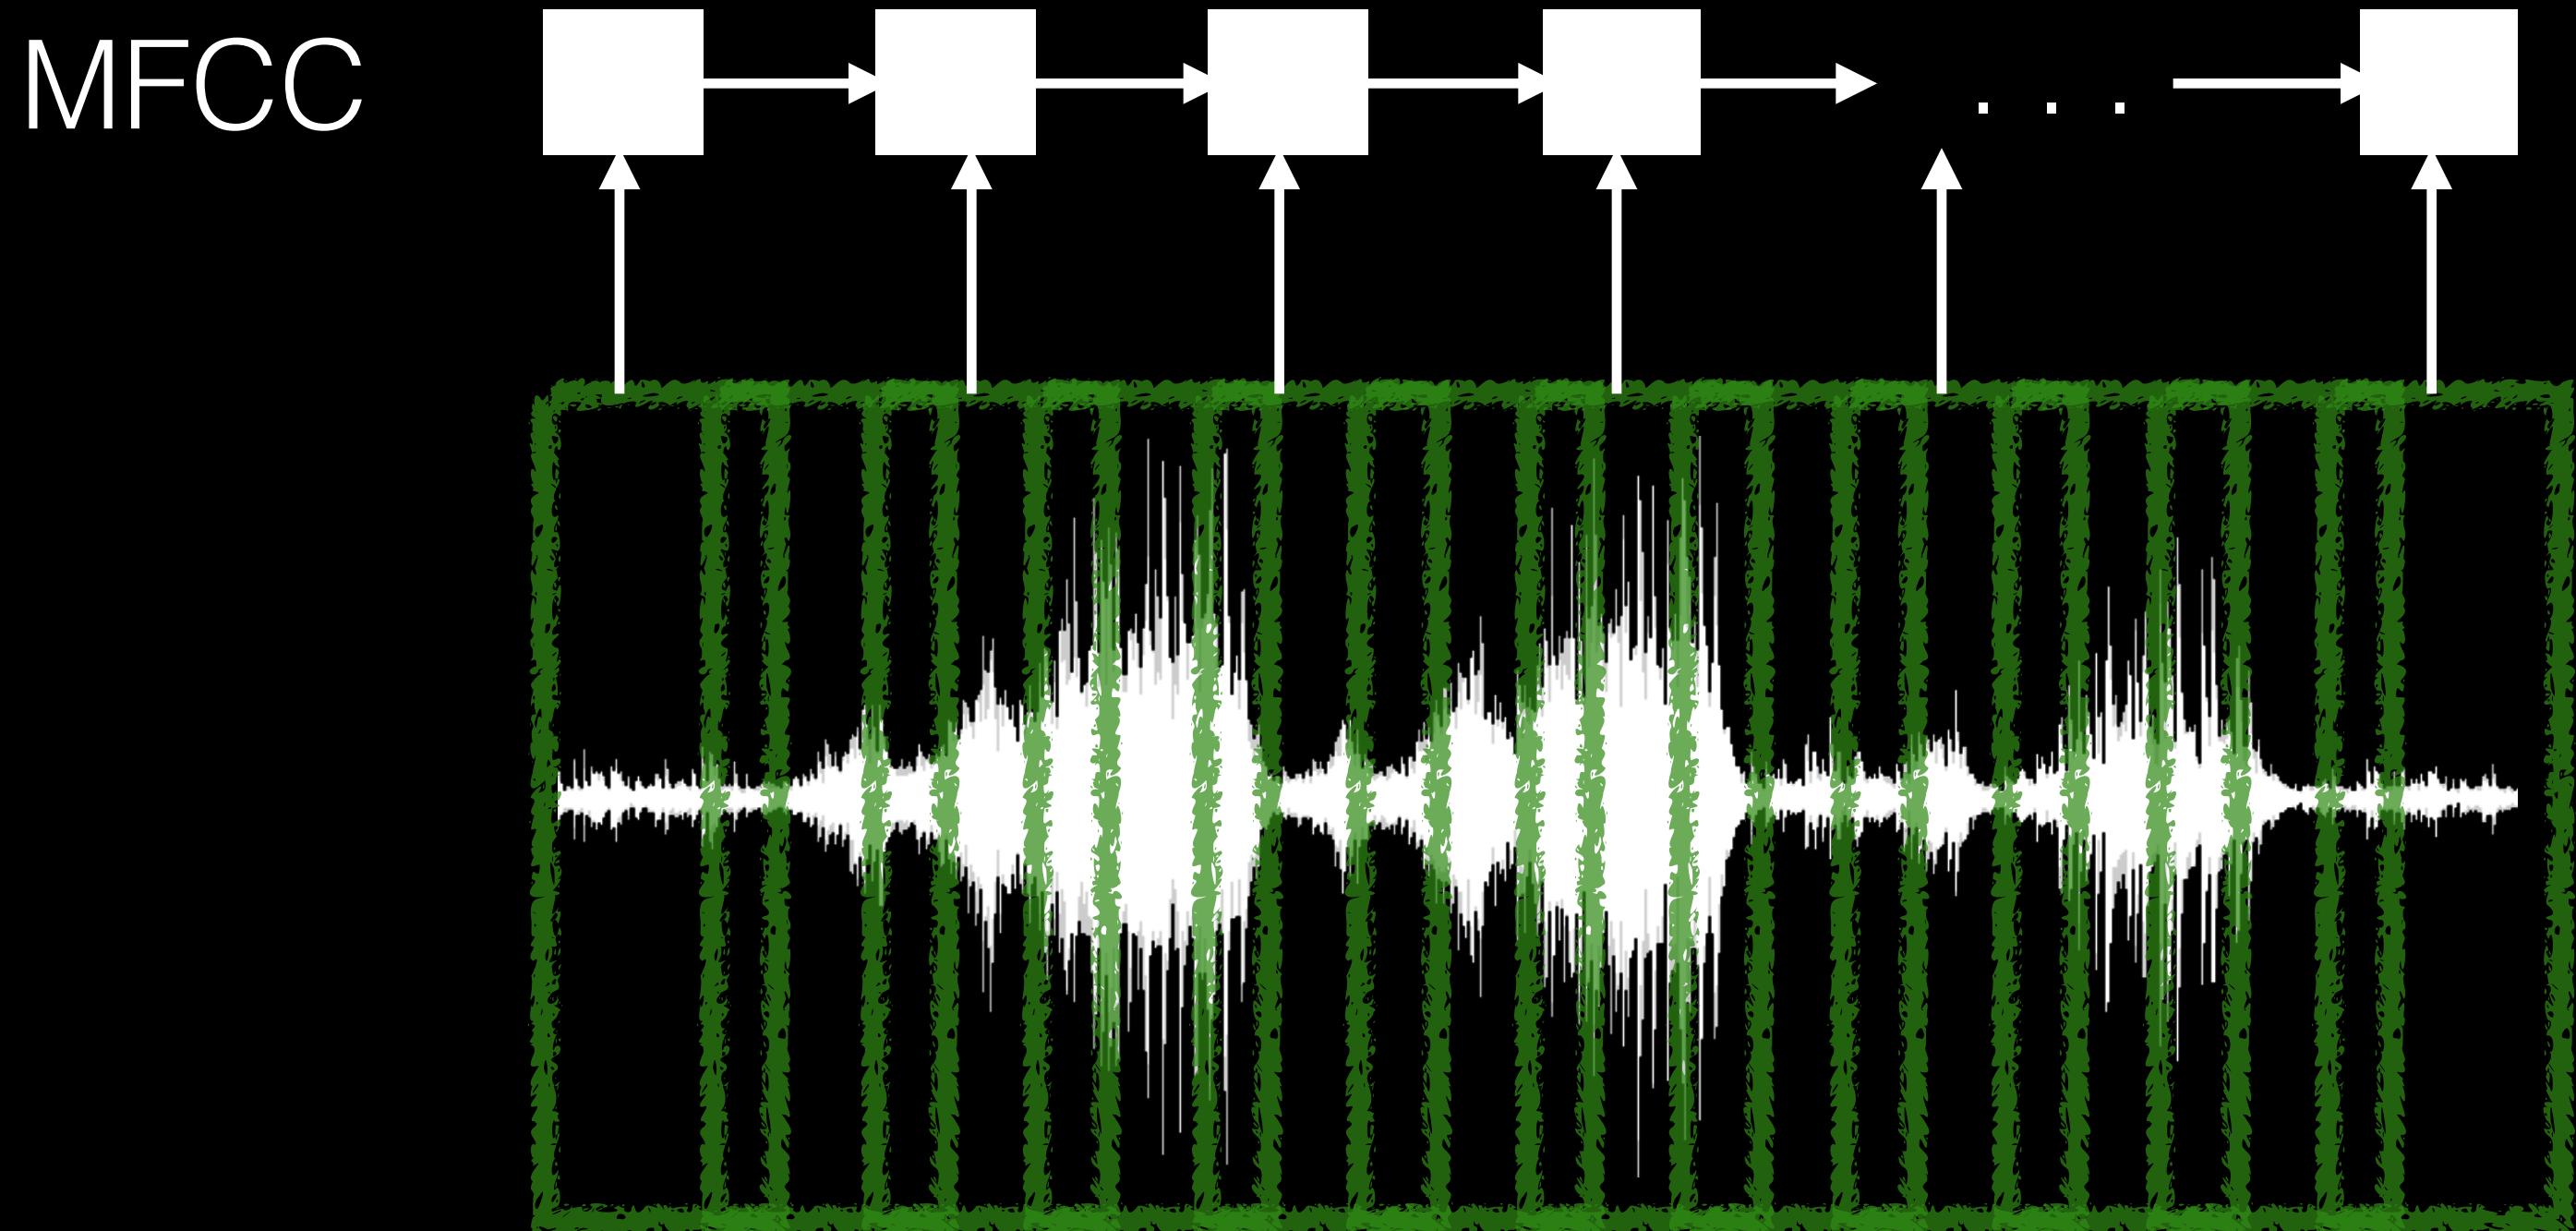


→ Hi how's it?

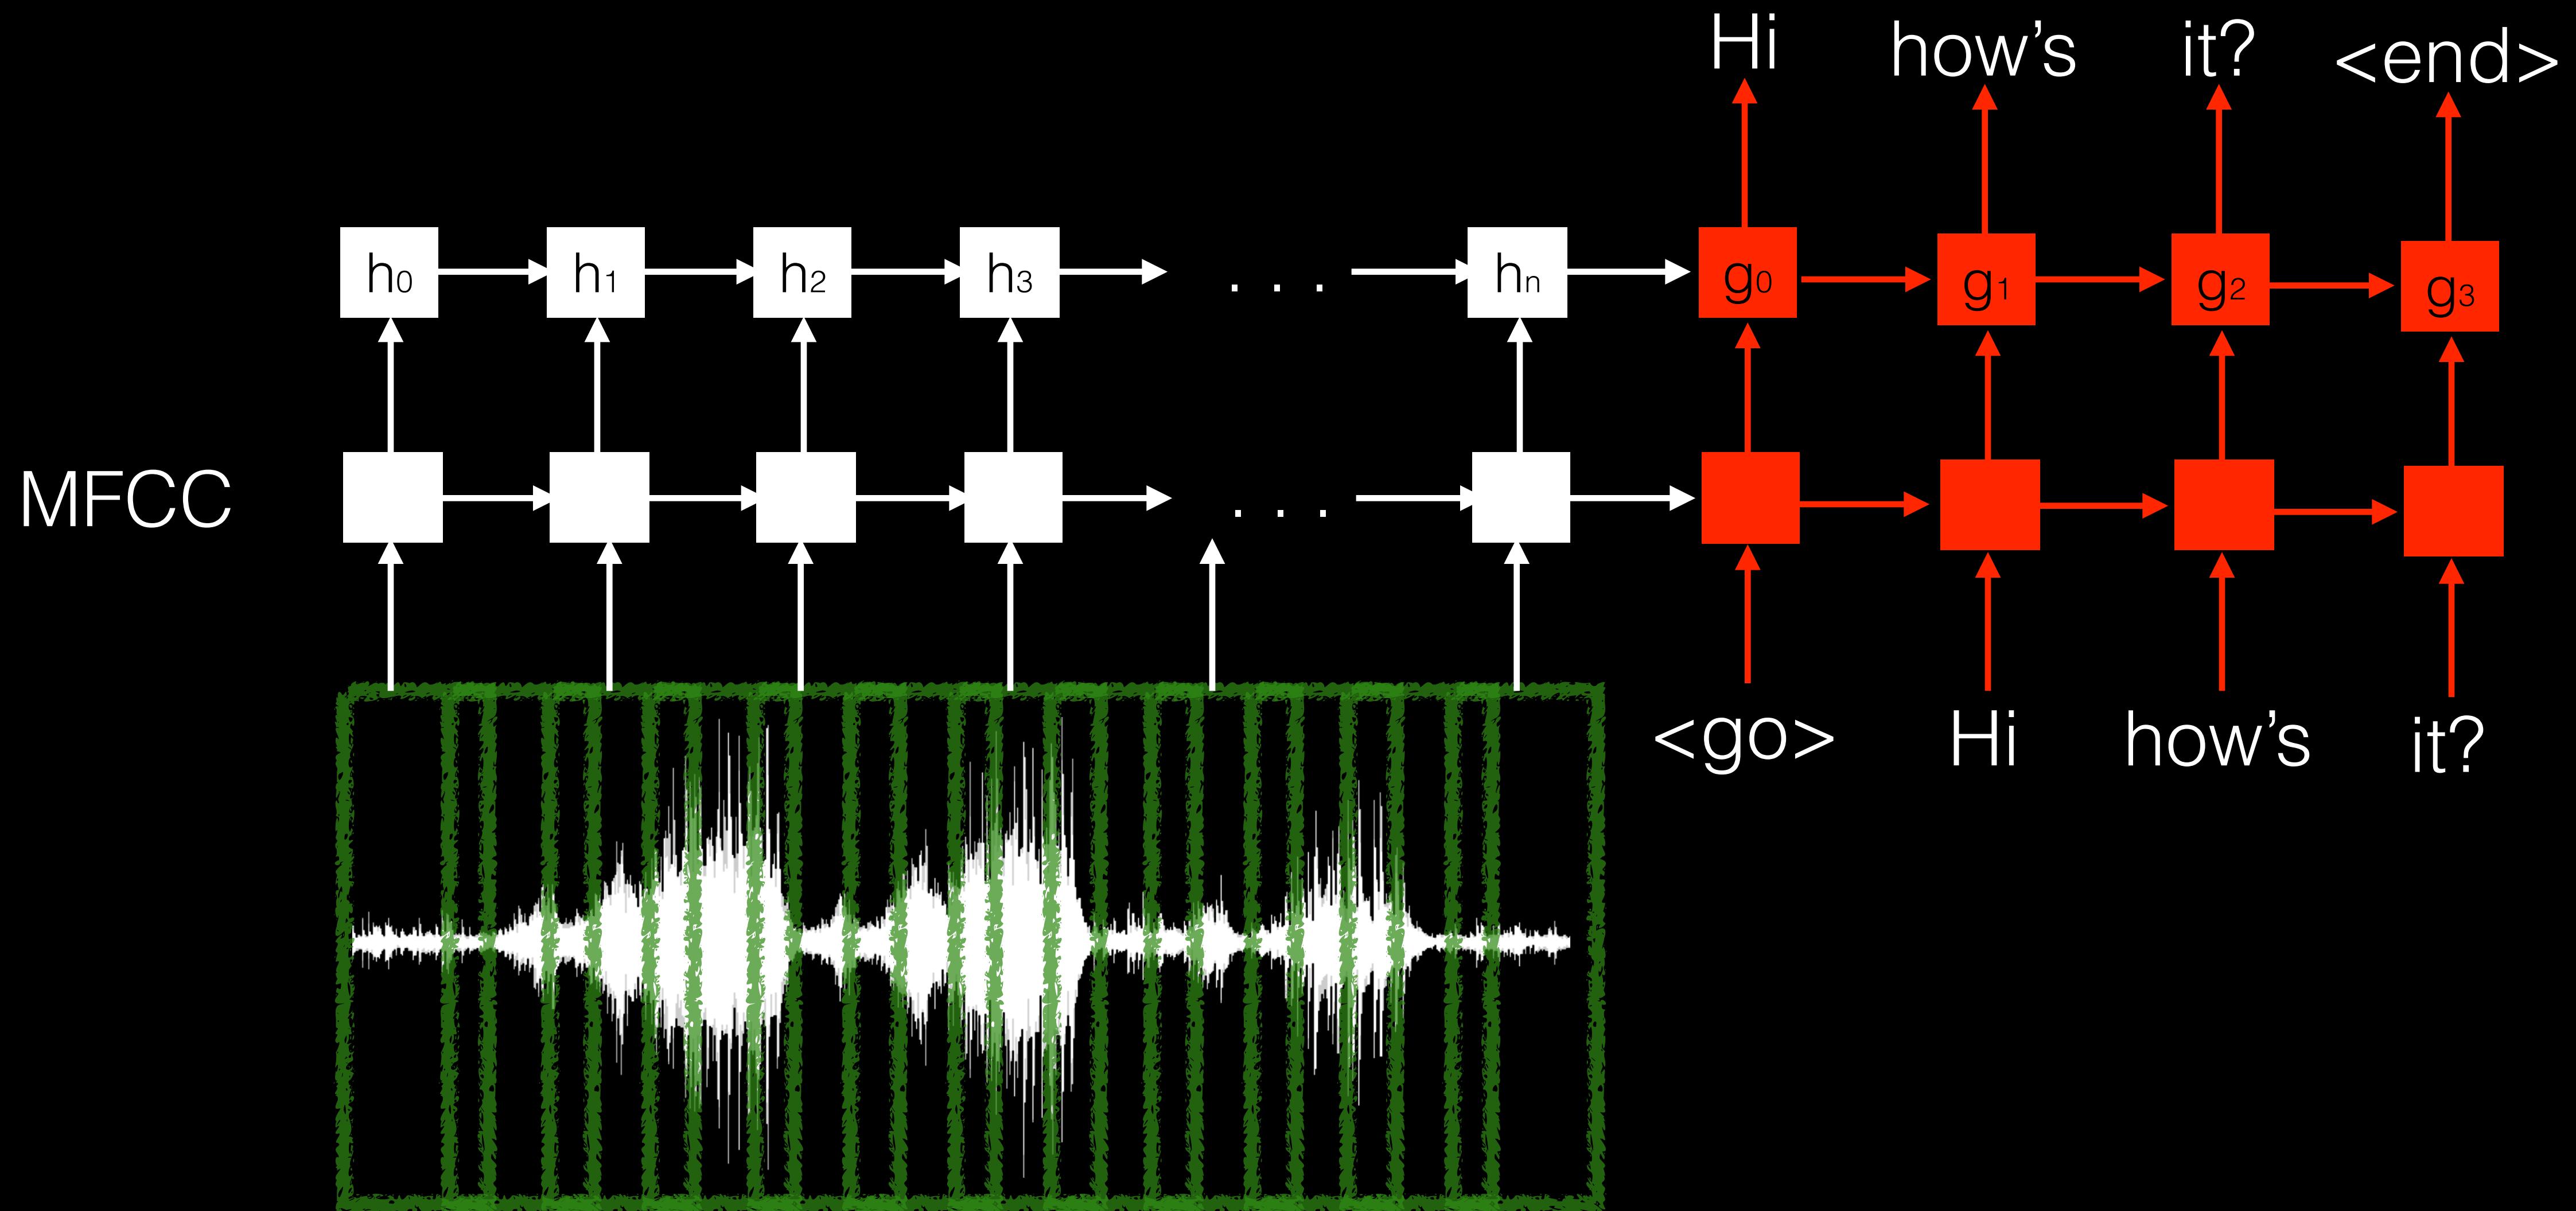
seq2seq for Speech



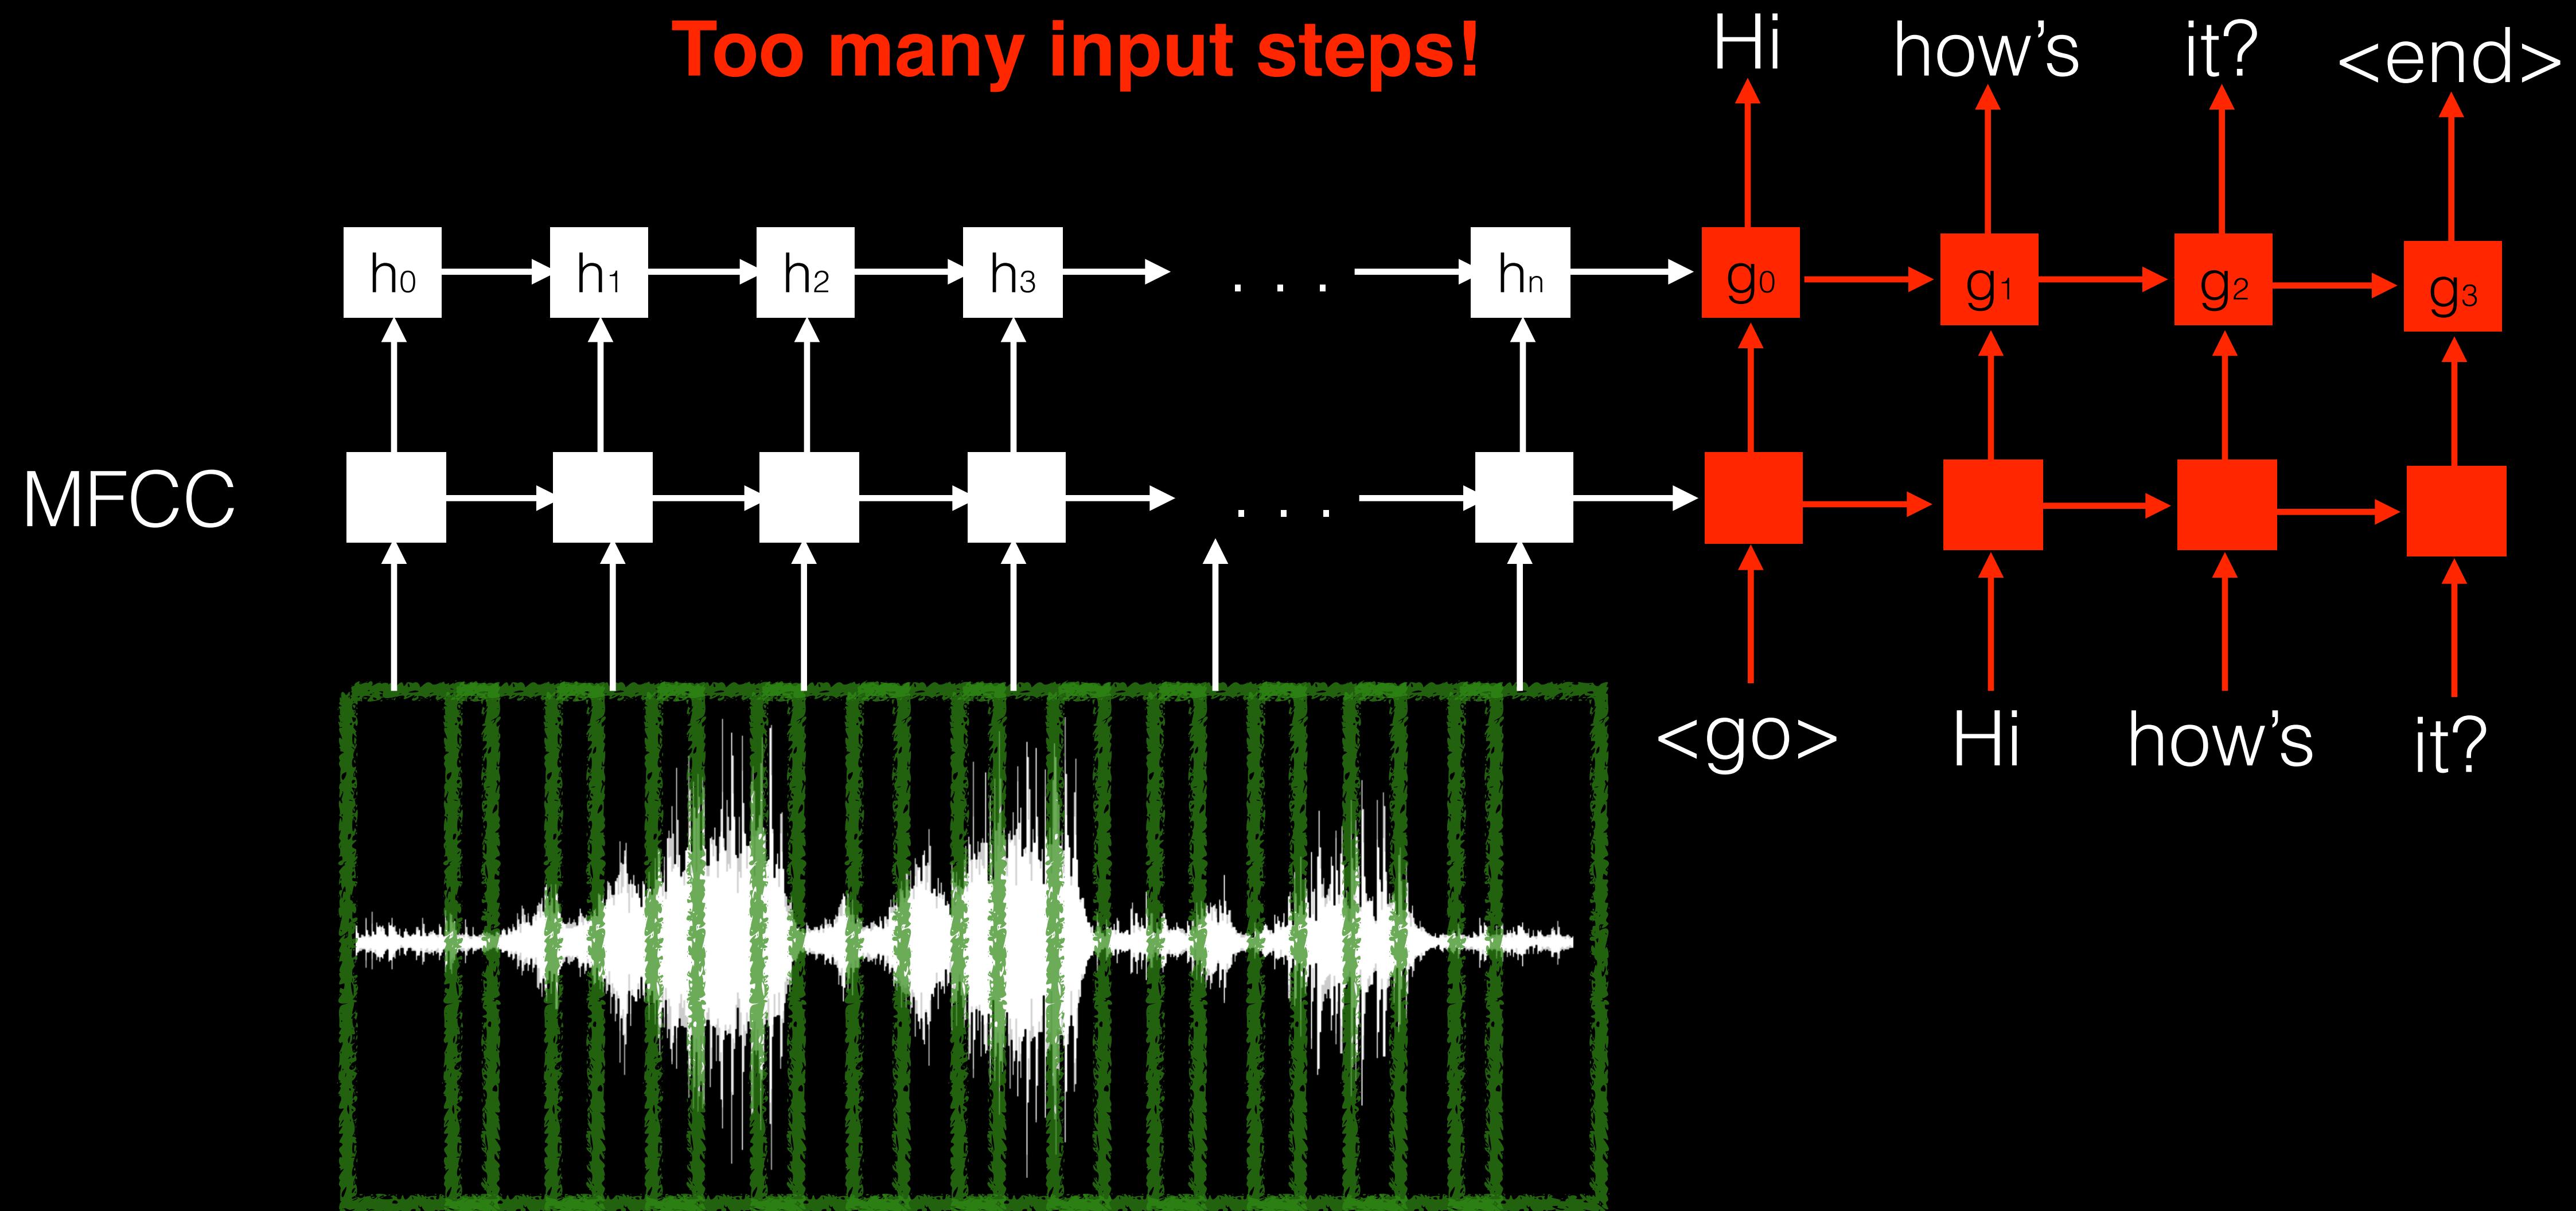
seq2seq for Speech



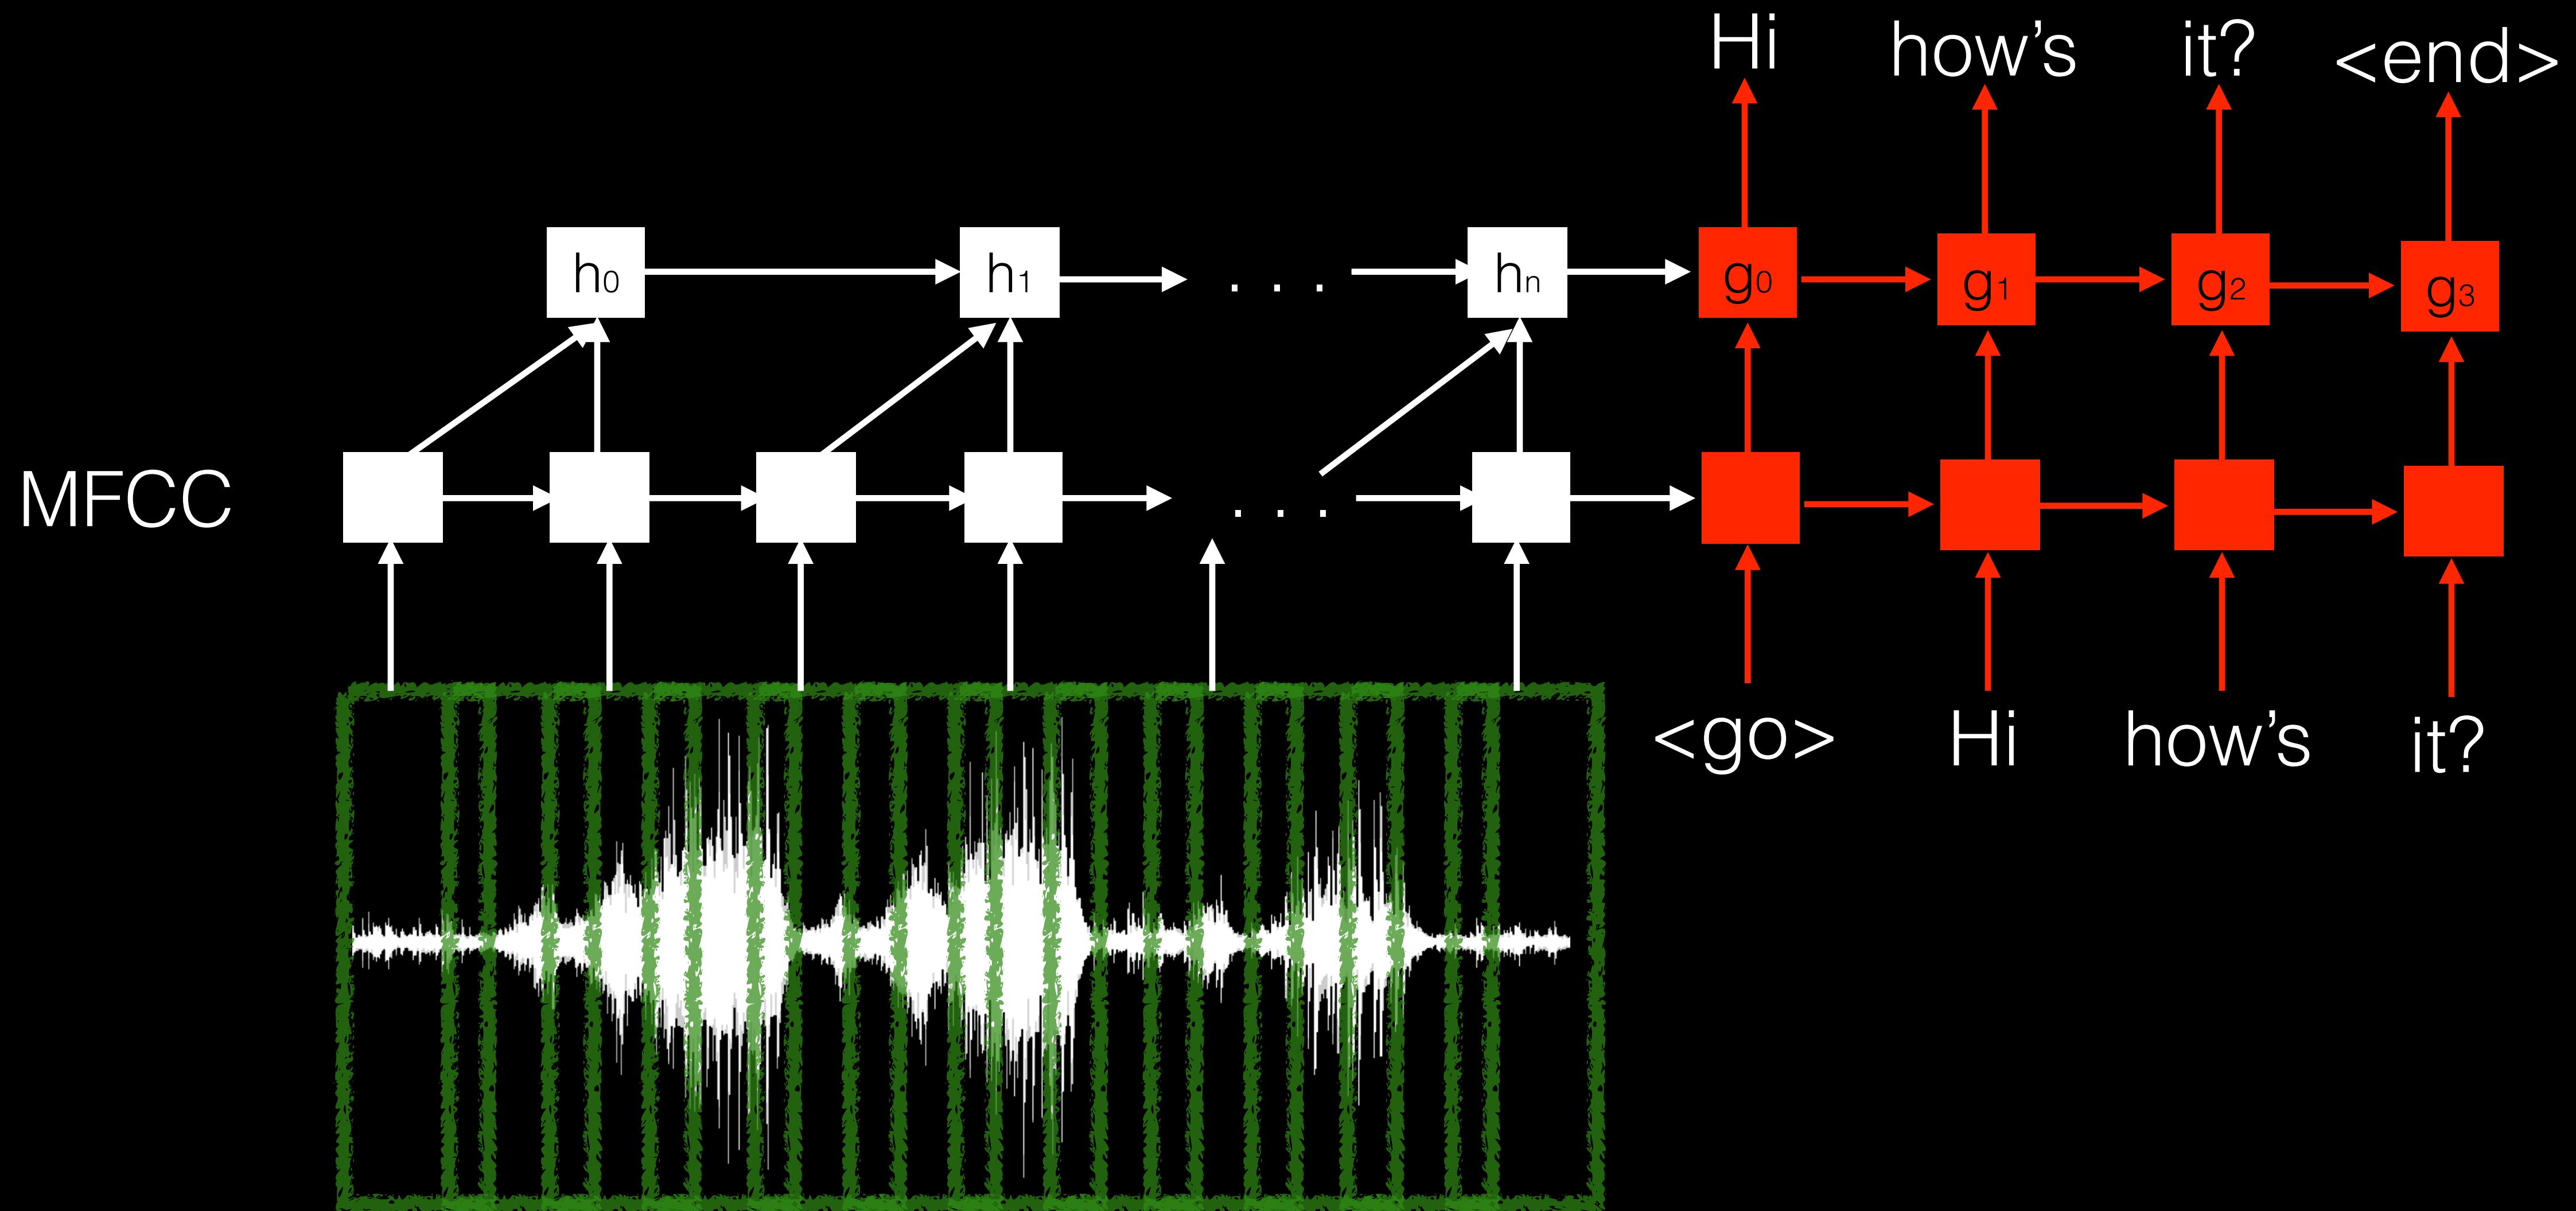
seq2seq for Speech



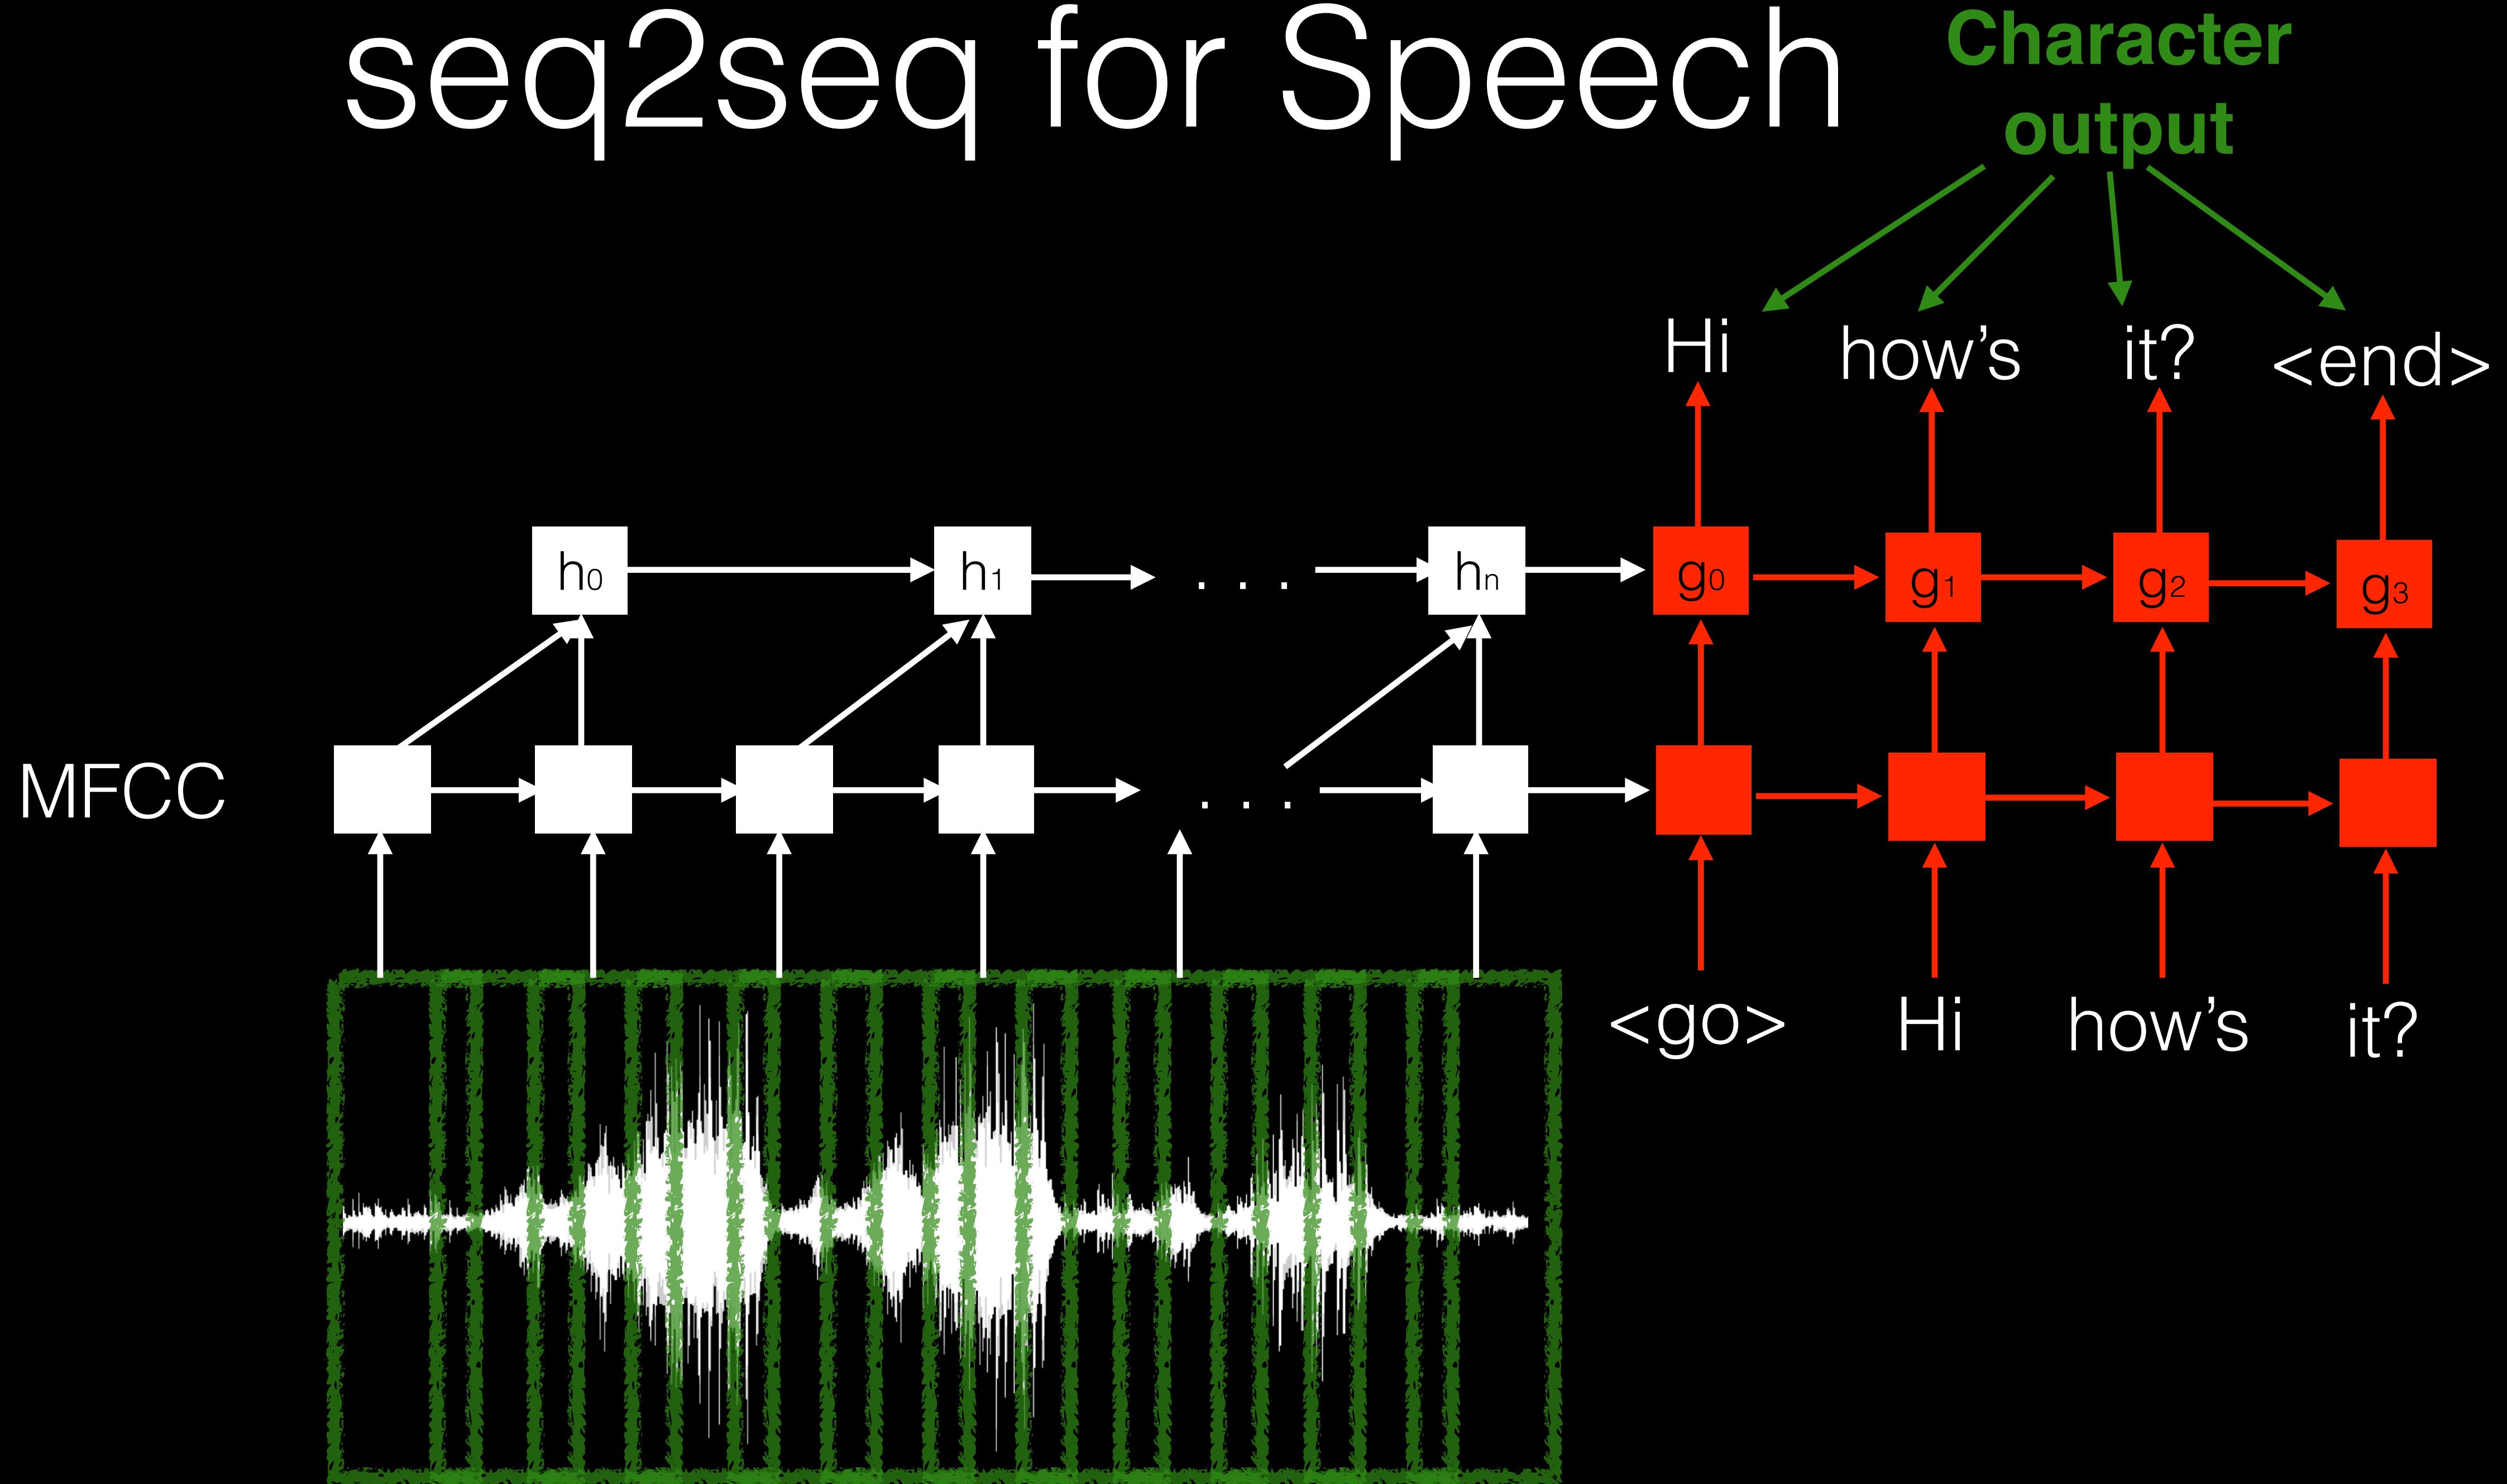
seq2seq for Speech



seq2seq for Speech



seq2seq for Speech



Sequence to Sequence With Attention for Speech

- Implicit language model
- “Offline” beam search decoding
- Not as good as
 - CTC (Adam Coates’ talk)
 - HMM-DNN hybrid (most widely-used speech systems)

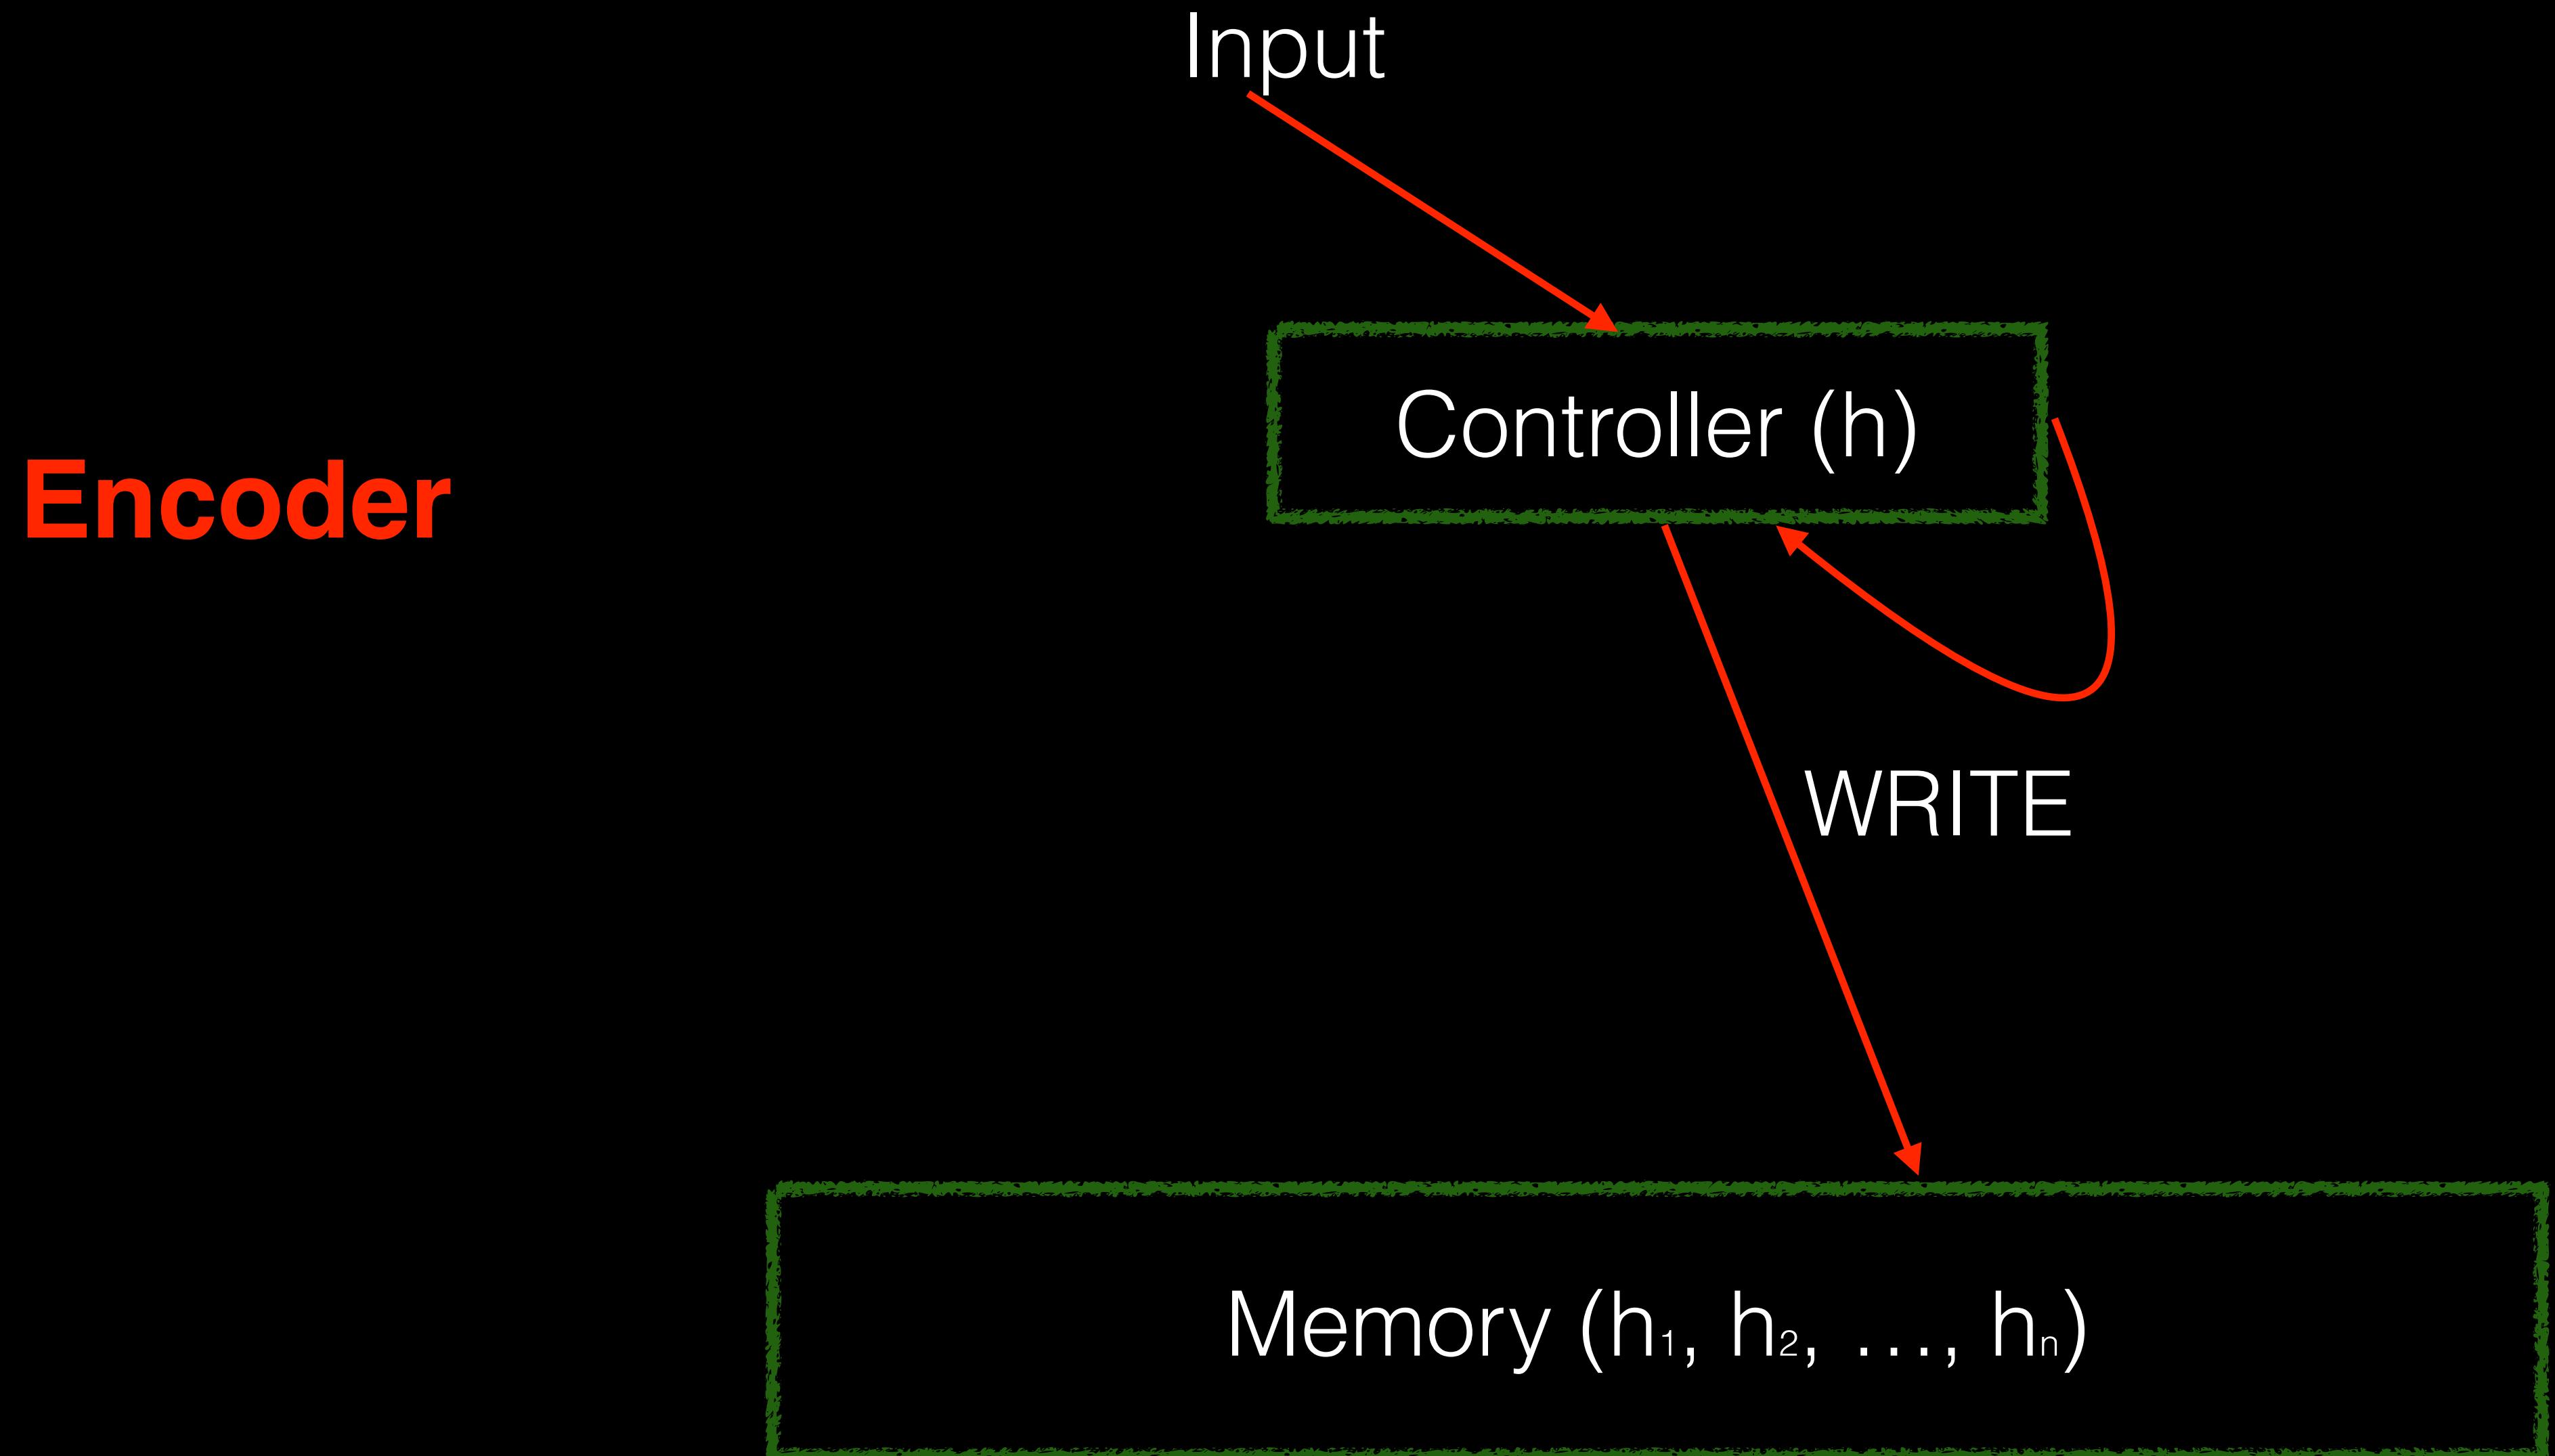
The Big Picture

- Sequence to sequence is an “End-to-end Deep Learning” algorithm
- It’s very general, so it should work with most NLP-related tasks **when you have a lot of data**
- If you don’t have enough data:
 - Consider dividing your problem into smaller problems, and train seq2seq on each of them.
 - **Train jointly with many other tasks**
- What I present next is an active area of research

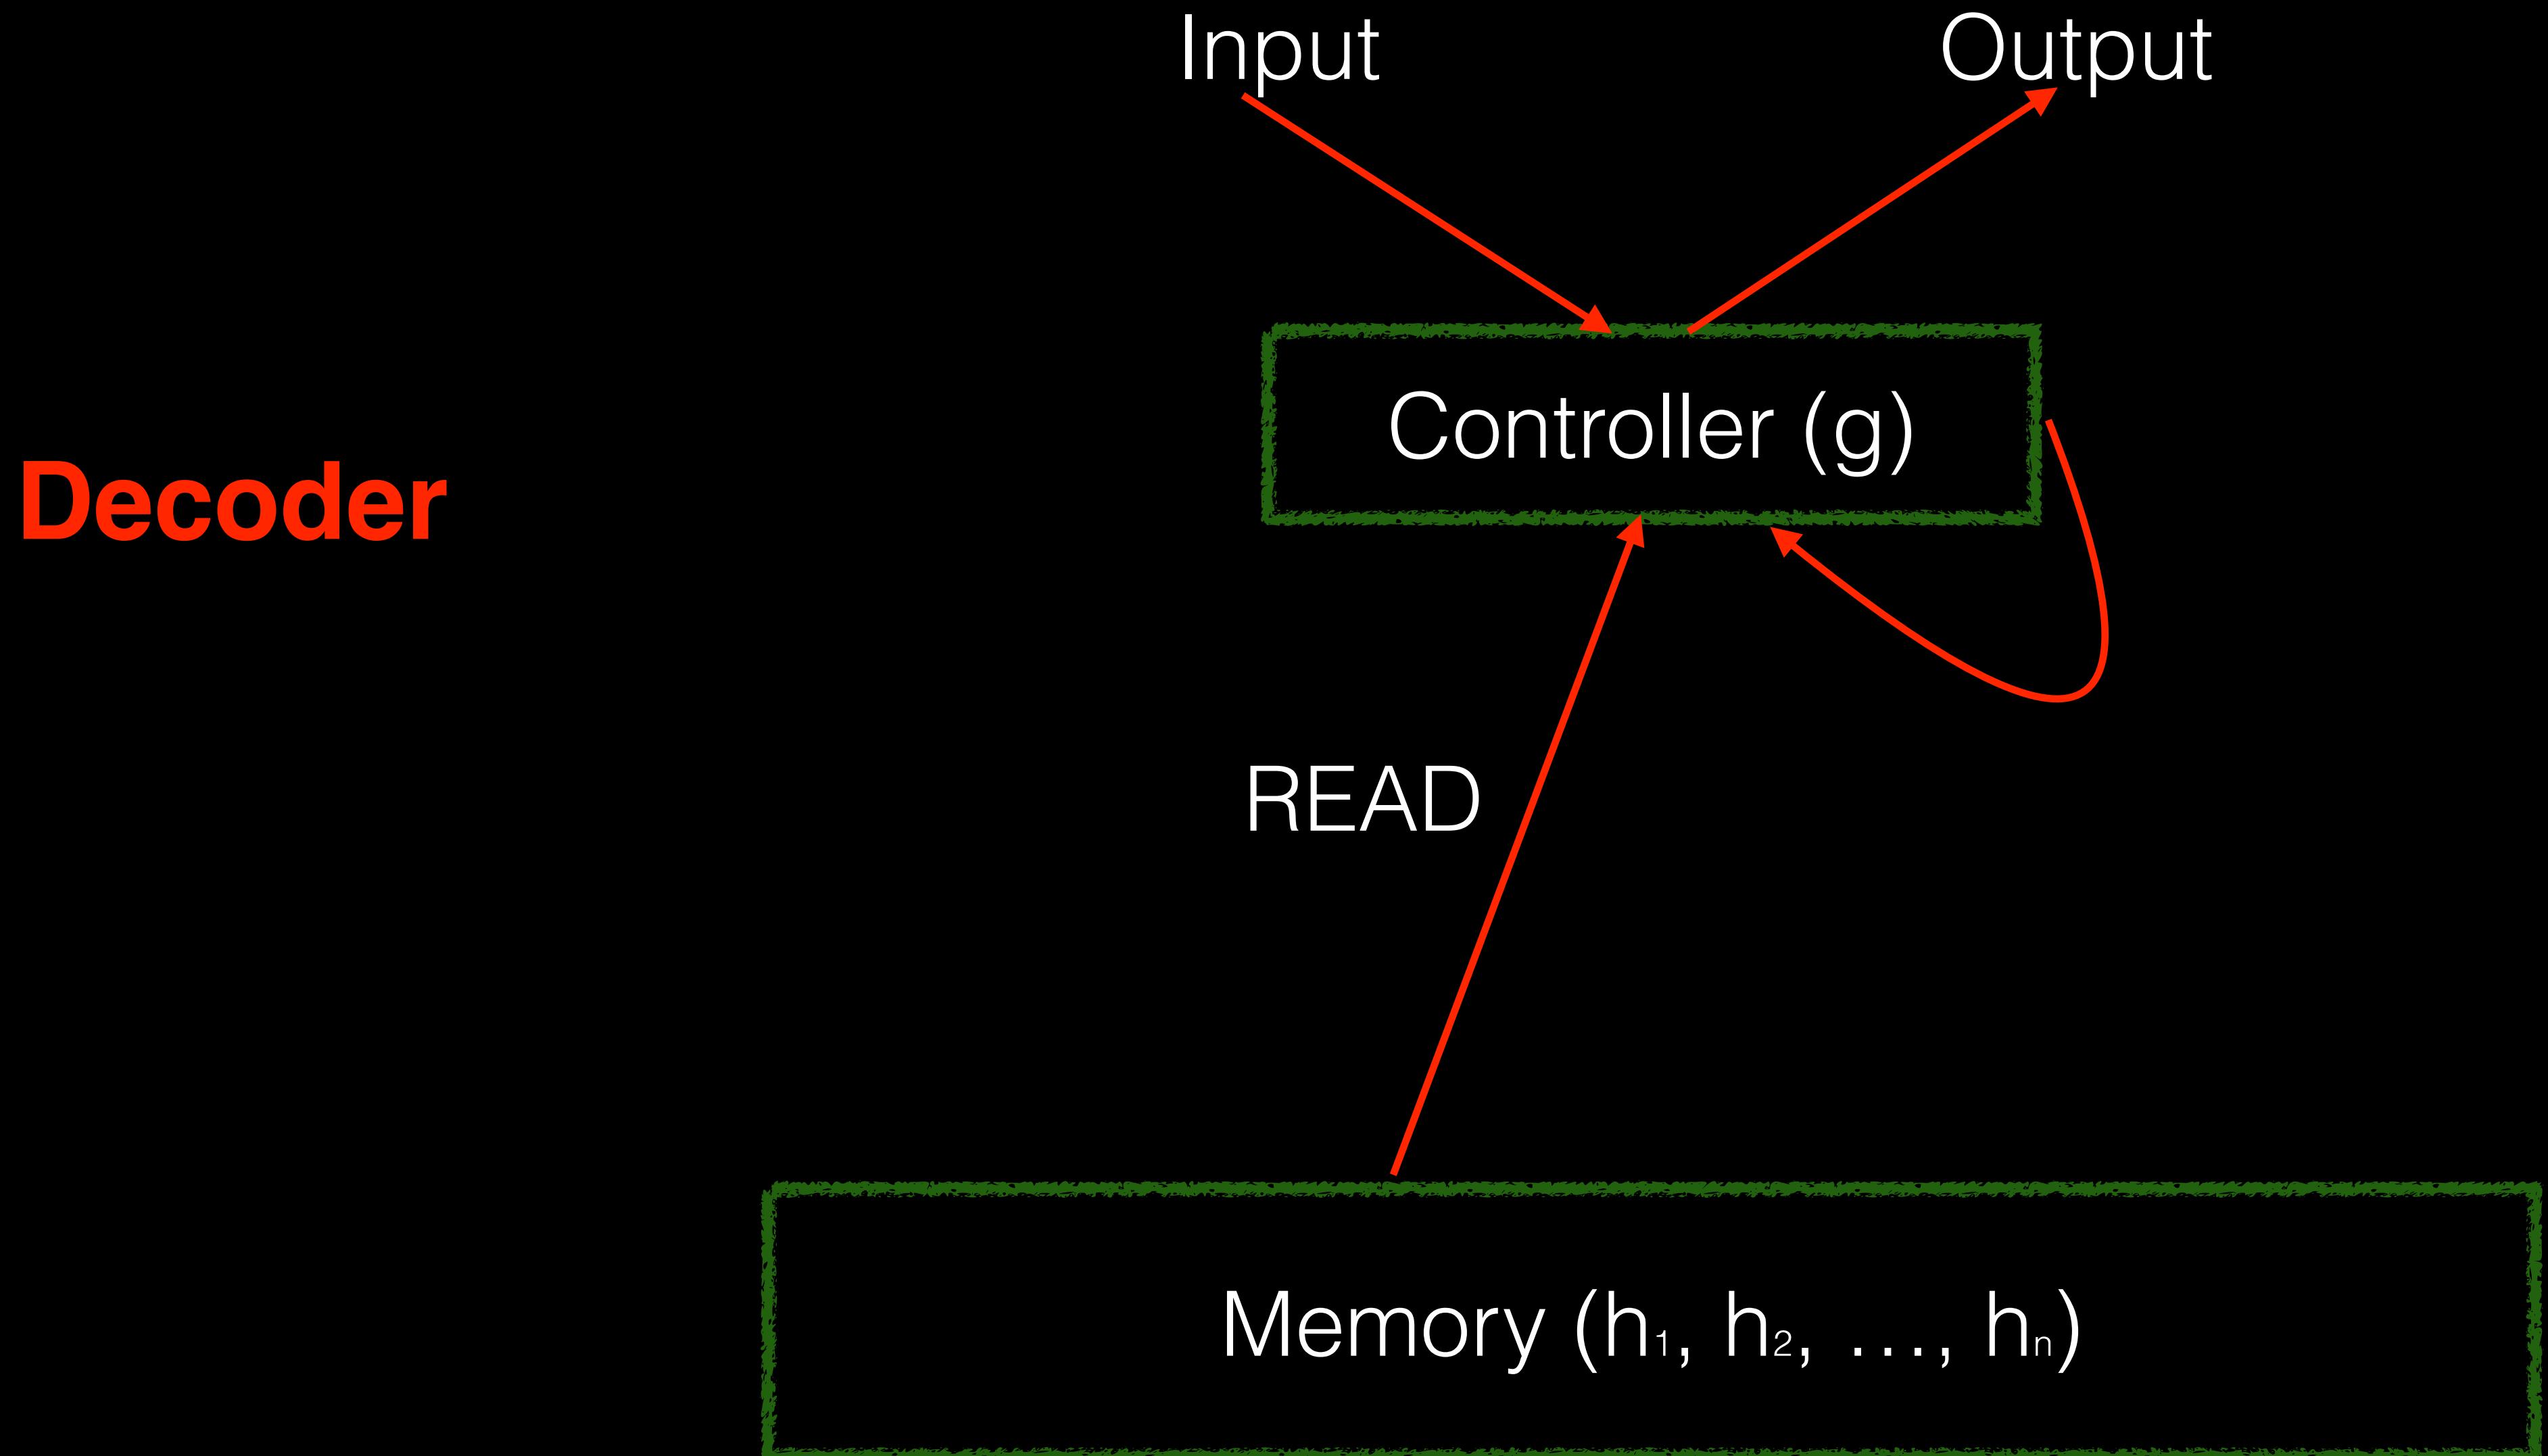
Automatic Q&A

- Reading a book and answer a question
 - Seq2seq with attention: Read the book, then read the question, then revisit all pages in the book.
- > Augmented RNNs with memory (Memory Networks, Neural Turing Machines, Dynamic Memory Networks, Stack-augmented RNNs etc.)

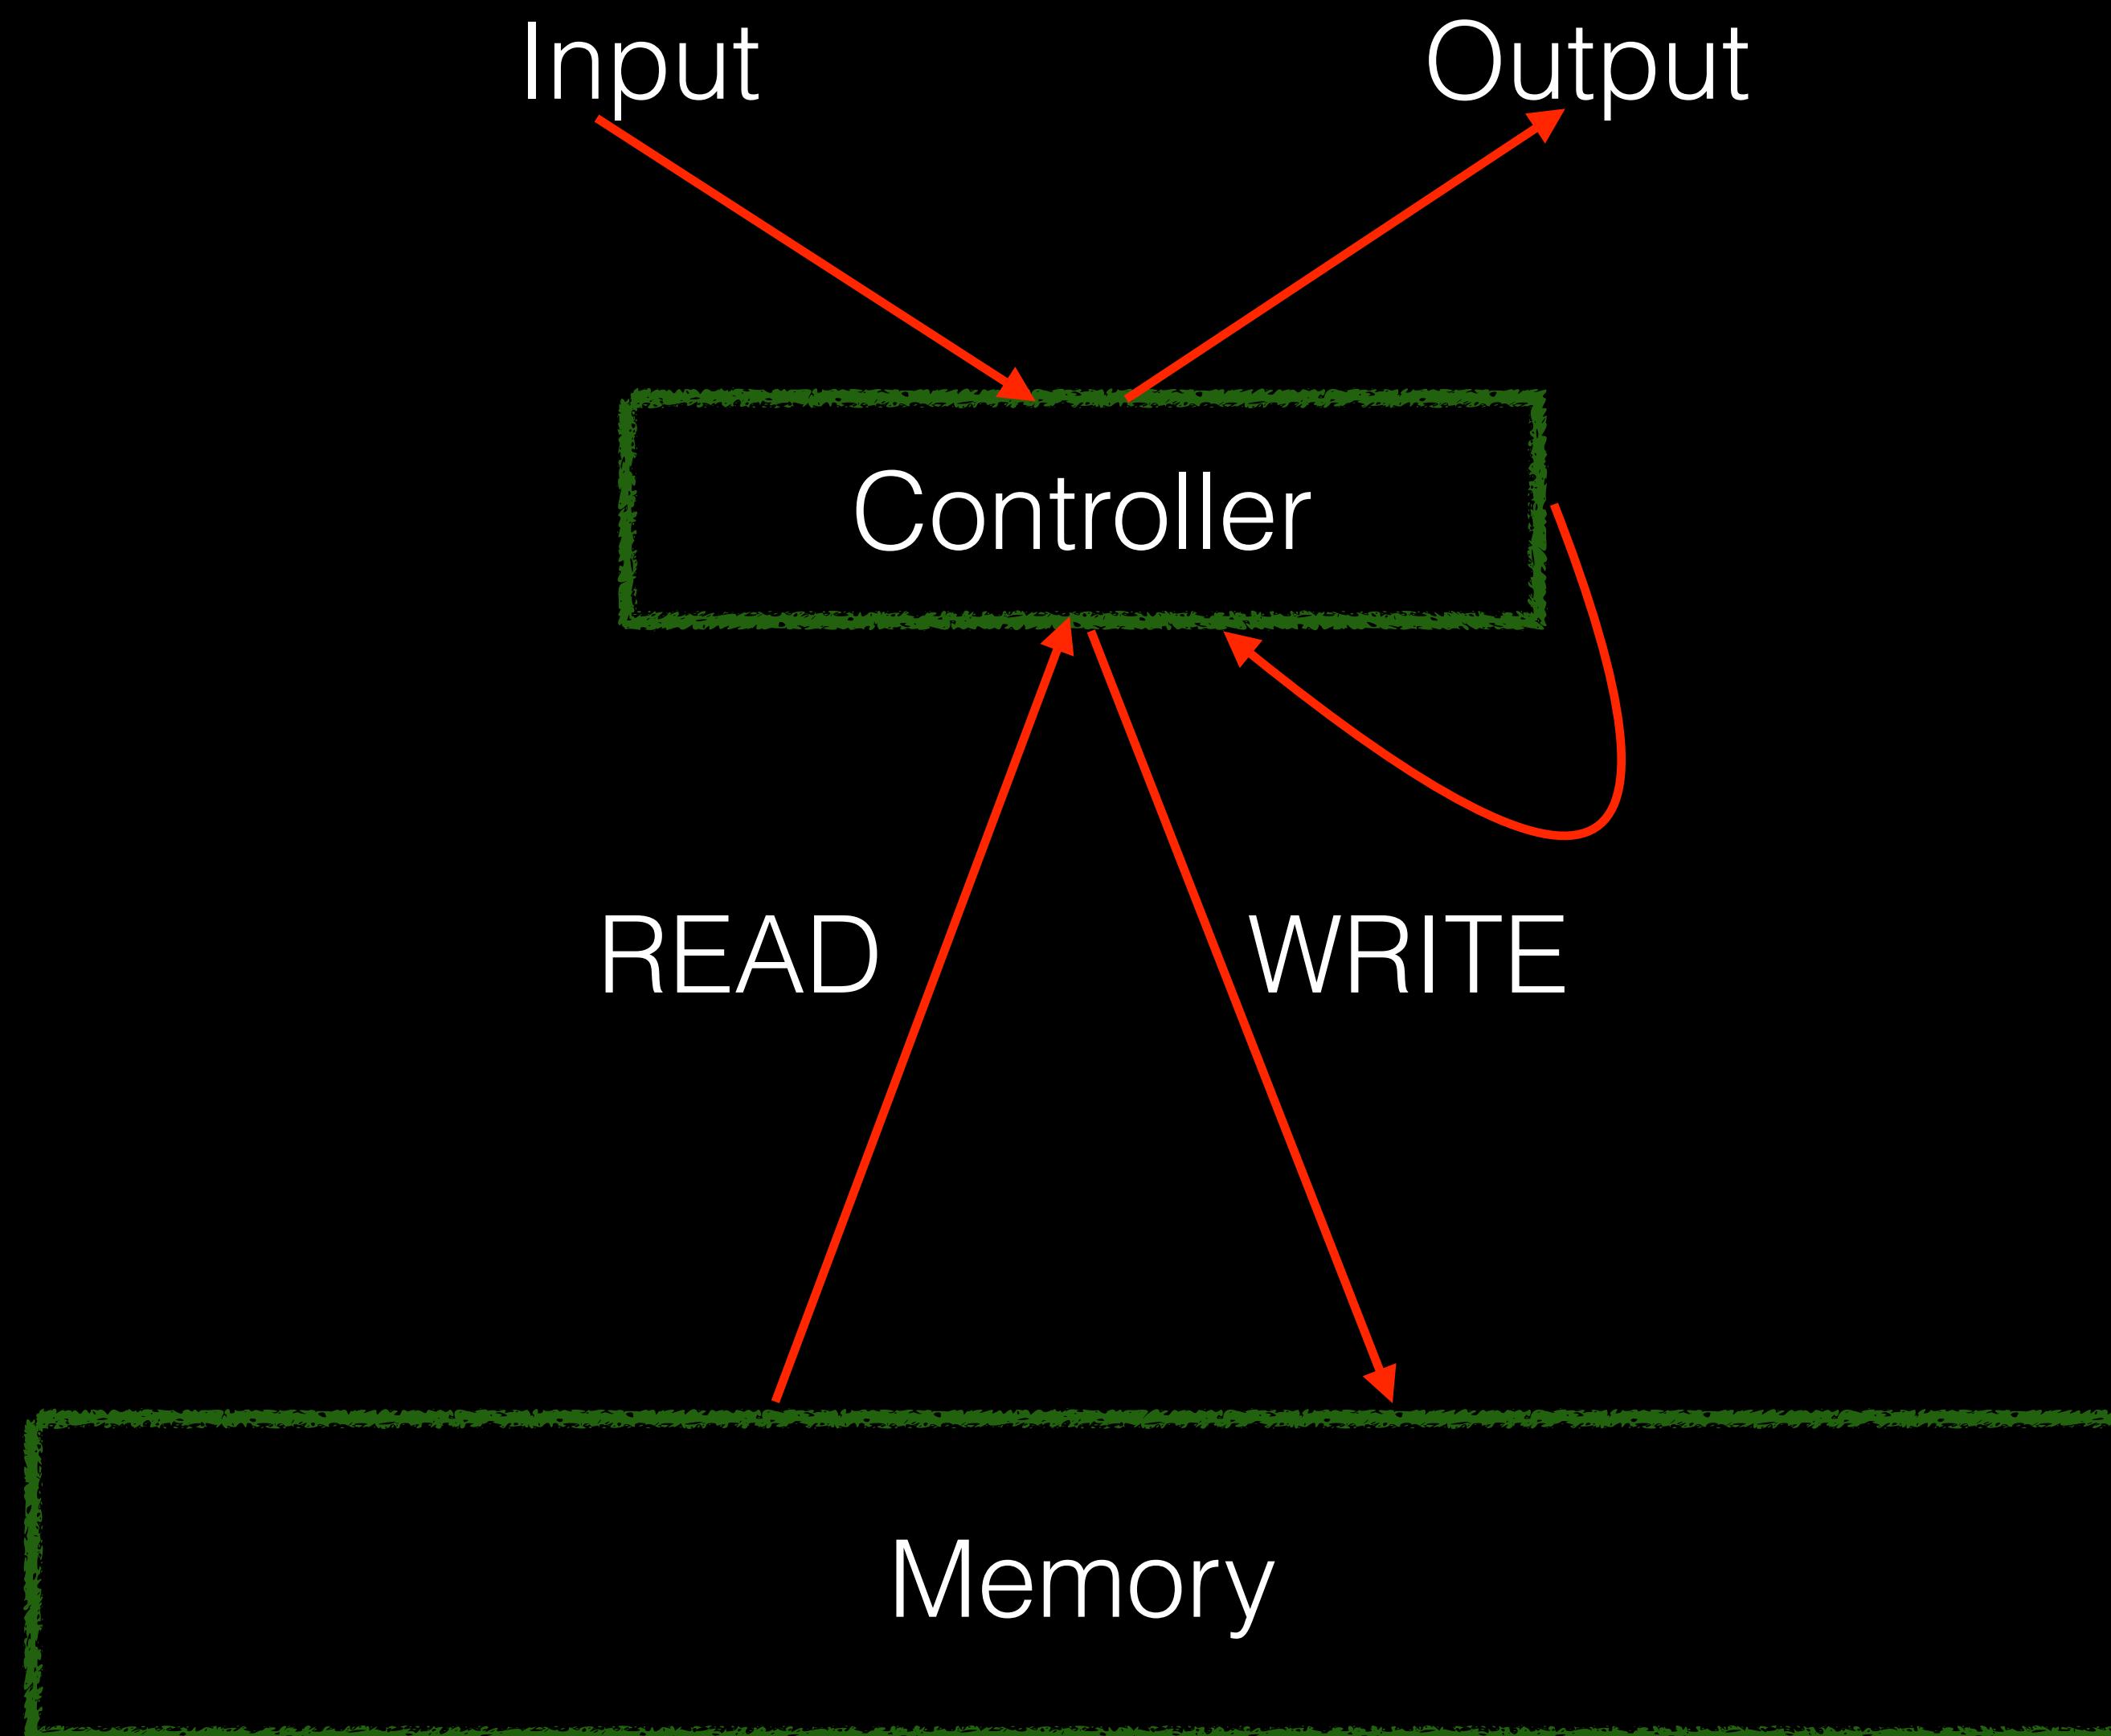
Revisit Attention Mechanism



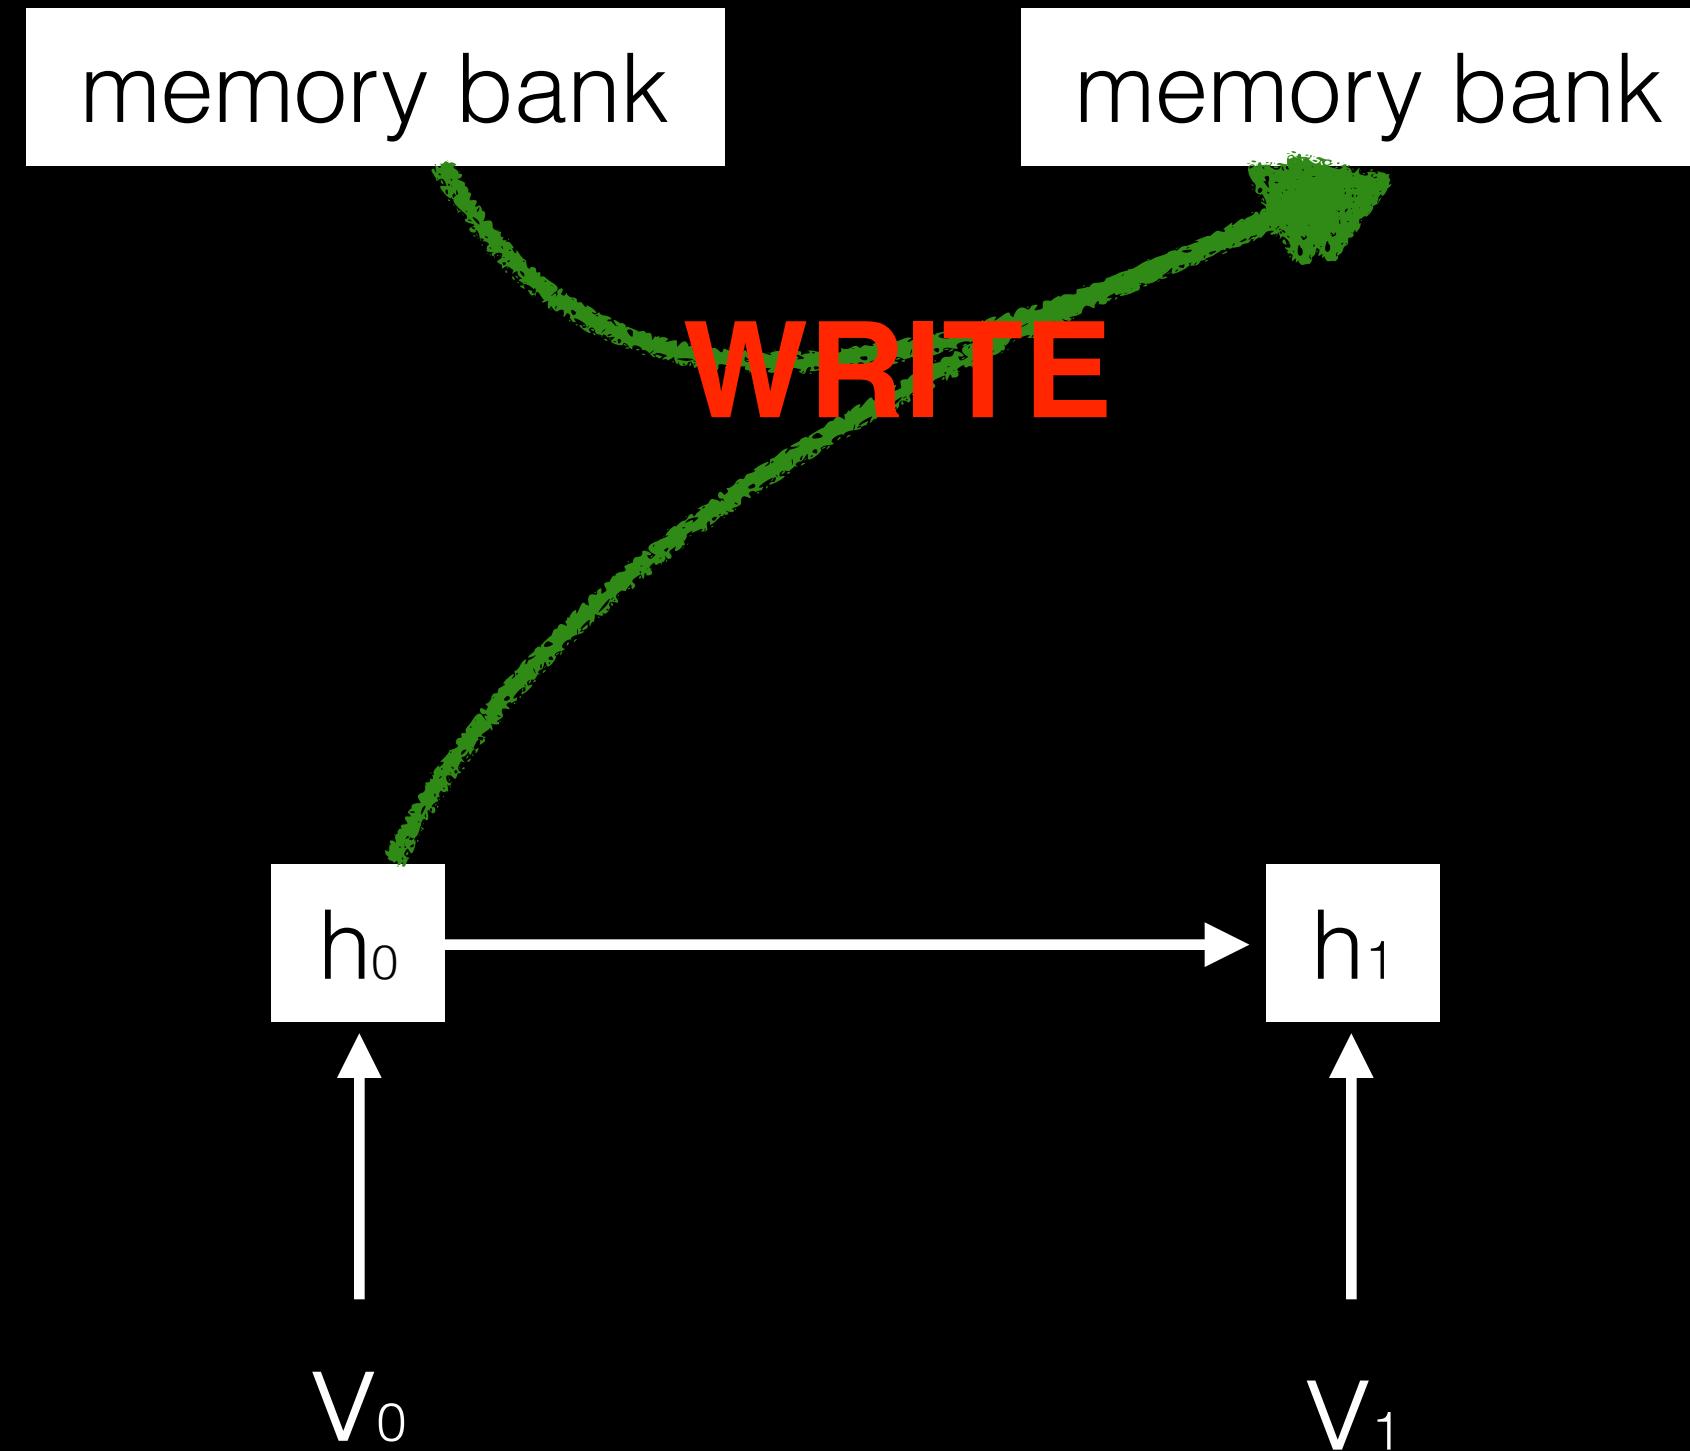
Revisit Attention Mechanism



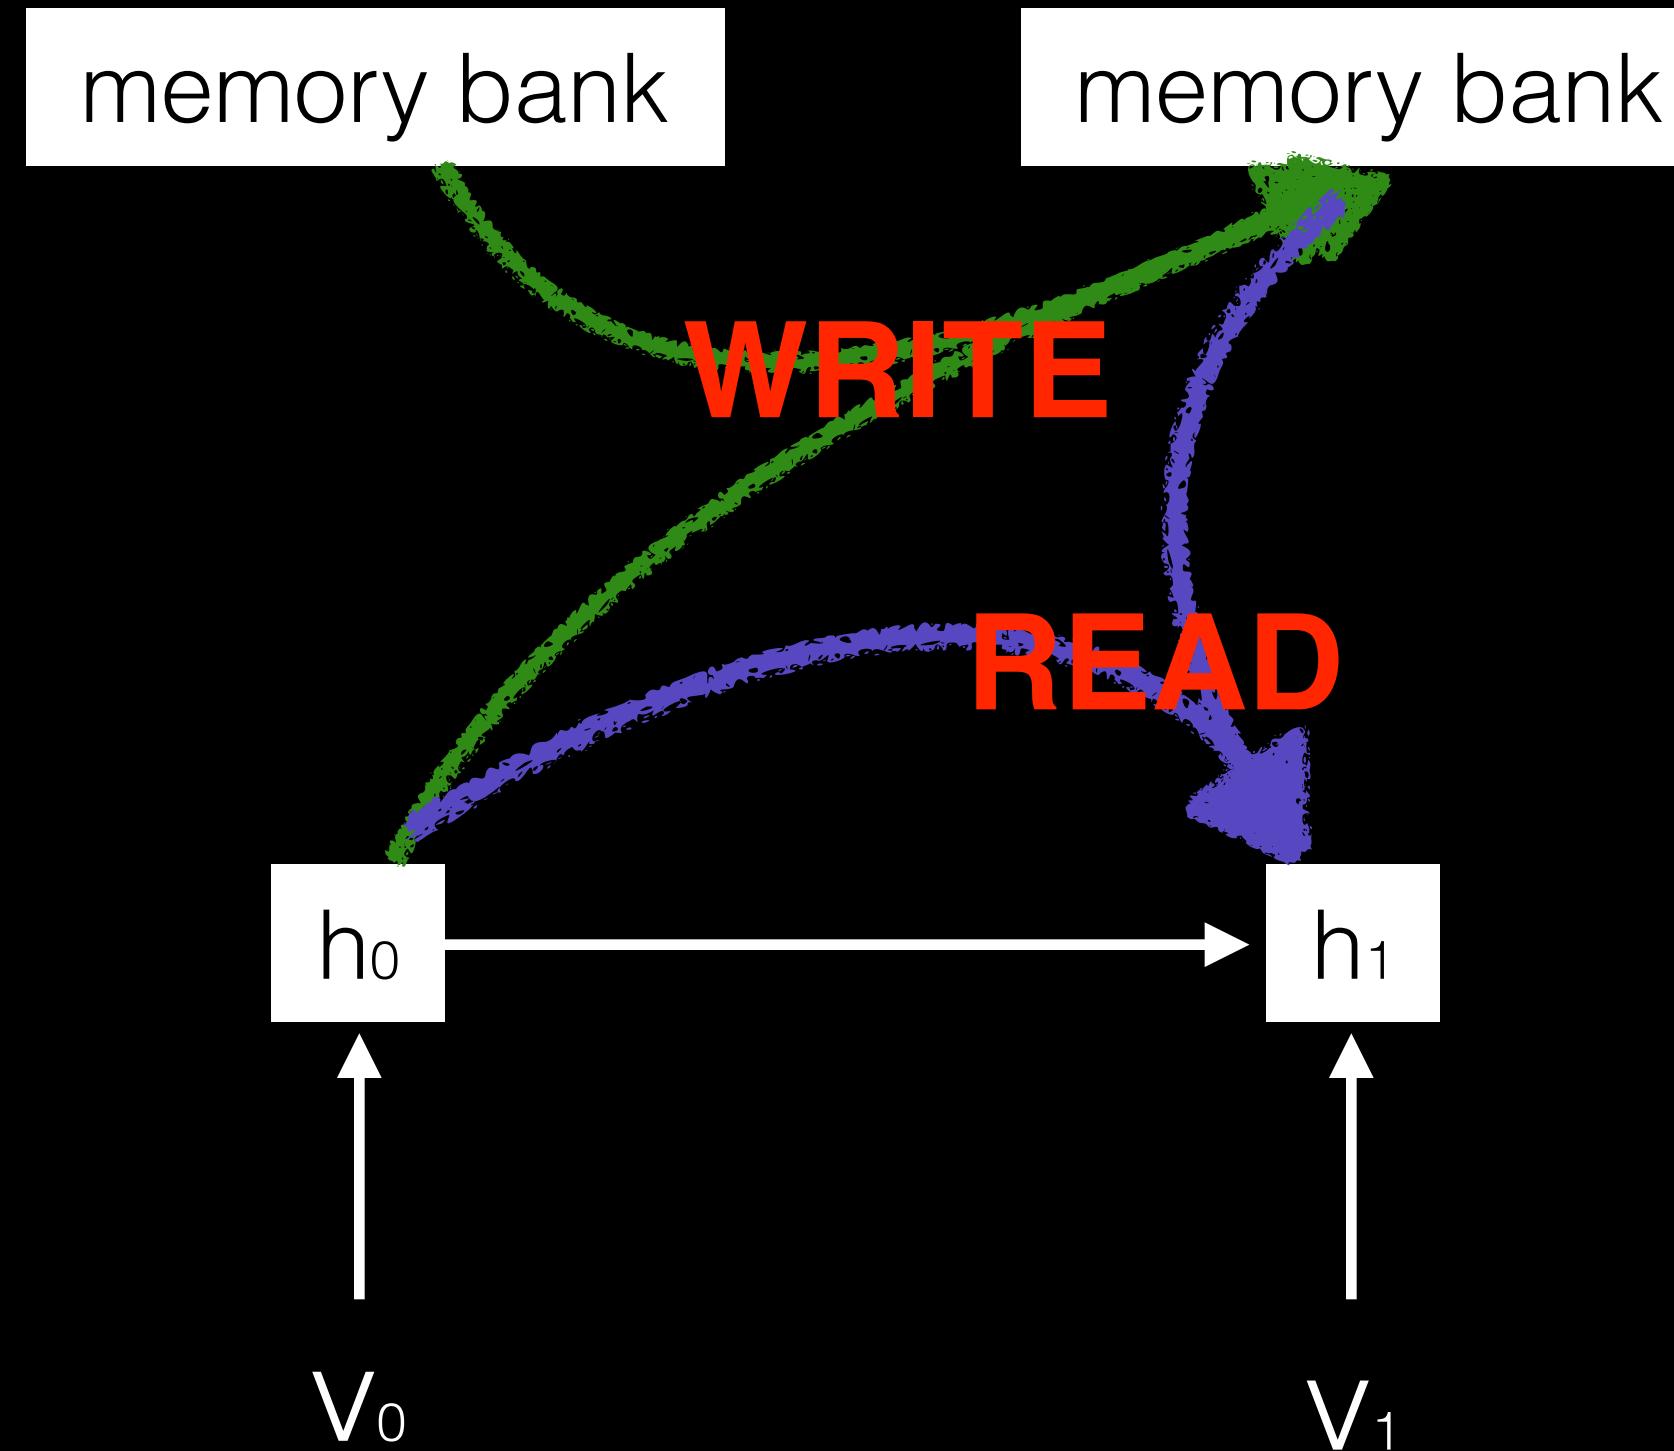
Differentiable Memory (Neural Turing Machines, Memory Networks, Stack-Augmented RNNs)



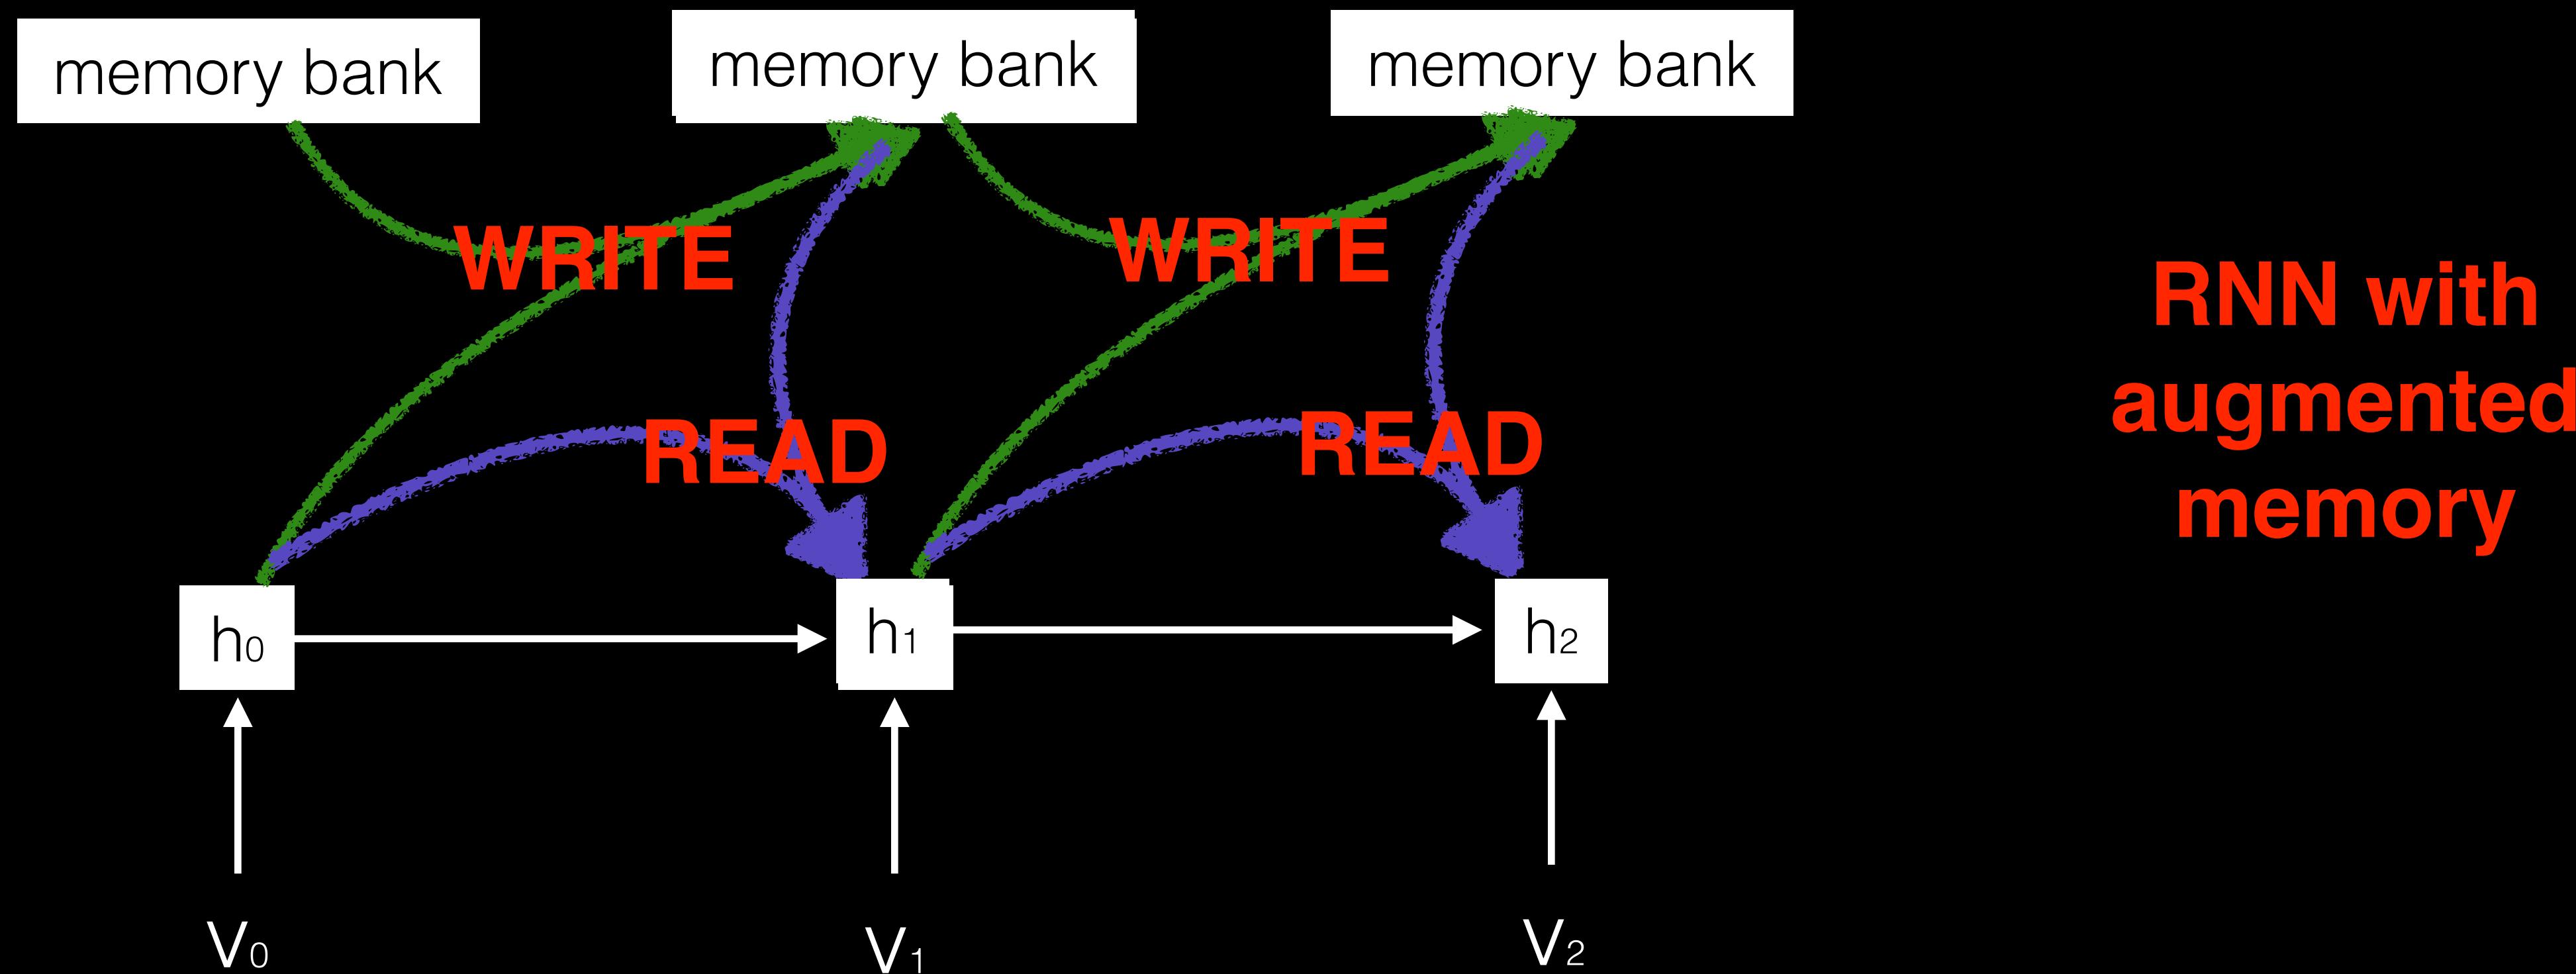
Differentiable Memory



Differentiable Memory



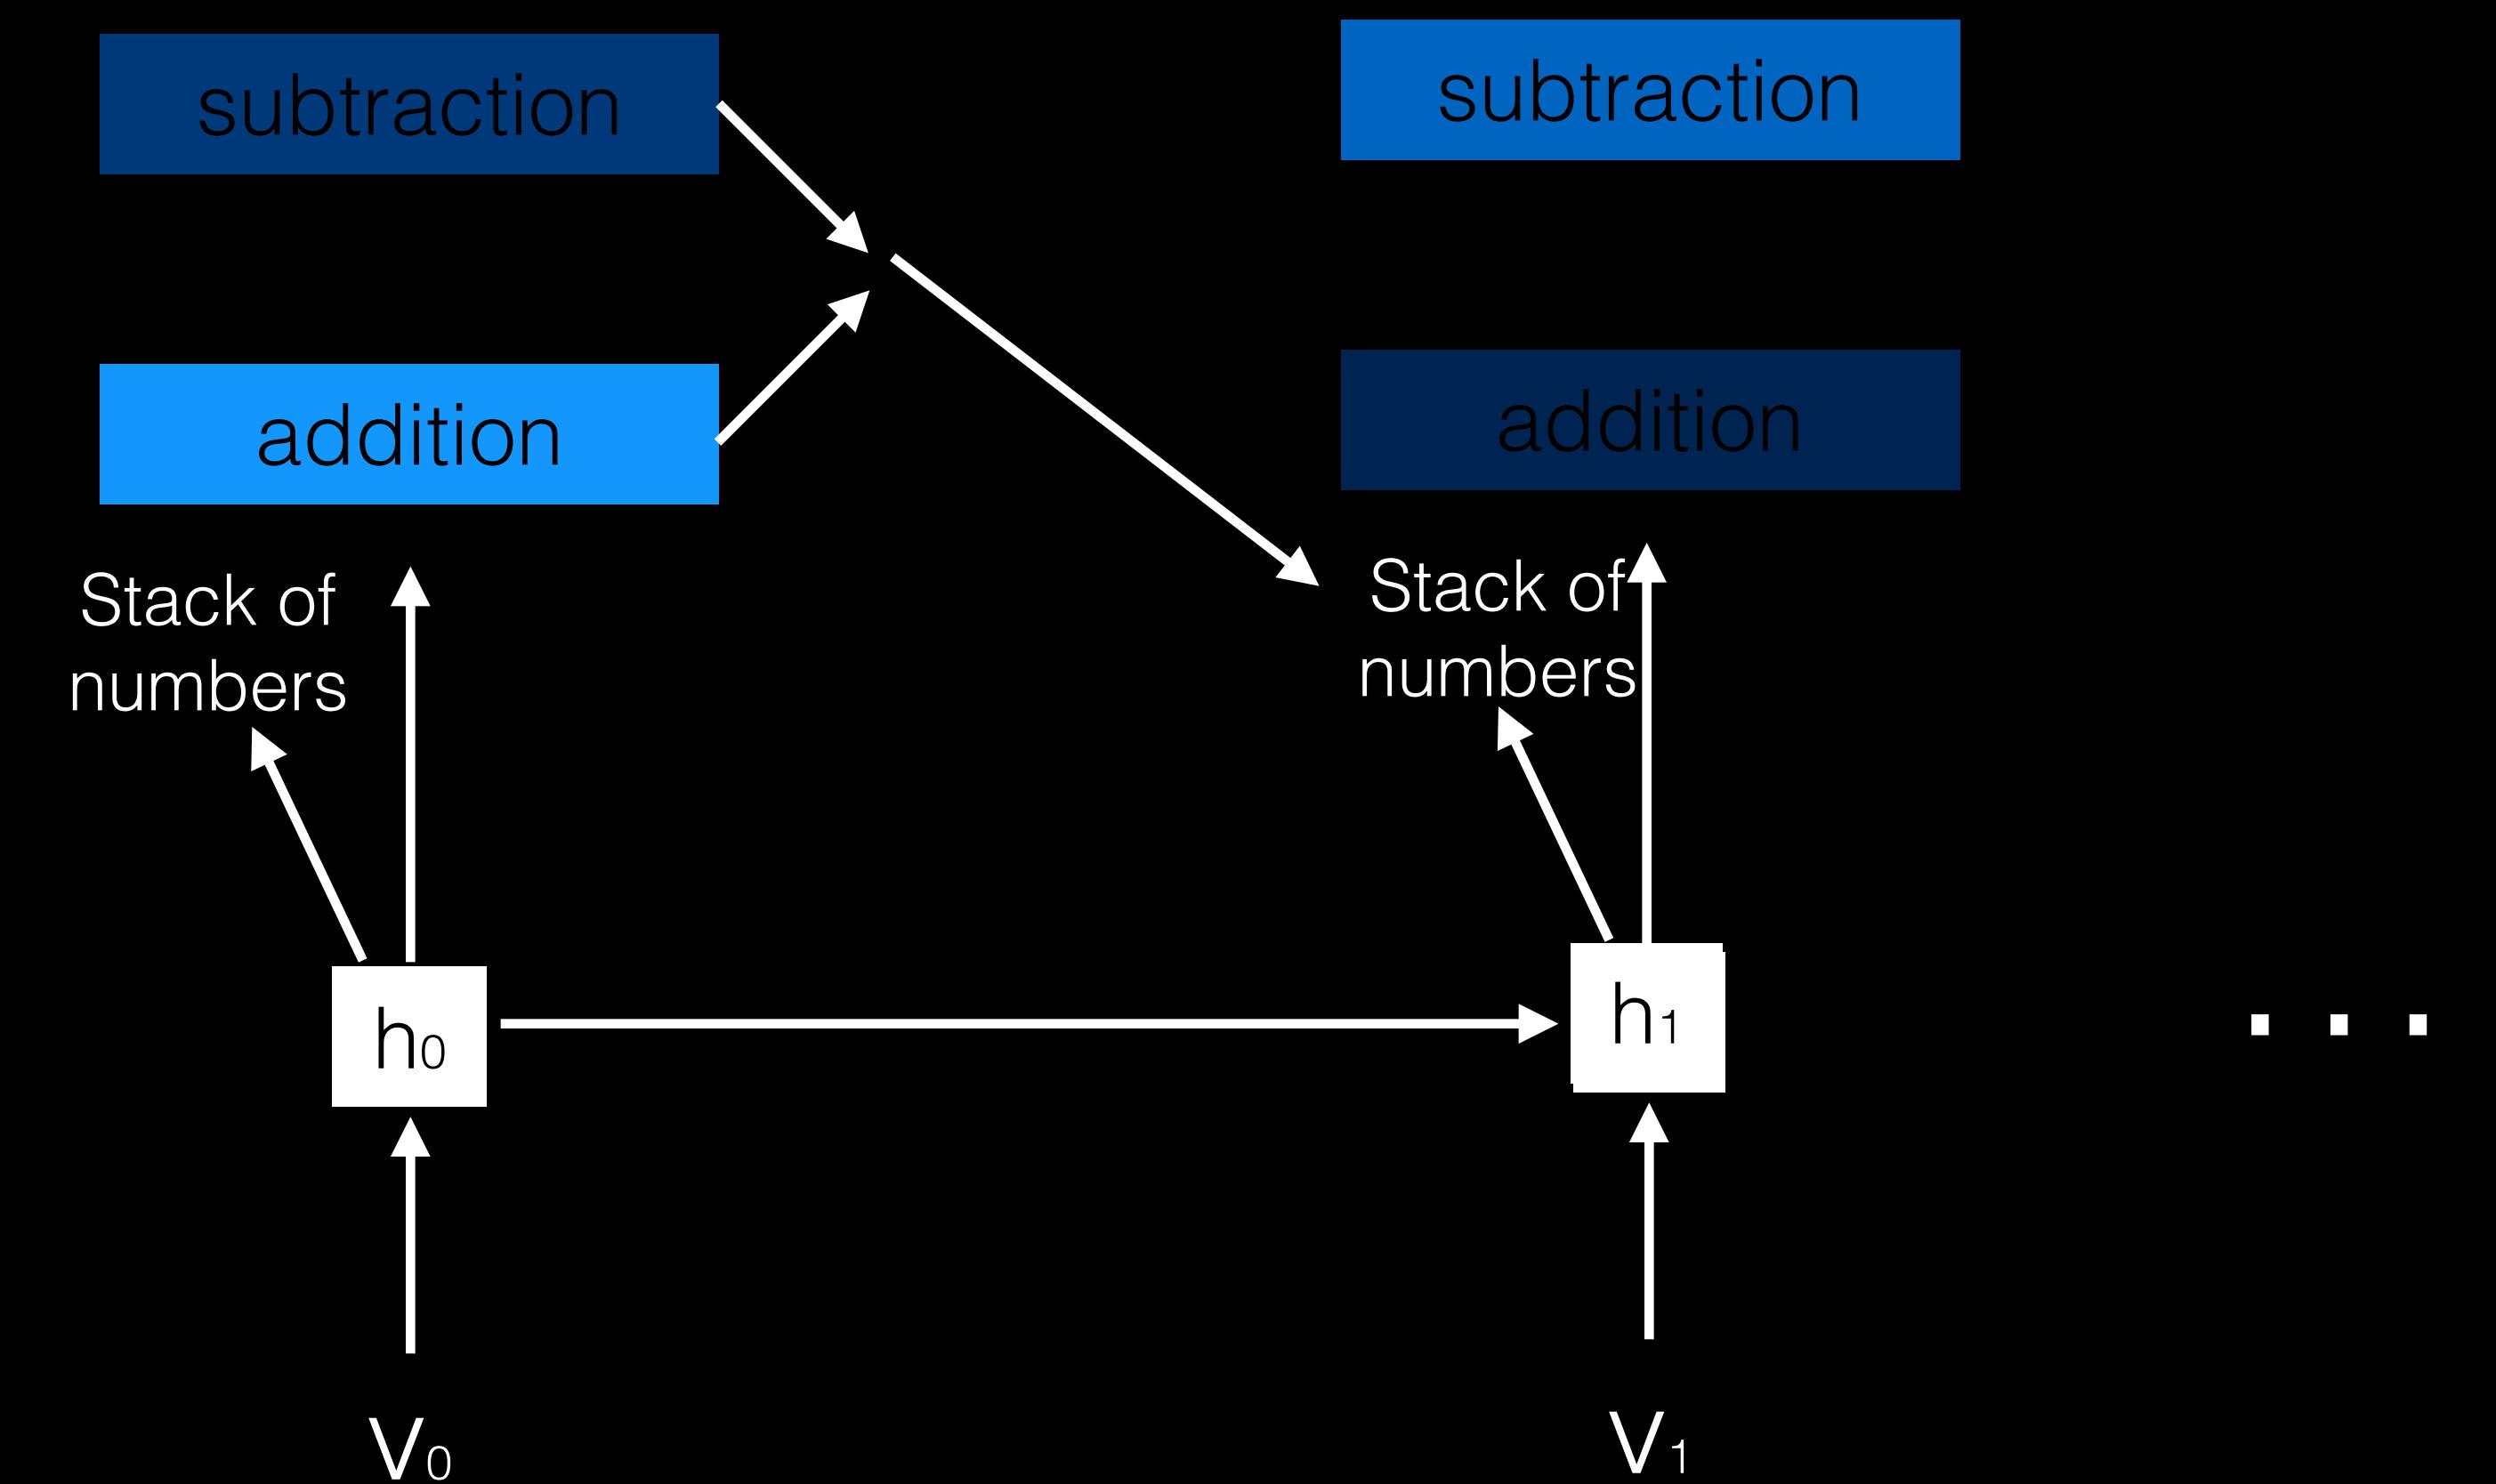
Differentiable Memory



RNN with augmented operations

- **Context:** The building was constructed in 2000 . . . It was destroyed in 2010 . . .
- **Question:** How long did the building survive?
- **Answer:** 10 years .

Neural Programmers



The Big Picture

- Sequence to sequence is an “End-to-end Deep Learning” algorithm
- It’s very general, so it should work with most NLP-related tasks **when you have a lot of data**
- If you don’t have enough data:
 - Consider dividing your problem into smaller problems, and train seq2seq on each of them.
 - **Train jointly with many other tasks**
- RNN with memory, or operation augmentation are exciting work in progress

Additional Reading

- Chris Olah's blog: Attention and Augmented Recurrent Neural Networks
- My own tutorials: <http://ai.stanford.edu/~quocle/tutorial2.pdf>
- Seq2seq in TensorFlow: <https://www.tensorflow.org/versions/r0.10/tutorials/seq2seq/index.html>

References

- Modeling
 - Sequence to Sequence with Neural Networks by Sutskever, Vinyals, Le. NIPS, 2014
 - Neural machine translation by jointly learning to align and translate by Bahdanau, Cho, Bengio. ICLR, 2015
 - Neural Turing Machines, by Graves, Wayne, Danihelka. arXiv, 2014
 - End-to-End Memory Networks by Sukhbaatar, Weston, Fergus. NIPS, 2015
 - Scheduled Sampling for Sequence Prediction with Recurrent Neural Networks, by Bengio, Vinyals, Jaitly, Shazeer. NIPS, 2015
 - Inferring Algorithmic Patterns with Stack-Augmented Recurrent Nets by Joulin and Mikolov. NIPS, 2015.
- Applications
 - Show and Tell: A Neural Image Caption Generator, by Vinyals, Toshev, Bengio, Erhan. CVPR, 2015
 - Grammar as Foreign Language by Vinyals, Kaiser, Koo, Petrov, Sutskever, Hinton. NIPS, 2015
 - Neural Conversational Model, by Vinyl and Le. ICML Workshop, 2015
 - A neural network approach to context-sensitive generation of conversational responses, by Sordoni, Galley, Auli, Brockett, Ji, Mitchell, Gao, Dolan, Nie. NAACL, 2015.
 - Neural responding machine for short-text conversation by Shang, Lu, Li. ACL, 2015.
 - Attention-Based Models for Speech Recognition. Chorowski, Bahdanau, Serdyuk, Cho, Bengio. NIPS, 2015
 - Listen, Attend and Spell. Chan, Jaitly, Le, Vinyals. ICASSP, 2016